arXiv:submit/7556205 [cs.CL] 6 May 2026

Beyond Vibe Coding: Intent Invariance and Structured Prompt Language

Wen G. Gong
Independent Researcher
wen.gong.research@gmail.com

Abstract

"Vibe coding" — the direct synthesis of code from natural language — has brought unprecedented velocity to AI application development. Yet velocity without verifiability is alchemy, not engineering. Natural language can be ambiguous: identical descriptions yield divergent implementations across invocations, with no inspectable intermediate state between human intent and machine execution. We call this the **intent gap**.

To close this gap, we present **Structured Prompt Language (SPL)**, a declarative intermediate representation (IR) that inserts a human-verifiable, machine-executable checkpoint into the agentic SDLC. SPL is a SQL-like language whose primitives (`GENERATE`, `EVALUATE`, `WHILE`, `CALL PARALLEL`) map directly to the semantics of LLM orchestration; its `.spl` files are readable by non-programmers, compilable to languages like Python/Go, and executable across LLM adapters without modification.

We formalize **Intent Invariance** as the correctness criterion: a pipeline satisfies it iff $T^{-1}(T(I)) \approx I$, where T is the forward compilation (NL→SPL→Code) and T^{-1} is the empirical reverse extraction (Code→NL spec). **Intent drift** $\Delta S = d(T^{-1}(T(I)), I)$ is the measurable quantity, operationalized by the `spl compare` command.

To measure ΔS, we performed 15 round-trip tests across 3 model tiers (Claude Sonnet, Gemini Flash, Qwen 3.6 Plus) and 5 real-world workflow patterns. Mean round-trip scores are 8.15, 7.90, and 7.30 respectively, with individual scores ranging from 6.0 to 9.0. The primary empirical finding concerns not the models but the pipeline: intent drift traces to ambiguous Mermaid topology diagrams at the human checkpoint, not to model capability — isolating where in the SDLC the vibe component resides and where human judgment is most needed. Code released under Apache 2.0 at `https://github.com/digital-duck/SPL.py`.

1 Introduction

Agentic AI development today has no stable intermediate representation. The dominant paradigm — "vibe coding," the stochastic synthesis of code from natural language — is an open loop: the practitioner describes intent, the LLM generates code, and the output is accepted or discarded. There is no inspectable state between human intention and machine execution, no formal criterion for correctness, and no reproducible path from description to working system.

We identify three structural failure modes:

Intent Drift. A specification in natural language that seems precise to its author contains latent degrees of freedom that the LLM resolves stochastically. The generated code satisfies

Preprint.

a plausible interpretation of the intent — not necessarily the author's. With no intermediate checkpoint, this drift is invisible until runtime.

Coordinate Dependence. Generated code is runtime-specific. The same logical intent must be re-encoded from scratch for each target framework — a brittle, expensive process that treats the execution environment as fundamental rather than incidental.

Irreproducibility. The same prompt can produce different code on each invocation, dependent on LLM choice. There is no stable artifact to version-control, audit, or transfer between practitioners.

We address all three with a structured intermediate representation and a measurable correctness criterion:

1. **Agentic SDLC Pipeline** (§3): A structured pipeline that separates topological intent from syntactic realization by inserting explicit human checkpoints. The round-trip test protocol operationalizes intent drift ΔS via the `spl compare` command with LLM-powered semantic scoring (0–10 scale).
2. **Structured Prompt Language (SPL)** (§4): A declarative, SQL-like IR whose primitives — `SELECT/GENERATE` (SQL-style data selection and LLM generation, extending SQL's `SELECT ... INTO` pattern), `CALL` (invoke a tool or API deterministically), `EVALUATE` (semantic branching), `WHILE` (iterative refinement), `CALL PARALLEL` (concurrent dispatch) — together model LLM orchestration semantics natively. As the stable, human-verifiable checkpoint between natural language and executable code, SPL enforces a hard separation between the syntactic layer (control flow, verifiable by tooling) and the semantic layer (prompt content, only observable at runtime) — a decomposition that makes intent drift locatable.
3. **Round-trip Tests** (§5): 15 tests across 3 model tiers and 5 workflow patterns. The pipeline stress-test reveals that intent drift enters at the human topology checkpoint, not at model capability limits — a process-structural finding with direct implications for where practitioners should invest verification effort.

2 Related Work

Vibe coding and epistemic grounding. Vibe coding has been formalized as a socio-technical framework in which LLMs excel at artifact generation but lack bounded, verifiable structure [1]. The transition to agentic software development requires *epistemic grounding* — explicit documents encoding hard constraints and validity invariants [1]. Agentic Entropy [2] names the process-level drift where autonomous updates erode global architectural intent; SPL's checkpoint pipeline directly addresses this. SPL occupies a pragmatic middle ground between prose and formal verification tools: more deterministic than the former, more accessible than the latter.

Declarative orchestration and intent drift measurement. DSPy [3] takes a declarative approach to prompt *optimization*; SPL takes a declarative approach to workflow *specification*. Unlike DSPy, which optimizes prompt weights, SPL enforces structural topology verification via a human-approved IR checkpoint — the `.spl` file is the invariant, not the prompt. LMQL [4] constrains at the token level; SPL constrains at the workflow level. A declarative LLM orchestration language [5] reports 60% development time reduction and 3× deployment velocity. ReusStdFlow [6] proposes dual-form IRs (graph + semantic summary) prioritizing reusability; SPL prioritizes intent integrity and cross-runtime execution. Recent work on Mermaid-based visual workflow generation [7] and workflow benchmarking [8] explores related visual IR patterns; SPL distinguishes itself by coupling the visual topology with a compilable declarative language and a formal round-trip test for intent preservation. SPL is an IR specification language — its `.spl` files compile to runtime frameworks such as LangGraph or PocketFlow; those frameworks are execution environments, not specification languages. Kuhn et al. [9] define Semantic Entropy as single-response LLM uncertainty; our intent drift measure ΔS (§3.2) quantifies *round-trip preservation across a full compilation pipeline* — a distinct empirical concept. DeepContext [11] formalizes intent drift

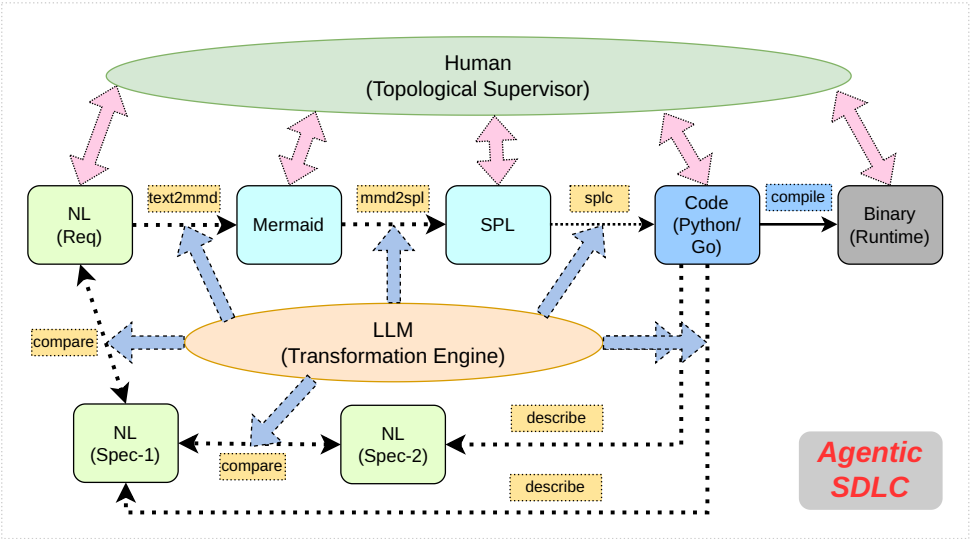

Figure 1: Agentic SDLC Pipeline

as embedding-space displacement across dialogue turns; our ΔS measures drift across a *compilation round-trip*, not a dialogue session — a structurally distinct scope.

Model-Driven Engineering (MDE) and round-trip engineering. The principle of generating executable artifacts from a higher-level model has a long history in software engineering (UML, Model-Driven Architecture). Classic MDE pipelines suffered a persistent synchronization problem: manual edits to generated code would diverge from the model, breaking round-trip fidelity. SPL addresses this for the LLM era: by treating the `.spl` script as the authoritative model and relying on LLM *regeneration* rather than manual code editing, resynchronization becomes a generation step rather than a maintenance burden. SGLang [12] pursues structured generation at the execution layer for throughput; SPL pursues structured specification at the design layer for auditability — complementary concerns at different points in the pipeline.

3 From Vibe to Verification

3.1 The Agentic SDLC Pipeline

Traditional vibe coding jumps directly from natural language to imperative code, collapsing two orthogonal concerns into a single unconstrained generation step: *topological intent* (what nodes and edges should exist in the workflow) and *syntactic realization* (how should each node be implemented). The LLM resolves both stochastically, with no checkpoint between them. SPL's Agentic SDLC pipeline (Figure 1) separates these concerns by inserting explicit human checkpoints between verified states.

Figure 1. The Agentic SDLC pipeline. The Human (Topological Supervisor, top) validates states at each checkpoint (pink bidirectional arrows). The LLM (Transformation Engine, center) executes mechanical translations between states (blue arrows). The forward path (top row) transforms natural language requirements through Mermaid, SPL, and compiled code to a running binary. The reverse path (bottom row) extracts specifications from the compiled artifact back to natural language, enabling round-trip testing via `spl compare`. The dotted lines trace the round-trip test loop.

We locate the "vibe" step precisely within this pipeline: it is the NL → Mermaid → SPL leg, where the LLM transforms free-form intent specification into a structured visual topology and then into executable declarative logic. This leg is where intent drift enters the pipeline. The human checkpoint sits between the LLM-generated Mermaid diagram and the SPL

script — the one moment where the designer can validate that the topology matches their intent *before any code is generated*.

The pipeline has two roles for the human and two roles for the LLM:

- **Human as topology validator**: (1) reviews the Mermaid diagram before any SPL or code is generated; approves the SPL script before compilation; (2) validates the running binary. These are state validations, not translation tasks.
- **LLM as transformation engine**: (1) executes text2mmd, mmd2spl, splc, and describe — all mechanical translations between verified states; (2) the same LLM is also invoked interactively to diagnose and repair failures at each checkpoint, which is vibe coding's strength now channeled through a structured process.

This separation reduces reliance on "prompt engineering" as a primary activity: instead of crafting better prompts, practitioners validate better states.

3.2 Intent Invariance and Round-trip Testing

We formalize the correctness criterion for the Agentic SDLC using two definitions.

Definition 1 (Intent Invariance). A pipeline satisfies Intent Invariance iff:

$$T^{-1}(T(I)) \approx I$$

where T: NL $\rightarrow$ SPL $\rightarrow$ Code is the forward compilation path and T^{-1}: Code $\rightarrow$ describe $\rightarrow$ NL spec is the reverse extraction path. A pipeline achieves Intent Invariance when the reverse-engineered specification sufficiently matches the original intent; intent drift ΔS measures the degree of deviation.

Definition 2 (Intent Drift). The intent drift of a pipeline is:

$$\Delta S = d\big(T^{-1}(T(I)),\ I\big)$$

where $d(\cdot,\cdot)$ is a semantic distance measured by the `spl compare` command[1] — an LLM judge scoring alignment across structure, logic, and quality dimensions (0–10 scale, higher means less drift).

T is not injective, so no true inverse exists. We use T^{-1} to denote empirical reverse extraction (code $\rightarrow$ natural language spec), not a mathematical inverse. The distance d is operationalized by an LLM judge with no closed-form derivation; it does not formally satisfy metric axioms. Intent drift ΔS is a practical measurement tool, not a theorem. Its value is that it makes the intent gap *locatable* — the score signals whether the round-trip preserved intent, and the judge's commentary identifies which layer (structure, logic, or quality) introduced drift.

3.3 Round-trip Test Protocol

The round-trip test (S1–S6) operationalizes ΔS measurement for the Agentic SDLC:

Step	Tool	Input $\rightarrow$ Output
S1	`spl3 splc describe --include-docs`	Original code + README $\rightarrow$ `S1-*-spec.md`
S2	`spl3 text2mmd`	Spec $\rightarrow$ `S2-*.mmd` (Mermaid diagram)
Human	**Visual review and approval**	**Verify nodes and edges manually**
S3	`spl3 mmd2spl`	Diagram $\rightarrow$ `S3-*.spl` (SPL script)
S4	`spl3 splc compile --llm`	SPL $\rightarrow$ `S4-*.py` (compiled code)
S5	`spl3 splc describe`	Reconstructed code $\rightarrow$ `S5-*-spec.md`
S6	`spl3 compare`	S1 spec vs S5 spec $\rightarrow$ round-trip score ΔS

[1] `spl3` and `spl` are both CLI aliases for the same SPL runtime.

Step	Tool	Input → Output

** S4 targets deterministic transpilation; `--llm` fallback is used for novel patterns not yet in the pattern library, reintroducing limited stochasticity at that step. All five recipes (R1–R5) matched established patterns (`react`, `self_refine`, `map_reduce`) and used the deterministic path; the `--llm` fallback was not invoked in any of the 15 test runs.*

The S2 human checkpoint is a key intervention: once the Mermaid topology is approved, all downstream LLM steps (S3–S5) operate within those topological constraints. A wrong or ambiguous topology at S2 propagates through the entire pipeline.

4 SPL: The Logical Intermediate Representation

4.1 Design Thinking

SPL is guided by 4 design considerations:

Auditable IR. The `.spl` file is the checkpoint between intent and execution. Practitioners review and approve the SPL before any code runs. Intent drift is caught at design-time, not at runtime.

LLM-native primitives. `GENERATE` (LLM call), `EVALUATE` (semantic branching), `EXCEPTION WHEN` (typed recovery), `CALL PARALLEL` (concurrent dispatch) — these are first-class constructs modeling the semantics of agentic systems natively, not library calls.

Runtime agnostic. The same `.spl` file compiles to languages like Python/Go and executes across LLM adapters without modification. Provider, runtime, and deployment environment are resolved at execution time — the `.spl` is the invariant; the runtime is a coordinate choice.

SQL as foundation. SQL has been the lingua franca of data engineering for 50 years — proven abstractions that outlasted every storage engine beneath them. SPL builds on this directly: `WORKFLOW` extends stored procedures; `GENERATE` extends `SELECT`; `EVALUATE` extends `CASE`. The paradigm is legible to data professionals familiar with SQL.

4.2 Core Primitives

We illustrate SPL syntax through the self-refine workflow pattern, which exercises the core primitives in a single coherent example.

```
WORKFLOW self_refine
  INPUT  @task TEXT, @max_rounds INT := 3
  OUTPUT @result TEXT
DO
  GENERATE Draft(@task) INTO @draft;
  @round := 0;
  WHILE @round < @max_rounds DO
    GENERATE Critique(@draft, @task) INTO @feedback;
    EVALUATE @feedback
      WHEN is_satisfactory THEN RETURN @draft;
      ELSE
        GENERATE Refine(@draft, @feedback) INTO @draft
          EXCEPTION WHEN refine_failed THEN
            RETURN @draft WITH status = 'fallback';
          END;
        @round := @round + 1;
    END;
  END;
  RETURN @draft;
END
```

The semantic layer — the natural language instructions inside each `CREATE FUNCTION` body — is where the practitioner expresses *what* each step should do:

```
CREATE FUNCTION Critique(draft TEXT, task TEXT)
RETURNS TEXT AS $$
You are a quality reviewer. Given the task: {task}
And the current draft: {draft}

Assess whether the draft fully satisfies the task. Respond with YAML only:
  verdict: satisfactory
or:
  verdict: needs_revision
  feedback: <specific issues to address>
$$;
```

The WORKFLOW body (syntactic layer) controls *when* each function is called; the `$$...$$` body (semantic layer) controls *what* the LLM does. This separation is what makes SPL auditable: structure is verifiable by tooling, prompt quality is inspectable by humans.

The `.spl` file is human-readable and requires no language-specific framework knowledge to review. Additional workflow patterns spanning the full primitive set are available in the R1–R5 recipes (§5) and the extended cookbook in the project repository.

4.3 The Syntactic/Semantic Layer Decomposition

SPL is designed around a hard boundary between two fundamentally different sources of variance in agentic systems.

Layer 1 — Syntactic (WORKFLOW body): Control flow, node wiring, variable routing, loop guards, and exception handlers, expressed in SPL keywords (`WHILE`, `EVALUATE`, `CALL`, `RETURN`). Fully deterministic; SPLc translates them mechanically with zero semantic ambiguity. Bugs at this layer are *mechanically detectable* — caught by the SPL parser.

Layer 2 — Semantic (CREATE FUNCTION bodies): The prose instructions inside `$$...$$` delimiters — the actual prompt text governing LLM behavior at runtime. Inherently stochastic: different phrasings of the same intent produce measurably different outcomes, even when the surrounding WORKFLOW structure is identical. Bugs at this layer are *only observable through runtime behavior*.

Without an explicit IR, both layers are entangled inside imperative code: a structural bug (wrong branching logic) and a prompt-quality bug (wrong instruction) are indistinguishable at a glance. SPL makes the boundary explicit — the WORKFLOW body is the deterministic skeleton; `$$...$$` bodies are the stochastic prompts. Intent drift at Layer 1 is fixable with tooling; intent drift at Layer 2 requires human judgment. SPL tells you which kind you have.

5 Round-trip Tests

To validate intent preservation across model tiers and workflow complexity, we run the S1–S6 protocol (§3) on 5 workflow recipes × 3 LLM backends — 15 tests total. The goal is not to rank models but to stress-test where ΔS varies and trace the structural cause.

5.1 Test Design

Workflow recipes. Five workflows from the PocketFlow cookbook [10], selected to span control-flow complexity and stress-test specific SPL constructs (`WHILE`, `EVALUATE`, `CALL PARALLEL`, nested loops):

Recipe	Pattern	Complexity
`R1-agent`	ReAct loop (decide → search → accumulate)	Medium
`R2-rag`	Linear pipeline (embed → retrieve → generate)	Low
`R3-judge`	Evaluator-optimizer (generate → score → revise)	Medium
`R4-thinking`	Chain-of-thought (think → answer, nested loops)	Medium

Recipe	Pattern	Complexity
`R5-research`	Recursive map-reduce (plan → parallel search → synthesize)	High

LLM backends. Three models selected from a candidate pool of seven based on syntactic and semantic SPL generation quality, representing three distinct provider tiers:

	Short name	Adapter	Model ID
M1	`sonnet`	`claude_cli`	`claude-sonnet-4-6`
M2	`gemini`	`openrouter`	`google/gemini-3-flash-preview`
M3	`qwen`	`openrouter`	`qwen/qwen3.6-plus`

The S6 judge is fixed at `claude-opus-4-6` across all runs to eliminate cross-judge variance; using a Claude-family judge alongside M1 introduces a potential intra-provider self-preference risk, a limitation for future work. Each configuration is run once; results are indicative rather than statistically definitive.

Metrics. The S6 judge scores each run on three dimensions (0–10, higher means less drift), aggregated into an overall score: Structure (organization, flow, architecture alignment), Logic (decision points, control flow, completeness), Quality (sophistication, robustness, prompt conventions).

5.2 Score Matrix

Recipe	Complexity	M1: Claude	M2: Gemini	M3: Qwen
R1-agent	Medium	8 / 8 / 8 / **8.0**	8 / 8 / 8 / **8.0**	8 / 9 / 8 / **8.5**
R2-rag	Low	8 / 8 / 8 / **7.75**	8 / 8 / 8 / **8.0**	8 / 8 / 8 / **8.0**
R3-judge	Medium	8 / 8 / 9 / **8.0**	7 / 8 / 7 / **7.5**	8 / 7 / 7 / **7.0**
R4-thinking	Medium	8 / 9 / 8 / **8.0**	8 / 8 / 7 / **7.5**	7 / 6 / 6 / **6.0**
R5-research	High	9 / 8 / 9 / **9.0**	8 / 9 / 8 / **8.5**	7 / 6 / 7 / **7.0**
Mean		**8.15**	**7.90**	**7.30**

*Structure / Logic / Quality / **Overall** from `spl compare` (S6 judge: claude-opus-4-6).*

5.3 Observations and Structural Failure Analysis

Mean scores form a gradient (8.15 → 7.90 → 7.30) across tiers. Within each tier, scores vary by recipe: R2-rag (lowest workflow complexity) scores 7.75 for Claude while R5-research (highest complexity) scores 9.0 for Claude and 8.5 for Gemini. Score and workflow complexity are not monotonically related across this small dataset.

In the lower-scoring cases, we can trace the specific divergence between S1 and S5 specs to the S2 diagram step. R2/Claude (7.75): the S2 diagram did not specify the embedding algorithm; S4 compiled code used a different algorithm family (TF-IDF bag-of-words) than the original (OpenAI dense vectors). R3/Gemini and R4/Qwen: the S2 diagram represented the WHILE loop exit as a graph-edge shorthand rather than an explicit termination condition; different models resolved the ambiguity differently at S3. In each of these cases, the S5 spec faithfully described what the compiled S4 code did — the divergence from S1 was introduced before S4, in the S2 diagram or S3 SPL.

This localization is a direct consequence of the IR structure. Without an inspectable intermediate state, the same failures would be indistinguishable from model error. SPL's pipeline makes the failure mode *addressable*: fixing the S2 diagram or the S3 SPL repairs the drift without touching the model or the prompt.

5.4 Limitations

This study establishes a methodology for quantifying intent drift ΔS within the structured SPL pipeline. We benchmark the round-trip integrity of the SPL compilation path (NL $\rightarrow$ SPL $\rightarrow$ Code) but do not include a direct "vibe coding" baseline (NL $\rightarrow$ Code) in this work. We prioritize the structural analysis of *where* drift enters the pipeline — identifying the S2 topology checkpoint as the critical failure mode — which is only locatable when an intermediate representation exists. While a comparative baseline would quantify the absolute magnitude of drift reduction, our primary contribution here is the *diagnostic capability* of SPL: enabling practitioners to pinpoint whether drift stems from topological ambiguity (S2) or semantic realization (S4). The `spl3 vibe` command (direct NL $\rightarrow$ Code generation) is implemented and will be used in a forthcoming ablation study. We also acknowledge that using a Claude-family judge (S6) alongside a Claude model (M1) introduces a potential intra-provider self-preference risk. We leave the quantitative baseline comparison and cross-judge validation for future work. More broadly, SPL is presented as an initial step toward verifiable agentic development, not a complete solution; the language, grammar, and toolchain are expected to evolve as the field matures and as community feedback and ablation results clarify where the IR adds the most value.

6 Conclusion

We have presented **Structured Prompt Language (SPL)** as a declarative IR that addresses the intent gap in agentic software development. By inserting a human-verifiable, machine-executable checkpoint, SPL separates the deterministic syntactic layer (control flow, verifiable by tooling) from the stochastic semantic layer (prompt content, requiring runtime observation) — making intent drift locatable.

Round-trip testing, operationalized by the `spl compare` command and measured by intent drift ΔS, provides a methodology for assessing whether a workflow preserves the original intent through compilation and back. Across 15 round-trip tests spanning 3 model tiers and 5 workflow patterns, mean round-trip scores are 8.15 (Claude), 7.90 (Gemini), and 7.30 (Qwen), with individual scores ranging from 6.0 to 9.0. In the lower-scoring cases, the divergence between S1 and S5 specs was traceable to specific ambiguities in the S2 diagram rather than to model capability differences on the same recipe.

Critically, the diagnostic capability of SPL — the ability to localize drift to the topology diagram (S2) or to the semantic prompt layer (S4) — is a property that purely stochastic vibe coding cannot offer, as it lacks inspectable intermediate states. Even if a direct baseline were to show comparable final code quality, SPL is strictly better at *finding where intent was lost*.

We present SPL, the round-trip test protocol, and these 15 measurements as a starting point. Whether the observed patterns generalise, whether the score gradient reflects pipeline properties or model properties, and what level of round-trip score is sufficient for a given application are open questions we leave to future work. All code is released under Apache 2.0 at `https://github.com/digital-duck/SPL.py`.

References

[1] Palmblad et al. Agentic AI-assisted coding offers a unique opportunity to instill epistemic grounding during software development. arXiv:2604.21744, 2026.

[2] Casserini et al. Beyond the 'Diff': Addressing Agentic Entropy in Agentic Software Development. arXiv:2604.16323, 2026.

[3] Khattab et al. DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines. arXiv:2310.03714, 2023.

[4] Beurer-Kellner et al. Prompting Is Programming: A Query Language for Large Language Models. PLDI 2023.

[5] Daunis Ivan. A Declarative Language for Building and Orchestrating LLM Workflows. arXiv:2512.19769, 2025.

[6] Zhang et al. ReusStdFlow: A Standardized Reusability Framework for Dynamic Workflow Construction in Agentic AI. arXiv:2602.14922, 2026.

[7] Huang et al. MermaidFlow: Redefining Agentic Workflow Generation via Safety-Constrained Evolutionary Programming. arXiv:2505.22967, 2025.

[8] Xiao et al. FlowBench: Revisiting and Benchmarking Workflow-Guided Planning for LLM-based Agents. arXiv:2406.14884, 2024.

[9] Kuhn et al. Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation. arXiv:2302.09664, 2023.

[10] PocketFlow. Minimalist LLM Framework (100 lines). `https://github.com/the-pocket/PocketFlow`

[11] Kovac et al. DeepContext: Stateful Real-Time Detection of Multi-Turn Adversarial Intent Drift. arXiv:2602.16935, 2026.

[12] Zheng et al. SGLang: Efficient Execution of Structured Language Model Programs. arXiv:2312.07104, 2023.

Appendix

A SPL Grammar Summary

```
program     ::= (workflow | procedure | prompt | function)*
workflow    ::= 'WORKFLOW' name params 'DO' statement* 'END'
statement   ::= generate | select | call | call_parallel | evaluate
              | while | assign | commit | raise | retry
              | do_exception | store | import | log
generate    ::= 'GENERATE' expr 'INTO' var
select      ::= 'SELECT' expr ('FROM' source)? ('WHERE' condition)? 'INTO' var
call        ::= 'CALL' name '(' args ')' 'INTO' var
call_par    ::= 'CALL' 'PARALLEL' (call_branch ',')* call_branch 'END'
evaluate    ::= 'EVALUATE' var when_clause* ('ELSE' block)? 'END'
while       ::= 'WHILE' condition 'DO' statement* 'END'
commit      ::= 'RETURN' var ('WITH' metadata)?
exception   ::= 'DO' statement* 'EXCEPTION' when_exc* 'END'
import      ::= 'IMPORT' string_literal
type        ::= 'TEXT' | 'INT' | 'FLOAT' | 'BOOL'
              | 'IMAGE' | 'AUDIO' | 'VIDEO'
              | 'MAP' | 'LIST' | 'SET' | 'EXCEPTION'
```

B SPLc — Transpiler

SPLc is an AST-driven transpiler. It detects the workflow pattern from the AST structure and routes to the appropriate code generator. For established patterns, compilation is deterministic (no LLM); for novel patterns not yet in the pattern library, an `--llm` fallback generates code probabilistically.

C Round-trip Test Command Sequences

Each recipe directory contains a `README-{model}.md` with the full S1–S6 command sequence. The canonical template (shown for R1/M1):

```
export RECIPE=agent  ADAPTER=claude_cli  MODEL_ID=claude-sonnet-4-6  MODEL=sonnet
export OUT=~/projects/digital-duck/SPL.py/NeurIPS-26-lab/R1-agent/tests/$ADAPTER/$MODEL
export SRC=~/projects/digital-duck/SPL.py/NeurIPS-26-lab/R1-agent/src/pocketflow-agent

# S1 - describe original code
spl3 splc describe $SRC --include-docs --adapter $ADAPTER --model $MODEL_ID \
  -o $OUT/S1-$RECIPE-$ADAPTER-$MODEL-1-spec.md

# S2 - Mermaid diagram   *** HUMAN CHECKPOINT: review $OUT/S2-*.mmd ***
spl3 text2mmd $OUT/S1-$RECIPE-$ADAPTER-$MODEL-1-spec.md \
  --adapter $ADAPTER --model $MODEL_ID -o $OUT/S2-$RECIPE-$ADAPTER-$MODEL.mmd

# S3 - Mermaid → SPL
spl3 mmd2spl $OUT/S2-$RECIPE-$ADAPTER-$MODEL.mmd \
  --adapter $ADAPTER --model $MODEL_ID -o $OUT/S3-$RECIPE-$ADAPTER-$MODEL.spl

# S4 - SPL → Python/PocketFlow
spl3 splc compile $OUT/S3-$RECIPE-$ADAPTER-$MODEL.spl \
  --lang python/pocketflow --llm --adapter $ADAPTER --model $MODEL_ID \
  --out-dir $OUT/targets/python_pocketflow --overwrite
cp $OUT/targets/python_pocketflow/S3-$RECIPE-$ADAPTER-$MODEL*.py \
   $OUT/targets/python_pocketflow/S4-$RECIPE-$ADAPTER-$MODEL.py

# S5 - describe reconstructed code
spl3 splc describe $OUT/targets/python_pocketflow/S4-$RECIPE-$ADAPTER-$MODEL.py \
  --adapter $ADAPTER --model $MODEL_ID -o $OUT/S5-$RECIPE-$ADAPTER-$MODEL-2-spec.md

# S6 - round-trip score (judge fixed at claude-opus-4-6)
spl3 compare $OUT/S1-$RECIPE-$ADAPTER-$MODEL-1-spec.md \
```

```
$OUT/S5-$RECIPE-$ADAPTER-$MODEL-2-spec.md \
--adapter claude_cli --model claude-opus-4-6 \
-o $OUT/S6-$RECIPE-$ADAPTER-$MODEL-spec-diff.md
```

Models used: `claude_cli/claude-sonnet-4-6` (M1), `openrouter/gemini-3-flash-preview` (M2), `openrouter/qwen3.6-plus` (M3). S6 judge fixed at `claude_cli/claude-opus-4-6` across all 15 runs.

D Momagrid Distributed Runtime

The `momagrid` adapter routes SPL inference tasks to a Momagrid Hub, dispatching them across a LAN grid of GPU nodes. No `.spl` changes are required — only `--adapter momagrid` differs from a local run. Source: `https://github.com/digital-duck/momagrid`

E Future Work

1. **Multi-judge evaluation:** Use multiple independent judges or a non-Claude judge when testing Claude models to address the intra-provider self-preference risk.
2. **Larger dataset:** Expanding to more recipes, more model tiers, and multiple runs per configuration would yield confidence intervals and statistical tests of the observed score gradient.
3. **Ablation — bypass IR:** Comparing round-trip scores for the full pipeline against a direct vibe-coded equivalent (`spl3 vibe`) would quantify the value-add of the Mermaid and SPL intermediate representation steps.
4. **SPLc deterministic coverage:** Extending the deterministic transpiler to cover linear pipeline, multi-agent, and map-reduce patterns would eliminate the `--llm` fallback at S4.
5. **Cross-runtime round-trip tests:** Running S1–S6 for Go and TypeScript targets (in addition to Python/PocketFlow) would test whether the round-trip score is runtime-invariant.