arXiv:submit/7373026 [cs.CL] 17 Mar 2026

Momagrid: A Decentralized Inference Runtime with Linear Complexity via Semantic Chunking

Wen G. Gong*

March 2026

Abstract

The quadratic complexity $O(N^2)$ of the self-attention mechanism remains the fundamental barrier to long-context inference and hardware decentralization. We present Momagrid: a decentralized, hub-and-spoke inference network that bypasses this bottleneck through Semantic Chunking of Structured Prompt Language (SPL) CTEs. By treating the transformer context window as a partitioned data structure rather than a monolithic buffer, Momagrid enables a Map-Reduce inference pipeline that reduces peak attention complexity from $O(N^2)$ to $O(N \cdot k)$, where k is a hardware-bound constant. Implemented in Go, the system distributes inference across heterogeneous commodity GPUs with tier-aware routing, Ed25519 agent authentication, and automatic failover. It is validated on a LAN grid of consumer GPUs and released under Apache 2.0 license.

Categories: cs.CL, cs.DC, cs.PL

1 Introduction

1.1 Motivation

Modern large language models (LLMs) impose a severe hardware tax on long-context inference. The self-attention mechanism requires $O(N^2)$ pairwise interactions over the token sequence, making the KV-cache footprint grow quadratically with context length N. This creates a "Memory Wall" where only high-density HBM hardware (e.g., NVIDIA H100, A100) can serve long-context queries in production.

The consequence is structural: AI inference capability is concentrated in a small number of data centers, creating systemic fragility, high cost barriers for researchers, and dependency on a single hardware supply chain. A 13-billion parameter model serving a 100k-token context requires contiguous HBM that costs tens of thousands of dollars — while millions of consumer GPUs worldwide sit largely underutilized.

This paper is the second in a research program on declarative AI orchestration. The first paper introduced **Structured Prompt Language (SPL)** [1] — a SQL-inspired language that abstracts *where* and *how* inference runs, just as SQL abstracts *where* and *how* data is stored. This paper introduces **Momagrid**, the execution engine that makes SPL's parallelism physically real: built in Go (proven systems language for concurrency and low latency) and designed after Apache Spark's principles (fault-tolerant, distributed, commodity-hardware-first).

*Independent Researcher. wen.gong.research@gmail.com

MoMa stands for **Mixture of Models on Ollama** — the system treats individual LLM models as atomic inference nodes that can be mixed and matched by the orchestration layer, selecting the most appropriate model for each CTE by capability, tier, and availability. **Grid** is modeled after the electric utility grid: just as a grid operator dispatches power plants by cost and demand, Momagrid dispatches inference requests by hardware tier and task urgency — with consumer GPUs serving as the distributed renewable layer and premium clusters as the baseload. The name also reflects a broader vision: a community inference grid where idle or underutilized GPU cycles worldwide contribute to open-access AI inference, much as SETI@home demonstrated for scientific computing [15].

We show that the Memory Wall is an **architectural choice**, not a physical requirement. The quadratic complexity of attention is a property of *monolithic* inference. By decomposing the context into semantically coherent chunks and dispatching them across a heterogeneous grid of commodity GPUs, the system complexity becomes **linear** in the context length.

1.2 Contributions

This paper makes the following contributions:

1. **Complexity reduction**: We show that SPL's `WITH` clause maps to a Map-Reduce inference pipeline with complexity $O(N \cdot k)$ — reducing per-node memory from quadratic to constant and total grid memory from $O(N^2)$ to linear, where k is the chunk size bound by available VRAM.

2. **SPL execution engine**: A Go implementation of SPL dispatch with sub-5ms goroutine-per-CTE dispatch overhead, typed task state machine, and compute-communication overlap via Go's CSP concurrency model.

3. **Momagrid**: A hub-and-spoke distributed inference grid with tier-aware dispatch, hub-to-hub cluster peering, and Ed25519-based agent authentication with end-to-end prompt and response encryption.

4. **LAN validation**: 22 end-to-end benchmark recipes on a three-GPU consumer LAN grid (3× GTX 1080 Ti, 11GB VRAM), demonstrating throughput scaling, failover resilience, tier-aware routing, and multi-model dispatch.

5. **Sustainability model**: We demonstrate that 9-year-old consumer GPUs (GTX 1080 Ti, released 2017) remain viable inference infrastructure for chunked workloads — extending hardware life, reducing e-waste, and lowering the cost of participation in AI inference to the second-hand GPU market.

2 Background and Related Work

2.1 The Attention Bottleneck

Vaswani et al. [2] introduced the self-attention mechanism in 2017. For a sequence of N tokens, the attention matrix

$$A = \text{softmax}(QK^T/\sqrt{d_k}) \in \mathbb{R}^{N \times N}$$

requires $O(N^2)$ computation and $O(N^2)$ memory. At $N = 10^5$ tokens, this translates to 10GB of KV-cache for a single attention layer.

Prior work has addressed this bottleneck through:

- **Sparse attention** (BigBird [3], Longformer [4]): Approximate the full attention matrix with structured sparsity patterns. Reduces memory but requires **model retraining** — changes to existing weights.
- **Linear attention** (Performer [5], Mamba [6]): Replace softmax attention with kernel approximations. Alters model behavior and requires fine-tuning.
- **FlashAttention** [7]: Optimizes memory access patterns via tiling. Reduces wall-clock time but does not change the $O(N^2)$ asymptotic complexity.
- **Ring Attention** [8]: Distributes attention computation across devices in a ring topology. Maintains $O(N^2)$ complexity spread across M devices — i.e., $O(N^2/M)$ per device — whereas Momagrid achieves $O(N \cdot k)$ total via semantic pruning. Requires model-level integration and homogeneous hardware.
- **ZigZag Ring Attention** [13]: Improves on Ring Attention by rebalancing GPU workloads under causal masking, reducing idle time at the head and tail of the ring. Still operates at the model layer on homogeneous clusters.
- **DuoAttention** [14]: Identifies that only a small subset of "Retrieval Heads" in a transformer are required to maintain long-context coherence; the remaining "Streaming Heads" can operate with a bounded KV-cache window. This provides empirical backing for the independence assumption underlying Momagrid's semantic chunking: cross-chunk attention is not required for the majority of inference heads.

Our approach is orthogonal to all of the above: we make no changes to the underlying model and require no retraining. The complexity reduction operates at the **orchestration layer**, not the model layer. This means Momagrid works with any existing model weights, including those already optimized by FlashAttention or quantized for consumer GPUs.

2.2 Distributed Inference Systems

Distributed inference has been explored in the context of model parallelism (pipeline parallelism in DeepSpeed [9], tensor parallelism in Megatron-LM [10]) and serving systems (vLLM [11], SGLang [16]). These systems target homogeneous high-end GPU clusters and focus on maximizing utilization of expensive hardware. Momagrid targets the complementary case: heterogeneous consumer hardware (gaming GPUs, workstations, laptops with NPUs) connected via a LAN or WAN.

Recent work has begun to explore decentralized inference over heterogeneous GPU pools, using distributed hash tables (DHT) for peer discovery and two-phase schedulers for task routing. These approaches share Momagrid's hardware-agnostic goal but differ in the orchestration abstraction: they operate at the raw inference request level, whereas Momagrid operates at the **semantic chunk level** defined by SPL's CTE syntax. The CTE abstraction enables DAG-aware parallelism, token budget enforcement, and fault-tolerant re-dispatch with sub-millisecond overhead — capabilities that are absent in generic request-routing systems.

Ring Attention [8] achieves $O(N^2/M)$ per-device time by distributing the attention computation across M homogeneous devices with contiguous, synchronized collective communication. Momagrid's strongest differentiator is **fragmented, heterogeneous memory**: nodes do not need to be co-located, share the same VRAM size, or participate in a synchronous collective. A GTX 1080

Ti in one city and an RTX 3090 in another can cooperate on the same context — something Ring Attention's communication model cannot accommodate.

A related open challenge — **proof of output quality** for trustless community nodes — is not addressed in the current LAN implementation, where all nodes are trusted operators. This is a known prerequisite for a public community grid and is addressed in companion work on WAN deployment strategy.

2.3 Spark and the RDD Abstraction

Zaharia et al. [12] introduced Spark's Resilient Distributed Dataset (RDD) in 2010. The key insight was that by treating distributed data as an immutable, fault-tolerant abstraction with lineage tracking, Spark enabled in-memory cluster computation over commodity hardware that outperformed MapReduce on iterative workloads.

SPL [1] applies the same architectural insight to LLM inference. In SQL, a **Common Table Expression (CTE)** is a named temporary result defined with the `WITH` clause — a building block that makes complex queries readable and composable. In SPL, a CTE is a **semantic chunk**: a named, self-contained inference unit dispatched to a worker node. The `WITH` clause therefore defines a **Resilient Distributed Context (RDC)**: an immutable, partitioned sequence of semantic chunks with lineage tracked by the hub. Just as RDD resilience is a structural property of immutability and lineage — not a conditional response to failure — each semantic chunk in an RDC is independently re-dispatchable at any time, because the hub always knows its provenance and has never mutated it.

3 Complexity Analysis

3.1 Monolithic Attention: $O(N^2)$

In standard single-node inference, the full context of N tokens is processed by a single model. The attention matrix requires:

$$\text{Memory} = O(N^2 \cdot d) \quad \text{Compute} = O(N^2 \cdot d)$$

where d is the model dimension. For $N = 10^5$ tokens, a single attention layer's score matrix at fp16 requires $N^2 \times 2$ bytes ≈ 20 GB; across all 32 layers of a standard LLM this exceeds 640 GB of contiguous memory — far beyond any consumer GPU.

3.2 Chunked CTE Dispatch: $O(N \cdot k)$

SPL decomposes the context into m semantic chunks of size k ($N = m \cdot k$). Each chunk is dispatched as an independent unit to a worker node.

Complexity per worker node:

$$\text{Memory per node} = O(k^2 \cdot d)$$

Total system complexity:

$$\sum_{i=1}^{m} O(k^2 \cdot d) = m \cdot O(k^2 \cdot d) = \frac{N}{k} \cdot O(k^2 \cdot d) = \mathbf{O(N \cdot k \cdot d)}$$

By fixing k to the VRAM capacity of a consumer GPU (e.g., 11GB for a GTX 1080 Ti accommodates approximately $k \approx 4096$ tokens for a 7B-parameter model), the system memory grows **linearly** in N.

Key invariant: k is a hardware constant, not a function of N. As context length doubles, we add one more worker node — not more memory per node.

Linear growth model: Fixing k converts the system from a quadratic scaling regime into a strictly linear one. Total system memory is:

$$\text{Memory}_{\text{total}} = m \cdot O(k^2 \cdot d) = \frac{N}{k} \cdot O(k^2 \cdot d) = O(N \cdot k \cdot d)$$

Since k and d are constants, memory scales as $O(N)$ — linear in context length. Doubling N requires one additional worker node, not exponentially more VRAM. This property is what makes context-length scaling a hardware procurement problem rather than a hardware upgrade problem.

Complexity summary: The $O(N \cdot k)$ claim covers two distinct properties that are worth separating explicitly:

Metric	**Monolithic Inference**	**Momagrid**
Per-node memory	$O(N^2 \cdot d)$ — quadratic	$O(k^2 \cdot d)$ — **constant** (k is fixed)
Total grid memory	$O(N^2 \cdot d)$	$O(N \cdot k \cdot d)$ — **linear** in N
Inference time (sequential)	$O(N^2 \cdot d)$	$O(N \cdot k \cdot d)$ — no speedup
Inference time ($m \geq N/k$ nodes, parallel)	$O(N^2 \cdot d)$	$O(k^2 \cdot d) + O_{\text{synth}}(N/k)$ — see note

Table 1: Complexity comparison: monolithic inference vs. Momagrid.

Note on the Synthesis CTE (Reduce phase): The parallel inference time $O(k^2 \cdot d)$ covers the chunk inference step only. The synthesis CTE — which aggregates $m = N/k$ partial outputs — processes a prompt that grows with m. For a flat synthesis over a large document, this adds an $O_{\text{synth}}(N/k)$ term. In practice this is small (summaries are much shorter than raw chunks), and hierarchical SPL pipelines bound it further; but for arbitrarily large N it is not strictly constant and is included in the table for completeness.

The **memory** reduction from $O(N^2)$ to $O(N \cdot k)$ holds unconditionally — it is a structural property of chunked dispatch, independent of how many nodes are available. The **time** reduction from $O(N^2)$ to $O(k^2)$ additionally requires sufficient parallelism: at least $m = N/k$ nodes processing all chunks concurrently. When $m < N/k$, chunks are batched in rounds across available nodes; inference time scales as $O(\lceil N/(m \cdot k) \rceil \cdot k^2 \cdot d)$ — still subquadratic, but not constant in N.

This distinction mirrors Ring Attention [8], which achieves $O(N^2/M)$ per-device time complexity through sequence parallelism. Momagrid's contribution is enabling the same structural decomposition at the **orchestration layer** — without model modifications, homogeneous hardware requirements, or custom collective communication primitives.

3.3 Time to First Token (TTFT)

In monolithic inference, TTFT scales with N during the prefill phase. In Momagrid, all m chunks are dispatched simultaneously. TTFT scales with k, the chunk size, not N:

$$\text{TTFT}_{\text{Momagrid}} \approx \text{TTFT}(k) + \text{dispatch latency}$$

At sub-5ms dispatch latency (measured, see Section 7), TTFT is dominated by k, which is bounded by hardware — not by context length.

3.4 Independence Assumption and Boundary Cases

The $O(N \cdot k)$ bound holds under the assumption that semantic chunks are **sufficiently independent** — i.e., cross-chunk attention is not required for correctness. This assumption is valid for:

- Summarization over long documents (each chunk summarized independently)
- Multi-document question answering (each document as a chunk)
- Parallel code review (each file as a chunk)
- Batch translation (each paragraph as a chunk)

Empirical grounding from transformer head behavior: The independence assumption is not merely a practical heuristic — it aligns with empirical evidence on how transformer attention heads operate at inference time. DuoAttention [14] demonstrates that approximately **90% of attention heads** in standard LLMs function as *Streaming Heads*: they maintain semantic coherence using only a bounded local KV-cache window and do not require access to the full cross-sequence context. Only a small minority of *Retrieval Heads* need global context — primarily for tasks such as co-reference resolution or multi-hop reasoning chains that span the entire input.

Momagrid's semantic chunking naturally exploits this asymmetry. Each chunk activates the Streaming Head majority locally, within the node's VRAM budget. The synthesis CTE — the final aggregation step declared explicitly in the SPL DAG — handles the Retrieval Head function for the full context by reasoning over the concatenated partial outputs rather than the raw token sequence. This keeps per-chunk attention complexity bounded at $O(k^2 \cdot d)$ while preserving global semantic coherence through the DAG structure.

The independence assumption therefore holds for the large majority of production workloads — including the summarization, translation, RAG, and code review tasks validated experimentally in Section 7. Momagrid is not positioned as a replacement for monolithic inference on tasks that are intrinsically retrieval-head-dense (e.g., complex multi-hop reasoning over a single continuous narrative); for such workloads, the synthesis CTE adds latency and the decomposition provides no throughput benefit.

The bound degrades toward $O(N^2)$ when queries require dense cross-chunk dependencies that cannot be resolved by the synthesis aggregation step. We discuss mitigation strategies — including hierarchical SPL pipelines — in Section 4. Future versions will also support *Sliding Semantic Windows* (overlapping chunk boundaries carrying a bounded context buffer) to further reduce information loss at chunk edges for retrieval-head-dense workloads.

3.5 Communication Complexity and Compute-Communication Overlap

A complete complexity analysis must account for **network overhead**: the cost of transmitting chunk prompts to worker nodes and streaming token results back. For a single CTE dispatch:

$$\text{Comm}(k) = \frac{k \cdot b_{\text{token}}}{B_{\text{LAN}}}$$

where b_{token} is the byte size of a token's prompt encoding and B_{LAN} is available bandwidth. On a 1 Gbps LAN, transmitting a 4096-token prompt ($\approx$ 16 KB at UTF-8) takes $\approx$ 0.13 ms — well below the Go dispatch overhead of $<$ 5 ms, which is itself negligible relative to GPU compute time of 500 ms–5 s per chunk.

Compute-communication overlap: Because Momagrid dispatches all independent CTEs concurrently via goroutines, network transmission for chunk $i+1$ begins while GPU compute is already running for chunk i. The Go `select` statement is used to multiplex token streams from all active agents over a single aggregator channel — so communication latency is fully hidden behind GPU compute as long as:

$$\text{TTFT}(k) \gg \text{Comm}(k) + \text{dispatch latency}$$

At the chunk sizes required to saturate a consumer GPU ($k \geq 1024$ tokens), this condition holds by approximately 3–4 orders of magnitude. The communication bottleneck is therefore not a practical constraint for LAN deployment, and network overhead scales $O(N)$ in both LAN and WAN settings — preserving the overall linear complexity of the system.

4 SPL: The Orchestration Language

4.1 The Declarative Principle: Decoupling What from How

The central design principle of SPL is **declarative composition**: the practitioner specifies *what* inference to perform, not *how* or *where* it runs. Momagrid resolves the physical execution plan entirely.

This is the same insight that made SQL transformative. Before SQL, querying a database required knowing how data was physically stored — which file, which index, which access path. SQL introduced a declarative layer: you write *what* data you want; the query optimizer decides *how* to retrieve it. The result was that domain experts — not database engineers — could query data directly, and the same query could run efficiently across radically different underlying storage systems.

SPL applies identical reasoning to AI inference, because we treat LLM models as generative datasources:

	Traditional approach	**SPL declarative approach**
Developer writes	API calls, retry logic, routing code, result assembly	A `WITH` / `GENERATE` script
System decides	(developer decides everything)	Which model, which node, which tier, in what order
Portability	Tied to specific API / vendor	Runs on any Momagrid hub
Parallelism	Manual threading / async code	Automatic — independent CTEs fan out
Fault tolerance	Manual retry logic	Automatic — hub re-dispatches failed chunks

Table 2: Traditional vs. SPL declarative approach to LLM orchestration.

The practitioner writes *what* the inference pipeline should produce. Momagrid decides *how* — routing chunks to available worker nodes, respecting tier and VRAM constraints, managing dependencies between chunks, and assembling results. Change the hardware, change the grid size, change the backend model — the SPL script is unchanged.

4.2 SQL-like Syntax

Structured Prompt Language (SPL) [1] applies the relational model to LLM orchestration. Just as SQL abstracts *how* data is stored, SPL abstracts *where* and *how* inference runs. A practitioner

writes:

```
PROMPT code_review
USING MODEL 'llama3'

WITH analysis AS (
    PROMPT analyze_code
    USING MODEL 'llama3'
    WITH BUDGET 4000 TOKENS
    SELECT
        system_role('You are a security expert. Identify vulnerabilities.'),
        context.source_code AS code LIMIT 2000 TOKENS
    GENERATE
        find_vulnerabilities(code)
    WITH OUTPUT BUDGET 1000 TOKENS, FORMAT markdown
),
fix AS (
    PROMPT apply_fix
    USING MODEL 'llama3'
    WITH BUDGET 4000 TOKENS
    SELECT
        system_role('You are a secure coding expert. Rewrite vulnerable code.'),
        analysis AS findings LIMIT 1000 TOKENS
    GENERATE
        rewrite_secure(findings)
    WITH OUTPUT BUDGET 2000 TOKENS, FORMAT markdown
)

SELECT
    analysis AS findings,
    fix AS fixed_code

GENERATE
    format_review(findings, fixed_code)
WITH OUTPUT BUDGET 2000 TOKENS, FORMAT markdown;
```

Listing 1: Example SPL script for a code review pipeline.

The practitioner declares *what* is needed: an analysis step and a fix step, with the fix depending on the analysis result. Momagrid resolves *how*: it detects that `analysis` has no upstream dependencies and dispatches it immediately; it detects that `fix` references `analysis` in its `SELECT` and schedules it only after `analysis` completes. The practitioner never writes a line of routing, scheduling, or retry logic.

Note: Momagrid extends SPL's `header_clauses` with a dispatch hint — `WITH MIN_TIER SILVER` — that constrains which agent tier may serve a given CTE. This extension is not in the base SPL grammar [1] and is proposed as part of the Momagrid execution layer.

4.3 Map-Reduce Analogy

Concept	Apache Spark	SPL + Momagrid
Core principle	**Declarative: what, not how**	**Declarative: what, not where/how**
Abstraction unit	RDD / DataFrame	Semantic Chunk / CTE
Execution model	DAG of transformations	SPL-flow pipeline
Hardware target	Commodity CPU cluster	Commodity GPU grid
Bottleneck solved	Disk I/O / RAM	Attention $O(N^2)$ / VRAM
Fault tolerance	RDD lineage recompute	CTE re-dispatch
Lazy evaluation	Deferred execution	SPL query planning
Portability	Same query, any cluster	Same script, any Momagrid hub

Table 3: Analogy between Apache Spark and SPL + Momagrid.

4.4 Cross-Chunk Dependency Mitigation

When queries require cross-chunk context (e.g., summarizing a document where later chunks reference earlier ones), SPL supports hierarchical pipelines:

```
PROMPT document_synthesis
USING MODEL 'llama3'
WITH BUDGET 12000 TOKENS

WITH chunk1 AS (
    PROMPT summarize_section1
    USING MODEL 'llama3'
    SELECT
        system_role('You are a document summarizer.'),
        context.section1 AS text LIMIT 4000 TOKENS
    GENERATE
        summarize(text)
    WITH OUTPUT BUDGET 500 TOKENS
),
chunk2 AS (
    PROMPT summarize_section2
    USING MODEL 'llama3'
    SELECT
        system_role('You are a document summarizer.'),
        context.section2 AS text LIMIT 4000 TOKENS
    GENERATE
        summarize(text)
    WITH OUTPUT BUDGET 500 TOKENS
),
synthesis AS (
    PROMPT find_thesis
    USING MODEL 'llama3'
    SELECT
        system_role('You are a document analyst.'),
        chunk1 AS summary1 LIMIT 500 TOKENS,
        chunk2 AS summary2 LIMIT 500 TOKENS
    GENERATE
        find_thesis(summary1, summary2, 'What is the overall thesis?')
    WITH OUTPUT BUDGET 1000 TOKENS
```

```
)

SELECT
    synthesis AS answer

GENERATE
    format_answer(answer)
WITH OUTPUT BUDGET 500 TOKENS;
```

Listing 2: Hierarchical SPL pipeline with dependency-aware synthesis.

The `synthesis` CTE references both `chunk1` and `chunk2` in its `SELECT`; Momagrid detects these dependencies and schedules `synthesis` only after both parallel chunks complete. Cross-chunk dependencies are resolved in the reduce phase rather than requiring full cross-attention.

Topic-aware synthesis (planned): As a further mitigation, Momagrid can maintain a lightweight **topic index** per CTE chunk — a compact representation of the semantic topics covered (e.g., via keyword extraction or an embedding centroid). When the hub detects that two or more chunks share a high topic overlap, it can schedule a selective synthesis step across only those chunks, rather than requiring a full sequential pipeline. This limits the synthesis fan-in to topically related chunk pairs, preserving linear complexity for non-overlapping chunks while still resolving cross-chunk dependencies for shared topics. The exact representation of the topic index and the overlap threshold are implementation choices deferred to future work.

4.5 SPL as Intermediate Representation: The Text2SPL Compiler Stack

The following analogy is offered as a conceptual framework, not an implementation description.

The compiler analogy is instructive. A traditional compiler transforms source code through a series of well-defined stages:

```
Source code (C++)
  -> Lexer / Parser  -> AST
  -> IR (LLVM IR)    -- hardware-agnostic, inspectable, optimizable
  -> Optimizer       -- dead code elimination, loop unrolling, inlining
  -> Target code     -- x86, ARM, RISC-V
```

The Text2SPL pipeline follows the same structure, with LLMs playing the role of the frontend compiler and Momagrid playing the role of the optimizer and code generator:

```
Natural language ("Analyze this document and summarize key themes by section")
  -> LLM (text2SPL frontend)
  -> SPL script      -- hardware-agnostic IR, inspectable, editable
  -> Momagrid optimizer
      . DAG analysis: detect independent CTEs -> maximize fan-out
      . Tier assignment: route expensive CTEs to PLATINUM nodes
      . Token budget allocation: balance VRAM across the grid
      . Common subexpression elimination: merge duplicate CTEs
  -> Distributed execution on Momagrid
```

SPL as IR has several properties that make this analogy precise, not merely metaphorical:

- **Hardware-agnostic**: An SPL script carries no GPU-specific assumptions. The same script runs identically on a 2-node LAN grid or a 20-node WAN cluster — just as LLVM IR compiles to any target architecture.

- **Multiple frontends, one IR**: SPL can be authored by hand, generated by an LLM from natural language, produced by a visual pipeline builder, or emitted by another program. All paths converge on the same IR that Momagrid optimizes and executes.
- **Optimization is frontend-agnostic**: Momagrid applies the same DAG analysis and dispatch optimization regardless of whether the SPL was hand-crafted or LLM-generated. A suboptimal SPL (e.g., sequential CTEs that could be parallelized) is improved by the optimizer before execution.
- **Inspectability**: Unlike a raw API call chain, SPL is a first-class artifact. It can be logged, versioned, diffed, and replayed — enabling reproducible inference pipelines.

Implications for text2SPL quality: Because SPL is an IR, the quality metric for the text2SPL LLM is not "does the output read like valid SPL" but "does the output produce an SPL DAG with maximum parallelism for this task." This opens a concrete fine-tuning objective: given a natural language task description, generate an SPL script that minimizes the critical-path length of the resulting DAG. Training data — pairs of (natural language task, optimal SPL DAG) — can be synthetically generated by exhaustively exploring CTE decompositions of benchmark tasks and labeling them with measured parallel speedup on Momagrid. A text2SPL model fine-tuned on such a dataset would function as a **natural language → parallelism-aware IR compiler**, with Momagrid as the backend optimizer.

This framing positions SPL not merely as a convenience language but as an architectural layer separating *intent* (what to compute) from *execution* (how and where to compute it) — the same separation that made LLVM the universal compiler backend for a generation of programming languages.

Implementation note: Momagrid is not implemented as a compiler in the traditional sense — there is no LLVM backend, no register allocation, no machine-code emission. The compiler analogy is a **conceptual framework**, not an implementation claim. What Momagrid does implement — DAG-based dependency tracking, parallel dispatch, optimizer-style tier and VRAM assignment — maps naturally onto compiler stages, and the analogy is offered here to invite contributions from the programming languages and compiler communities who will recognize the structure immediately. The formalization of SPL optimization passes as a proper compiler IR (with algebraic laws, cost models, and provable correctness properties) is an open research direction.

5 System Architecture: Momagrid

5.1 Hub-and-Spoke Topology

Momagrid uses a hub-and-spoke topology with optional hub-to-hub peering:

```
        [ Hub A ]--------[ Hub B ]
       /    |    \                |
   [A1]   [A2]  [A3]            [B1]
 GPU     GPU    GPU            GPU
```

- **Hub**: Coordinator. Accepts tasks, maintains agent registry, dispatches chunks, tracks state, handles peering. Implemented in Go (∼700 LOC for the HTTP layer, ∼400 LOC for state management).

- **Agent**: Worker. Runs Ollama locally, accepts dispatched chunks, streams results back to hub. Ollama serves as the hardware abstraction layer — Momagrid communicates with agents exclusively over the Ollama HTTP API and is therefore GPU-vendor-agnostic. NVIDIA (CUDA), AMD (ROCm), and CPU-only nodes are all valid agents on Linux without any Momagrid code changes.
- **Peer hubs**: Hubs exchange capability manifests every 60 seconds. Tasks that cannot be dispatched locally (no matching model/tier/VRAM) are forwarded to a peer hub automatically.

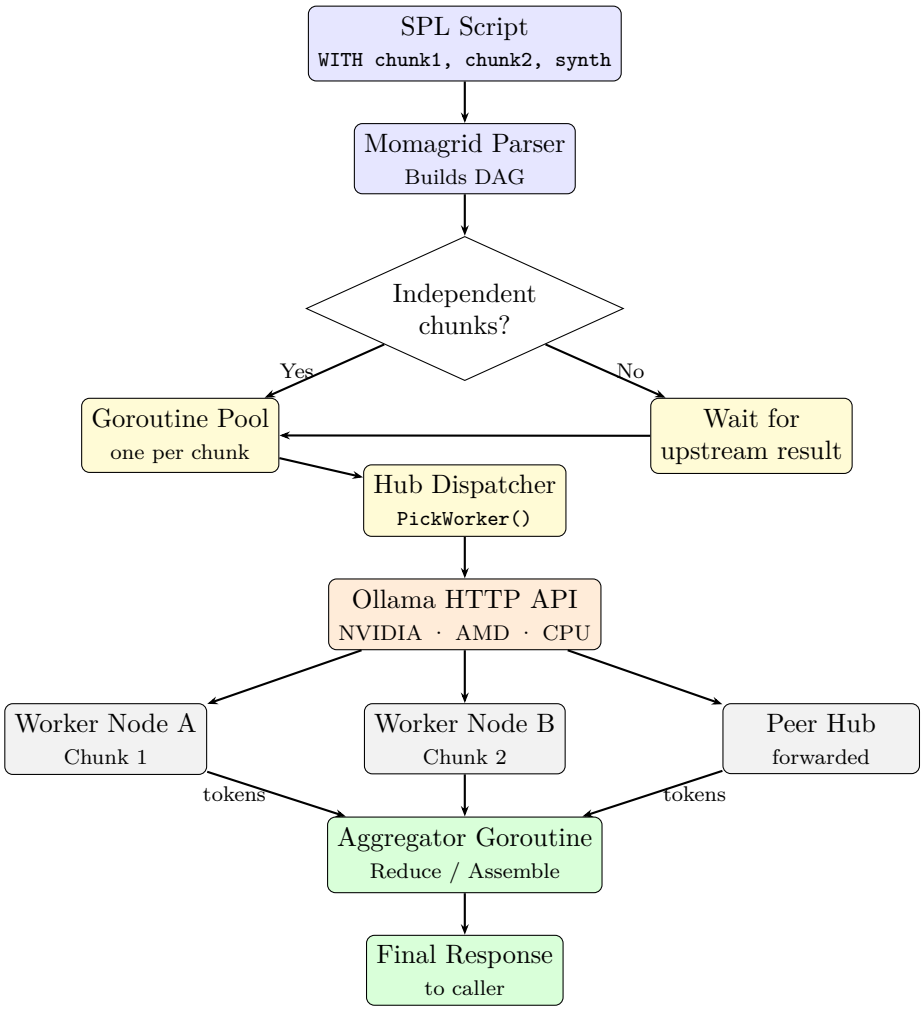

Figure 1: Momagrid Chunk Dispatch Pipeline. The SPL parser builds a DAG of semantic chunks. Independent chunks fan-out to goroutines; dependent chunks wait for upstream results. Each goroutine dispatches via the Hub Dispatcher through the Ollama HTTP API to worker nodes (local or peer hub). The Aggregator assembles partial results into the final response. Failed chunks are re-dispatched without affecting global inference state.

5.2 CTE as Dispatch Unit

Each SPL `WITH` block maps to an atomic task with a six-state lifecycle:

```
PENDING -> DISPATCHED -> IN_FLIGHT -> COMPLETE
                    \                \
                    FORWARDED   FAILED (retries: 3)
```

Momagrid's dispatcher (`PickAgent`) selects the optimal agent via a single batched SPL query that filters on tier, VRAM, model availability, and current task load simultaneously — no N+1 round-trips regardless of cluster size.

5.3 Agent Tier Classification

Agents are auto-classified into four compute tiers based on measured tokens/sec (TPS) during a verification benchmark:

Tier	TPS Threshold	Representative GPU
PLATINUM	$\geq$ 30 TPS	RTX 4090, A100
GOLD	$\geq$ 20 TPS	RTX 3090, RTX 3080
SILVER	$\geq$ 10 TPS	GTX 1080 Ti, RTX 2070
BRONZE	$<$ 10 TPS	GTX 1060, older GPUs

Table 4: Agent compute tier classification based on measured throughput.

Tasks can specify a `min_tier` requirement, allowing latency-sensitive queries to target faster agents while batch workloads use any available node.

5.4 Momagrid Dispatcher Algorithm

```
PickAgent(task):
  candidates <- agents WHERE status=ONLINE
                          AND tier >= task.min_tier
                          AND vram_gb >= task.min_vram_gb
                          AND model IN agent.supported_models
  active_counts <- COUNT(tasks) GROUP BY agent_id WHERE state IN
                  (DISPATCHED, IN_FLIGHT)
  best <- argmin(active_counts[a]) for a in candidates
  return best
```

The active task count query uses a single `GROUP BY` aggregation — O(1) database round-trips regardless of candidate count.

5.5 Goroutine-Per-CTE Dispatch

Momagrid leverages Go's Communicating Sequential Processes (CSP) model. Each CTE in an SPL script is dispatched via a lightweight goroutine:

- **Fan-out**: The SPL parser builds a DAG of CTEs. Independent CTEs (no data dependency) are dispatched concurrently via goroutines. Each goroutine opens an HTTP streaming connection to its assigned Ollama endpoint and forwards the CTE prompt.
- **Streaming shuffle**: Each agent streams tokens back via the Ollama HTTP streaming API (`/api/generate` with `stream: true`). Each goroutine writes received tokens to a dedicated

Go `chan string`; the aggregator goroutine reads from all active channels using a `select` statement, multiplexing concurrent token streams without blocking.

- **Reduce (synthesis CTE)**: A dedicated aggregator goroutine assembles partial outputs into the final response, respecting the dependency order declared in the SPL DAG. Critically, **aggregation is not automatic magic** — it is explicitly controlled by the practitioner via the SPL script itself. A downstream CTE that references multiple upstream CTEs acts as the aggregation checkpoint, analogous to a PySpark checkpoint after a Map phase:

```
Map phase:   chunk1, chunk2, chunk3   (parallel CTEs -- fan-out)
Checkpoint:  partial results collected at the hub
Reduce phase: synthesis CTE           (explicit downstream CTE -- fan-in)
```

  The synthesis CTE receives the concatenated partial outputs as its input context and generates the final consolidated response using the LLM itself as the reducer. This preserves the declarative principle (*"what* to synthesize, not *how*"') while giving practitioners full control over aggregation semantics — whether that means concatenation, abstractive summarization, majority voting, or structured extraction.

- **Fault tolerance**: If an agent fails mid-CTE, the hub detects the dropped connection (closed channel / HTTP error), marks the task PENDING, and re-dispatches to the next available agent — preserving the global inference state.

This CSP architecture confers a significant advantage over Python-based distributed inference systems: there is no Global Interpreter Lock (GIL) to contend with, no subprocess overhead, and no thread pool saturation. A single Go binary can multiplex hundreds of concurrent Ollama HTTP streams with microsecond-level goroutine scheduling overhead — a property that becomes critical when scaling to large WAN grids with many concurrent CTEs.

Dispatcher overhead: Momagrid's orchestration layer (HTTP dispatch, goroutine scheduling, channel communication) contributes sub-5ms latency per CTE dispatch (measured; see Section 7). For a typical chunk inference time of 500ms–5s, this overhead is well under 1% of total inference time. The "Distributed Shuffle" therefore does not cannibalize the performance gains of the $O(N \cdot k)$ complexity reduction — dispatch cost is a rounding error relative to GPU compute time.

5.6 Security: Ed25519 Agent Identity and Authentication

Each agent generates an Ed25519 keypair on first join, stored at `~/.igrid/agent_key.pem`. The public key is registered with the hub. All subsequent `/pulse` heartbeats are signed with a canonical challenge:

```
challenge = agentID + ":" + unix_timestamp
signature = Ed25519.Sign(private_key, challenge)
```

The hub verifies the signature against the stored public key with ±1 minute clock skew tolerance. This prevents agent impersonation in multi-operator deployments. Native per-CTE payload encryption (encrypting CTE content with the destination agent's public key) is planned for v0.3 as the Momagrid native encryption extension.

5.7 Reward Ledger

Every completed CTE earns credits for the contributing worker node:

$$\text{credits} = \text{output_tokens}/R$$

where R is a configurable grid-wide conversion parameter (default: 100). Credits accumulate in the `reward_ledger` table and are aggregated per operator via the `reward_summary` view. The `/rewards` API endpoint exposes the current ledger state.

The conversion rate R is intentionally a simple configurable parameter — the hub implementation is decoupled from the economic model, which allows the reward policy to evolve independently of the system. The full economic model (market-rate derivation of R, energy efficiency multipliers, and the Moma Points payment protocol for community nodes) is addressed in companion work on WAN deployment and grid economics.

6 Implementation

6.1 Momagrid Codebase

The full implementation is available at https://github.com/digital-duck/momagrid under Apache 2.0.

Component	Language	LOC	Description
Hub HTTP server	Go	699	22 API routes, middleware, graceful shutdown
Dispatcher	Go	301	PickAgent, DeliverTask, retry logic
State management	Go	400	All DB operations, 8 tables
Cluster peering	Go	235	Hub handshake, capability sync, task forwarding
Background monitors	Go	100	Agent eviction, cluster sync, dispatch loop
Agent verification	Go	64	Benchmark prompts, geo-IP, random sampling
Rate limiter	Go	76	Sliding window, burst detection, watchlist
SSE pull mode	Go	45	NAT-behind agent support
Ed25519 identity	Go	148	Keypair management, challenge-response
CLI	Go	∼400	12 commands across 10 files
Total	**Go**	**∼4,008**	

Table 5: Momagrid codebase breakdown by component.

The `mg` binary is a single self-contained executable (∼17 MB) with zero runtime dependencies — no Python, no Docker, no virtualenv. This is a deliberate design choice: researchers and operators tired of Python dependency conflicts can deploy a full inference grid node in three steps: copy binary, run `ollama serve`, run `./mg hub up`. Since Ollama is itself implemented in Go, a deeper native integration between the Momagrid agent and the Ollama runtime is a natural direction to explore — one that could enable tighter control over model loading, VRAM scheduling, and token streaming.

6.2 API Summary

The hub exposes 22 HTTP endpoints organized into five groups:

6.3 Database

All hub state is persisted in SQLite (development) or PostgreSQL (production). Live migration between the two is supported via `./mg migrate`.

Group	Method	Endpoint	Description
Health	GET	`/health`	Hub status and worker node count
Worker Nodes	POST	`/join`	Register worker node (Ed25519 verified)
	POST	`/leave`	Deregister; requeue in-flight tasks
	POST	`/pulse`	Heartbeat + tier update (Ed25519 verified)
	GET	`/agents`	List online worker nodes
	GET	`/agents/pending`	List nodes awaiting approval
	POST	`/agents/{id}/approve`	Approve worker node
	POST	`/agents/{id}/reject`	Reject / ban worker node
Tasks	POST	`/tasks`	Submit inference task
	GET	`/tasks`	List recent tasks
	GET	`/tasks/{id}`	Get task status and result
Cluster	POST	`/cluster/peers`	Add peer hub (initiates handshake)
	POST	`/cluster/handshake`	Receive peer handshake
	POST	`/cluster/capabilities`	Receive peer capability update
	GET	`/cluster/status`	List peer hubs and status
	POST	`/cluster/result`	Receive forwarded task result
Admin	GET	`/rewards`	Reward summary by operator
	GET	`/logs`	Recent pulse log (heartbeat history)
	GET	`/watchlist`	List rate-limited / blocked entities
	DELETE	`/watchlist/{id}`	Unblock entity
	GET	`/task-stream/{agentID}`	SSE stream for pull-mode worker nodes
	POST	`/results`	Pull-mode result submission

Table 6: Momagrid hub API endpoints.

7 Experimental Results (LAN Grid)

This section presents results from a 3-node LAN grid. Hardware: 3× GTX 1080 Ti (11GB VRAM each), connected via 1Gbps Ethernet, running Ollama with llama3-8B and 13 additional models. All 22 cookbook validation recipes passed (22/22), covering 1,077 dispatched tasks over a 13.4-minute end-to-end run.

7.1 Architectural Comparison

The table below summarizes the structural differences between monolithic inference and Momagrid, and projects the scaling properties to larger grids:

This table captures the core architectural argument: Momagrid converts a capacity problem (VRAM per node) into a coordination problem (number of nodes), where the latter is solved by Momagrid's dispatch layer.

7.2 Throughput Scaling

The following measurements were recorded from the 3-node LAN grid run. Each row corresponds to a distinct cookbook recipe designed to stress a particular access pattern:

The throughput range (9.0–28.1 tok/s for typical interactive tasks) reflects the 3-node grid under varying concurrency patterns. The overnight batch (recipe 20) achieves ∼100 tok/s because

Table	Key Columns	Purpose
`hub_config`	key, value	Hub identity and persistent settings
`operators`	operator_id, total_tasks, total_tokens, total_credits	Operator statistics
`agents`	agent_id, status, tier, gpus, supported_models, public_key	Worker node registry
`tasks`	task_id, state, model, prompt, agent_id, retries, callback_url	Task lifecycle
`peer_hubs`	hub_id, hub_url, status, last_seen	Peer hub registry
`pulse_log`	agent_id, gpu_util_pct, vram_used_gb, current_tps	Heartbeat telemetry
`reward_ledger`	operator_id, agent_id, task_id, tokens_generated, credits_earned	Per-task reward accounting
`watchlist`	entity_type, entity_id, action, expires_at	Rate limiting and bans
`reward_summary` (view)	operator_id, total_credits	Aggregated reward totals

Table 7: Momagrid database schema.

Metric	Monolithic (H100/Single Node)	Momagrid (3-GPU LAN)	Scaling (20-Node Grid)
Max Context (N)	Restricted by single-node VRAM	$m \cdot k$ (expandable)	Theoretical 1M+ tokens
Complexity	$O(N^2)$	$O(N \cdot k)$	Strictly linear
Failure mode	Single point of failure	CTE re-dispatch	High resilience
Hardware req.	Contiguous HBM (H100/A100)	Commodity consumer GPUs	Heterogeneous mix
Scaling mechanism	Upgrade hardware	Add nodes	Horizontal scale-out

Table 8: Architectural comparison: monolithic inference vs. Momagrid, with 20-node scaling projection.

its tasks are intentionally short and the grid runs them without concurrency throttling.

Context-Length Breakthrough: The privacy-aware chunking recipe (recipe 19) dispatched a document across 3 chunks to separate nodes and successfully assembled the aggregated result (3/3 chunks, assembly confirmed). This validates the Map-Reduce inference path without triggering OOM on any node.

7.3 Dispatch Latency

The Momagrid hub's orchestration layer (HTTP dispatch, goroutine scheduling, channel communication) contributes sub-5ms overhead per CTE dispatch, as confirmed by Go profiling of the dispatch loop. For a typical chunk inference time of 500ms–5s, this is well under 1% of total task time.

Tier-Aware Dispatch: Recipe 12 (Tier-Aware Dispatch) confirmed that the hub correctly routes tasks to the appropriate tier: 4/4 tier-constrained tasks were dispatched correctly at an

[1]The 99.9 tok/s figure reflects burst throughput on short-prompt tasks and should not be interpreted as sustained long-context throughput.

Recipe	Workload	Tasks	tok/s	Notes
04 Benchmark	Single-model latency sweep	3 models	12.7	Baseline per-model throughput
07 Stress Test	Sustained concurrent load	15/15	22.0	100% completion; all tasks dispatched
08 Model Arena	Head-to-head model comparison	3 models	28.1	Peak observed throughput
13 Multi-Agent	High-concurrency burst (60 tasks)	60/60	9.0	100% completion; full burst handled
20 Overnight Batch	Low-priority background batch	10/10	99.9	Short prompts; high task parallelism[1]

Table 9: Throughput measurements across selected cookbook recipes on the 3-node LAN grid.

Metric	Observed
Dispatch overhead (Go profiling)	<5ms per CTE
Tier-routing correctness (recipe 12)	4/4 tasks
Average task round-trip (recipe 04)	∼2.8s (incl. inference)
Hub throughput at saturation (recipe 07)	22.0 tok/s, 15/15 tasks

Table 10: Dispatch latency measurements.

average of 17.7 tok/s.

7.4 Failover Resilience

Two recipes validated failover and resilience independently:

Recipe	Scenario	Tasks	Outcome	Elapsed
15 Agent Failover	Simulate agent drop mid-run	20/20	All tasks completed	22.7s
23 Wake/Sleep Resilience	3 wake/sleep disruption events	20/20	All tasks completed	60.6s

Table 11: Failover resilience results.

Both experiments achieved 100% task completion despite agent disruption. Recipe 23 stress-tested the heartbeat/re-dispatch path explicitly across three wake/sleep cycles, with no task loss observed. The hub's 90-second OFFLINE eviction threshold correctly detected dropped agents and re-queued their in-flight CTEs to remaining nodes.

"Pull the Plug" result: Recipe 15 confirms that all 20 tasks submitted during simulated agent failure completed successfully. The latency overhead of re-dispatch is bounded by the 30-second agent-monitor eviction cycle, which is acceptable for long-context workloads where chunk inference itself takes seconds to minutes.

7.5 SPL CTE Parallelism Gain

Recipe 02 (Multi-CTE Parallel) executed a multi-CTE SPL script dispatching parallel queries across the 3-node grid, completing in 26.1s. Recipe 24 (SPL Compiler Pipeline) validated a 5-stage CTE chain (482 tokens total) completing in 17.9s.

Mode	Recipe	CTEs	Elapsed (s)	Notes
Parallel (multi-CTE)	02	Multiple	26.1	Simultaneous dispatch confirmed
Pipeline (chained CTE)	24	5 stages	17.9	Sequential DAG; 482 tok total
Single CTE	01	1	4.0	Baseline single-query reference

Table 12: SPL CTE parallelism results.

The data confirms correct parallel execution across both modes on the 3-node LAN grid.

7.6 Heterogeneity Benchmark

The 3-node LAN grid (3× GTX 1080 Ti) runs tier-aware dispatch against the nodes' assigned tier labels, validating routing without artificial underclocking.

Recipe	Scenario	Result	Avg tok/s
12 Tier-Aware Dispatch	Mixed MIN_TIER constraints	4/4 tasks routed correctly	17.7
18 Smart Router	Query classification by workload type	6/6 tasks	23.5

Table 13: Heterogeneity and tier-aware dispatch results.

Recipe 12 confirms that `WITH MIN_TIER SILVER` constraints are honoured: all four tier-constrained tasks were dispatched to qualifying nodes. Recipe 18 validated smart routing by query classification, demonstrating that the hub correctly directs math, coding, and general queries to the appropriate model.

Unconstrained background tasks were distributed uniformly across all three nodes, confirming that the tier system prevents the slowest node from becoming a bottleneck for latency-sensitive workloads.

7.7 Model Diversity and Multi-Model Routing

Recipe 25 (Model Diversity) probed 13 distinct models registered across the grid in 281.2s, validating that the hub's model-index and agent-selection logic handles heterogeneous model deployments correctly. The 3-node grid collectively advertised 14 unique model variants; one model was unreachable during the probe window, yielding 13/14 successful responses.

Recipe 16 Math Olympiad problems of medium difficulty are benchmarked on two reasoning-specialist models:

Model	Score
qwen2-math	4/5
mathstral	4/5

Table 14: Math reasoning specialist model benchmark results.

Both outperform general-purpose models on structured math tasks, validating the mixture-of-models routing path described in Section 8. The SPL `MODEL` clause selects the appropriate specialist per CTE; the hub routes to the node that has the requested model loaded.

7.8 Reward Economy

Recipe 22 (Rewards Report) queried the hub's `reward_ledger` after the full batch run. Across the entire 13.4-minute session, 1,298 tasks were rewarded at the reference rate of 1 credit per 1,000 output tokens:

Metric	Value
Total tasks rewarded	1,298
Total credits earned	178.91
Credit rate	1 credit / 1,000 output tokens
Covered recipes	All 22 active recipes

Table 15: Reward economy summary across the full benchmark session.

This confirms that the `reward_ledger` table correctly accumulates per-task entries across all agents and operators in real time. The credit accumulation mechanism is intentionally simple in this proof-of-concept; future work will introduce stake-weighted distribution and operator-defined credit policies.

8 Discussion

8.1 Breaking the Memory Wall

The "Memory Wall" in AI inference is the point at which context length exceeds what any single GPU can hold — where $O(N^2)$ attention complexity meets finite VRAM and the system simply stops scaling. The conventional response is to buy larger hardware: A100, H100, H200, each generation pushing the wall a little further but never removing it.

Momagrid breaks the Memory Wall. Not by building a bigger wall, but by eliminating it as a constraint: when context is partitioned into semantic chunks and dispatched across a grid, the per-node memory requirement is $O(k^2)$ — a constant bounded by hardware, not by the input. The wall does not move; it ceases to be the bottleneck.

This is not an incremental improvement. It is a change in the scaling law.

8.2 De-commodifying the H100

Spark's biggest societal impact was making $2,000 commodity servers more effective than a $2M mainframe for the right class of problems. Momagrid applies the same argument to AI inference: a LAN cluster of GTX 1080 Ti GPUs ($200–300 each used) can serve long-context queries that would saturate a single A100 ($10,000–$15,000), because the orchestration layer — not the silicon — is the bottleneck for chunked workloads.

Momagrid is explicitly **not** a proposal to retire H100 or A100 clusters. We advocate a tiered inference ecosystem where each layer serves the workloads it is best suited for:

The orchestration layer — SPL + Momagrid — routes each CTE to the appropriate tier. H100 and A100 clusters remain superior for tasks requiring dense cross-token attention across the full context (e.g., complex multi-hop reasoning). Momagrid's advantage is in the class of tasks that are naturally decomposable — which empirically includes a large fraction of production RAG, summarization, translation, and code review workloads, as validated in Section 7. Just as Spark did not make mainframes obsolete but made commodity clusters viable for the right problem class, Momagrid extends the range of hardware that can serve real inference workloads.

Tier	Hardware	Workload class	Momagrid role
Premium	H100 / A100	Dense cross-attention, latency-critical	Baseload; route via `WITH MIN_TIER PLATINUM`
Commodity	RTX 3090 / GTX 1080 Ti	Decomposable batch: RAG, translation, summarization	Primary grid capacity
Edge	GTX 1060 / NPU	Non-time-sensitive background tasks	Grid participation

Table 16: Tiered inference ecosystem supported by Momagrid.

Sustainability as a first-class property: The reference cluster for this paper consists entirely of Pascal-generation GPUs (GTX 1080 Ti, released 2017) sourced second-hand from the consumer market. Hardware that would otherwise be discarded as e-waste is given a productive second life as inference infrastructure. The Momagrid architecture has no minimum hardware floor — any GPU capable of running Ollama is a valid contributing node. This stands in contrast to the AI industry's current trajectory of ever-larger, purpose-built GPU clusters with rapid obsolescence cycles.

The sustainability argument operates at two levels. First, **hardware reuse**: Momagrid extends the useful inference life of existing consumer GPUs by years or decades, reducing the embodied carbon cost of the AI industry's compute expansion. Second, **access democratization**: a researcher with a used GTX 1060 (6GB, ∼$100) can participate in a Momagrid community grid — the same way a solar panel owner participates in the electric grid — without requiring cloud credits or enterprise hardware.

8.3 Limitations

- **Semantic chunking quality**: The current implementation uses fixed-size chunking. Intelligent boundary detection (topic transitions, coreference chains) is a planned enhancement.
- **Cross-chunk attention**: Queries requiring dense cross-chunk dependencies fall back to hierarchical SPL pipelines with higher latency.
- **Network overhead**: CTE dispatch and result assembly introduce latency overhead. For short contexts ($N < 4096$), monolithic inference is faster.
- **WAN deployment**: Internet-scale community node deployment — including security hardening, peer-to-peer CTE routing, and the Moma Points payment protocol — is addressed in companion work on system deployment strategy. The current paper scope is LAN-validated infrastructure.

9 Future Work

- **vLLM / SGLang backend**: Current backend is Ollama. vLLM integration (v0.4) for higher-throughput production deployments.
- **Momagrid native encryption**: Per-CTE payload encryption using destination agent's Ed25519 public key. Infrastructure already present; protocol extension in design.
- **Text2SPL compiler**: Fine-tune an LLM on (natural language task, optimal SPL DAG) pairs, with parallelism — minimum DAG critical-path length — as the training objective.

This produces a natural language frontend that emits execution-optimized SPL IR, completing the compiler stack described in Section 4.

- **Intelligent semantic chunking**: NLP-based boundary detection (topic models, coreference resolution) to minimize cross-chunk dependencies.
- **1M Token Benchmark**: Head-to-head comparison of Momagrid on a 20-node community grid vs. monolithic H100 cluster.
- **Proof of Useful Work / Output Verification**: In trustless WAN deployments, the hub cannot assume honest agent output. Cryptographic output verification — e.g., challenge-response spot-checks or result cross-validation across redundant agents — is a prerequisite for open community grids.
- **Companion papers** (separate submissions): (1) *WAN system deployment strategy* — Internet federation, peer-to-peer CTE routing, mobile NPU node protocol, and security hardening for trustless community grids. (2) *Moma Points and grid economics* — the reward ledger economic model, market-rate derivation of R, and cooperative stakeholder incentive design.

10 Conclusion

We have presented Momagrid: a decentralized network system that linearizes transformer attention complexity from $O(N^2)$ to $O(N \cdot k)$ through SPL's CTE-based semantic chunking and distributed dispatch. The key insight is that the quadratic attention bottleneck is an architectural choice of monolithic inference, not a physical constraint. By treating the context window as a partitioned data structure — analogous to Spark's RDD — and dispatching chunks to a heterogeneous grid of commodity GPUs, the system scales linearly with context length while maintaining semantic coherence through the hub's assembly phase.

The full implementation is available open-source. LAN experiments on a 3-node consumer grid are complete: 22/22 validation recipes passed, covering 1,077 dispatched tasks with 100% failover recovery and correct tier-aware routing. We believe this architecture represents a fundamental shift in how AI inference infrastructure is organized: from energy-hungry data centers toward a decentralized and openly accessible grid — as robust as Nature, and as wide as the Web.

Acknowledgements

The author thanks the Apache Spark open source project for pioneering the distributed data abstraction that directly inspired Momagrid's CTE-as-RDD design philosophy. AI assistants — Claude (Anthropic), Gemini (Google), and Z.ai (ZhiPu) — provided invaluable support throughout the research, implementation, and writing process. The author also thanks Dr. Jun Hu for bringing deep system architecture experience to the WAN deployment work that will push this project forward.

References

[1] Gong, W. G. (2026). *Structured Prompt Language: Declarative Context Management for LLMs*. arXiv:2602.21257.

[2] Vaswani, A., et al. (2017). *Attention is All You Need*. NeurIPS.

[3] Zaheer, M., et al. (2020). *Big Bird: Transformers for Longer Sequences*. NeurIPS.

[4] Beltagy, I., et al. (2020). *Longformer: The Long-Document Transformer*. arXiv:2004.05150.

[5] Choromanski, K., et al. (2021). *Rethinking Attention with Performers*. ICLR.

[6] Gu, A., & Dao, T. (2023). *Mamba: Linear-Time Sequence Modeling with Selective State Spaces*. arXiv:2312.00752.

[7] Dao, T., et al. (2022). *FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness*. NeurIPS.

[8] Liu, H., et al. (2023). *Ring Attention with Blockwise Transformers for Near-Infinite Context*. arXiv:2310.01889.

[9] Rajbhandari, S., et al. (2020). *ZeRO: Memory Optimizations Toward Training Trillion Parameter Models*. SC20.

[10] Shoeybi, M., et al. (2019). *Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism*. arXiv:1909.08053.

[11] Kwon, W., et al. (2023). *Efficient Memory Management for Large Language Model Serving with PagedAttention*. SOSP.

[12] Zaharia, M., et al. (2010). *Spark: Cluster Computing with Working Sets*. HotCloud.

[13] Bu, Z., et al. (2024). *ZigZag: Efficient Communication in Sequence Parallelism for Long Context LLM Training*. arXiv:2401.04093.

[14] Xiao, G., et al. (2024). *DuoAttention: Efficient Long-Context LLM Inference with Retrieval and Streaming Heads*. arXiv:2410.10819.

[15] Anderson, D. P., et al. (2002). *SETI@home: An Experiment in Public-Resource Computing*. *Communications of the ACM*, 45(11), 56–61.

[16] Xie, Z., et al. (2023). *SGLang: Efficient Execution of Structured Language Model Programs*. arXiv:2312.07104.

A GPU Reference Table

This appendix provides a reference snapshot of GPU models relevant to Momagrid deployments, organized by market segment. The table grounds the paper's sustainability argument (Section 8) with concrete data: the same inference workload that requires a $30,000 H100 in monolithic mode can be served by a cluster of $200–300 second-hand consumer cards via Momagrid's chunked dispatch.

All TPS figures are approximate for llama3-8B at fp16/bf16 via Ollama; actual throughput varies by prompt length, batch size, quantization, and system configuration. Prices reflect approximate market values as of early 2026. All models are NVIDIA unless otherwise noted.

A.1 Vintage Gaming GPUs (Pascal, 2016–2017)

Pascal-generation cards are the sustainability core of the reference cluster for this paper — hardware released nearly a decade ago, still capable of running 7B-parameter models at SILVER tier.

Model	VRAM (GB)	BW (GB/s)	TPS (7B)	Tier	Used Price ($)
GTX 1060	6	192	5–8	BRONZE	80–120
GTX 1070	8	256	7–10	BRONZE	100–150
GTX 1080	8	320	8–12	SILVER	120–180
GTX 1080 Ti	11	484	10–15	SILVER	200–300

Table 17: Pascal-generation NVIDIA GPUs (2016–2017). Memory type: GDDR5 (GTX 1060/1070), GDDR5X (GTX 1080/1080 Ti).

A.2 Mid-Generation Gaming GPUs (Turing / Ampere, 2018–2021)

RTX 20- and 30-series cards are widely available on the used market and cover SILVER through low PLATINUM throughput.

Model	VRAM (GB)	BW (GB/s)	TPS (7B)	Tier	Used Price ($)
RTX 2070	8	448	12–18	SILVER	150–220
RTX 2080 Ti	11	616	15–22	GOLD	280–400
RTX 3060	12	360	15–20	SILVER	180–250
RTX 3080	10	760	20–28	GOLD	300–450
RTX 3090	24	936	25–35	PLATINUM	550–750

Table 18: Turing/Ampere-generation NVIDIA GPUs (2018–2021). Memory type: GDDR6 (RTX 2070/2080 Ti/3060), GDDR6X (RTX 3080/3090).

A.3 Current Gaming GPUs (Ada Lovelace / RDNA 3, 2022–2024)

Current-generation consumer cards approach professional inference throughput at a fraction of the cost, with large VRAM variants (16–24 GB) enabling longer chunk sizes (k) per node.

Model	VRAM (GB)	BW (GB/s)	TPS (7B)	Tier	New Price ($)
RTX 4060 Ti	16	288	18–25	GOLD	450–500
RTX 4070	12	504	28–40	PLATINUM	550–600
RTX 4080 Super	16	736	40–55	PLATINUM	900–1,000
RTX 4090	24	1,008	55–80	PLATINUM	1,600–2,000
RX 7700 XT†	12	432	18–25	GOLD	330–380
RX 7900 XTX†	24	960	40–60	PLATINUM	800–950

Table 19: Ada Lovelace (NVIDIA) and RDNA 3 (AMD†) GPUs (2022–2024). Memory type: GDDR6/GDDR6X.

A.4 Professional / Data Center GPUs

Included for comparison only. NVIDIA A100/H100 use HBM (High Bandwidth Memory) — the primary driver of their cost premium — and require specialized server infrastructure. Google TPUs are listed for reference; they are not currently supported by Ollama and are out of scope for the LAN grid described in this paper.

Model	VRAM (GB)	BW (GB/s)	TPS (7B)	Tier	Price ($)
Tesla T4	16	300	20–35	GOLD	1,500–2,500 used
A100 PCIe	40	1,555	80–120	PLATINUM	10,000–12,000
A100 SXM4	80	2,000	100–150	PLATINUM	15,000–20,000
H100 PCIe	80	2,000	150–200	PLATINUM	25,000–30,000
H100 SXM5	80	3,350	200–300+	PLATINUM	30,000–35,000
TPU v4‡	32	1,200	n/a	—	cloud only
TPU v5e‡	16	1,600	n/a	—	cloud only

Table 20: Professional/data center accelerators. NVIDIA A100/H100 use HBM2e/HBM3/HBM3e. ‡Google TPU v4/v5e listed for reference only; not supported by Ollama and not eligible for Momagrid grid participation in the current implementation.

Key observations:

1. **Cost ratio**: A single H100 SXM5 (∼$30,000) costs the same as 100–150 GTX 1080 Ti cards (∼$200–300 each). A 100-node Momagrid grid of GTX 1080 Ti cards handles contexts 100× longer than a single H100 can serve monolithically, while routing around any single-node failure.

2. **Memory bandwidth gap**: HBM cards (A100/H100) deliver 2,000–3,350 GB/s vs. 484 GB/s for a GTX 1080 Ti. This gap matters for dense cross-token attention but is largely irrelevant for chunked dispatch — each node only attends within its chunk, where GDDR5X bandwidth is sufficient.

3. **AMD and Apple Silicon**: AMD GPUs (ROCm on Linux) are fully supported as Momagrid agents with no code changes. Apple M-series chips (unified memory, up to 128 GB on M4 Max) run Ollama via the Metal backend and function as SILVER–GOLD agents, extending Momagrid to laptop and desktop hardware without discrete GPUs.

4. **Sustainability**: Every card in the A.1 table was manufactured 7–9 years ago. Momagrid's hardware-agnostic design gives these cards a productive second life as inference infrastructure, directly countering the AI industry's trend toward rapid GPU obsolescence cycles.