# SPL 2.0: A Declarative Language for Agentic Workflow Orchestration

Wen G. Gong

2026-03-23

## Abstract

We present SPL 2.0, a declarative language for orchestrating LLM-powered agentic workflows. Unlike imperative frameworks (LangGraph, AutoGen, CrewAI) that require Python programming and manual orchestration, SPL 2.0 expresses workflow intent in SQL-like syntax — the specification layer between natural language and execution. The stack comprises four layers: Natural Language (interface), a text2SPL compiler, SPL 2.0 (declarative specification), and a pluggable runtime supporting local (Ollama), cloud (OpenRouter, Claude), and decentralized (Momagrid) execution. Three novel primitives drive the design: EVALUATE for hybrid deterministic-semantic branching, WHILE for loop control, and EXCEPTION for error-handling.

A fourth design principle distinguishes SPL 2.0 from prior declarative approaches: the explicit separation of **CALL** (deterministic Python functions) from **GENERATE** (LLM invocations). Every operation in an SPL workflow is unambiguously declared as either probabilistic or deterministic. This forces clarity at the language level that imperative frameworks cannot provide, and makes the `@spl_tool` decorator a Foreign Function Interface: any Python callable — database queries, REST APIs, ML inference, file I/O — becomes a first-class CALL target with a single decorator line. A 55-function Standard Library (`spl.stdlib`) provides SQL-equivalent built-ins (type conversion, string, numeric, JSON, date) auto-loaded without any `--tools` flag.

The declarative approach yields three compounding benefits beyond LOC reduction (5–7× versus imperative frameworks): **(1) human reviewability** — business stakeholders can read and approve an SPL script without executing it, because intent is explicit in the syntax rather than buried in control flow; **(2) declarative operations** — the same `.spl` file promotes unchanged through Dev, Test, and Production by swapping CLI flags (`--adapter`, `--model`, `--tools`, `--datasets`), mirroring the 12-factor app principle of environment-as-configuration; **(3) compiler-level optimization** — because intent is explicit, future passes can merge redundant LLM calls, cache contexts, and route to cost-optimal models in ways structurally unavailable to imperative code.

We validate the language through a cookbook of 40 recipes covering 8 agentic workflow categories, achieving a 100% success rate (40/40 recipes, total 1,493.4 seconds) on a single GTX 1080 Ti GPU running gemma3 via Ollama — including, for the first time, the text2SPL compiler recipe (recipe 22) that closes the natural-language-to-SPL loop end-to-end. The `spl-cookbook` package and Recipe Kitchen browser UI make the cookbook accessible as a pip-installable tool. The reference implementation is publicly available.

**Keywords**: declarative language, agentic workflows, LLM orchestration, structured prompt lan-

guage, semantic evaluation

---

# 1. Introduction

## 1.1 The Problem: Fragmentation in LLM Programming

Building agentic AI systems today demands a daunting intersection of skills. A practitioner must master prompt engineering for effective LLM interaction, Python programming for orchestration logic, API wrangling across disparate model providers, and manual state management for multi-step reasoning. Each popular framework — LangGraph [1], AutoGen [2], CrewAI [3], Auto-GPT [16] — introduces its own abstractions, API surfaces, and execution models, creating a fragmented ecosystem where knowledge transfers poorly and vendor lock-in is the norm.

This situation echoes a pattern the data industry has seen before. In the early days of relational databases, programmers wrote imperative C and COBOL code to traverse data structures, manage cursors, and handle errors — all in application-specific ways. The introduction of SQL [4] in the 1970s transformed this landscape by providing a declarative specification layer: users described *what* data they wanted, not *how* to retrieve it. This separation enabled decades of optimizer innovation underneath a stable surface language. The current state of LLM programming is analogous to the pre-SQL era: powerful capabilities buried under layers of imperative glue code.

The barrier to entry compounds the problem. To build even a simple self-refining agent — one that drafts, critiques, and revises its own output — a developer must write 80–150 lines of Python across multiple framework abstractions. This excludes the vast population of domain experts, analysts, and researchers who understand their workflows but lack the software engineering skills to implement them as agentic systems.

## 1.2 The Opportunity: Declarative Agentic Workflows

Consider the world an AI application runs in today: a zoo of models — Llama, Gemma, Mistral, Claude, GPT, Qwen, Phi — deployed across a zoo of infrastructure: local GPUs in a developer's office, cloud APIs on three continents, community nodes contributed by universities and individuals worldwide. Each model has different capabilities, different context limits, different cost profiles. Each deployment has different latency, availability, and authentication requirements.

A student new to AI who wants to build a simple self-refining essay agent faces an immediate question that has nothing to do with their goal: *which model, on which infrastructure, accessed through which API, orchestrated with which Python framework?* Before writing a single line of logic, they must make a dozen infrastructure decisions. Most give up. Those who persist write code that is tightly coupled to one model, one provider, one framework — and must be rewritten when any of those change.

The opportunity SPL 2.0 addresses is this: **what if you could program against the entire zoo declaratively, with two primitives?**

```
SELECT  -- assemble what goes into the model
    system_role('You are a helpful assistant.'),
    context.question AS question
GENERATE answer(question)   -- invoke the model, capture what comes out
```

This is the complete I/O interface to any model, on any infrastructure. SELECT assembles the input. GENERATE invokes the model and captures the output. The runtime — not the programmer — decides which model in the zoo executes the inference, and where. The same two lines run on a local Ollama instance, an OpenRouter cloud API, or a Momagrid community node, without modification.

A declarative approach to this zoo offers three structural advantages. First, it enables **compiler-level optimization**: when the system understands intent rather than following imperative steps, it can merge redundant LLM calls, cache repeated contexts, and route queries to cost-optimal models — optimizations invisible to the user but not achievable by an imperative framework, which by construction interleaves logic with execution order and thereby eliminates the system's visibility into intent. Second, it enforces **separation of concerns** between what the workflow accomplishes and how it executes, allowing the same specification to run on local hardware, cloud APIs, or decentralized compute grids without modification. Third, it achieves **broad accessibility**: the SELECT/GENERATE mental model is simple enough for a student encountering AI programming for the first time, and powerful enough for production agentic workflows.

The deeper structural insight underlying SPL 2.0 is a parallel with Oracle's database programming stack:

- **SELECT/GENERATE** are to SPL what **SELECT** is to SQL — the atomic I/O contract with the backend (database rows, or LLM inference).
- **PROMPT** is the named, executable unit that binds them — analogous to a single SQL query.
- **WORKFLOW** is to SPL 2.0 what **PL/SQL** is to Oracle — the procedural extension that adds variables, loops, branching, and exception handling around declarative queries.

This is not an analogy by convenience; it is a deliberate design choice. The SQL-to-PL/SQL evolution solved the identical problem in a different domain: how to add procedural control flow to a declarative query language without sacrificing the optimizer's ability to reason about individual queries. SPL 2.0 applies this proven architectural pattern to LLM orchestration — across a global zoo of models that did not exist when SQL was designed.

### 1.3 Our Contribution: The SPL Stack

Figure 1 illustrates the three-tier SPL stack. SPL 2.0 is an **additive extension** of SPL 1.0: every valid SPL 1.0 program runs unchanged inside the SPL 2.0 runtime, and the new WORKFLOW constructs are purely additive.

We make the following contributions:

1. **A four-layer architecture** for LLM workflow execution:
    - *Natural Language* (interface) — the natural input modality
    - *text2SPL* (compiler) — translates natural language intent into executable specification
    - *SPL 2.0* (specification) — a declarative language with a formal EBNF grammar
    - *Pluggable runtime* — adapters for Ollama (local), OpenRouter (cloud), Claude CLI (Anthropic), and Momagrid (decentralized)
2. **Three novel language primitives** for agentic workflows:

```
┌──────────────────────────────────────────────────────────┐
│  Momagrid                                                 │
│  Decentralized Inference Network as Runtime               │
│  (GPU nodes contributed by independent operators)         │
├──────────────────────────────────────────────────────────┤
│  SPL 2.0                                                  │
│  Declarative Agentic Workflow Orchestration               │
│  (WORKFLOW, EVALUATE, WHILE, EXCEPTION, PROCEDURE)        │
├──────────────────────────────────────────────────────────┤
│  SPL 1.0  ← foundation, fully contained in SPL 2.0       │
│  Declarative Context Management                           │
│  (PROMPT, SELECT, GENERATE)                               │
└──────────────────────────────────────────────────────────┘
```

Figure 1: The SPL stack. SPL 2.0 contains SPL 1.0 as its foundation layer. Momagrid provides a decentralized inference runtime that SPL 2.0 targets via the `--adapter momagrid` flag, alongside local (Ollama) and cloud (OpenRouter, Anthropic) runtimes.

- **EVALUATE** with hybrid deterministic-semantic condition evaluation, where the LLM itself serves as a runtime judge
- **WHILE** with semantic termination conditions, enabling loops that iterate until output quality is "satisfactory" rather than until a counter reaches N
- **EXCEPTION** for error-handling to address LLM-specific failure modes (HallucinationDetected, RefusalToAnswer, ContextLengthExceeded, BudgetExceeded, etc.)

3. **A reference implementation** consisting of a Python Click-based CLI, a recursive-descent parser, 5 runtime adapters, 231 automated tests, and a cookbook of 40 recipes spanning 8 categories of agentic workflow patterns.

4. **The CALL/GENERATE design principle**: an explicit language-level separation between deterministic Python operations (CALL) and probabilistic LLM invocations (GENERATE), with the `@spl_tool` decorator as a Foreign Function Interface to the Python ecosystem.

5. **SPL Standard Library** (`spl.stdlib`): 55 SQL-equivalent built-in functions (type conversion, string, numeric, conditional, JSON, date, hashing, list) auto-loaded in every workflow without configuration.

6. **The `spl-cookbook` package** and **Recipe Kitchen** Streamlit UI: a pip-installable distribution with `recipe-kitchen` and `bake-all` CLI entry points, making the full 40-recipe cookbook accessible without manual path configuration.

7. **Backward compatibility** with SPL 1.0 [5]: all existing SPL 1.0 PROMPT statements execute unchanged in SPL 2.0, and the new WORKFLOW constructs are purely additive extensions.

### 1.4 Paper Organization

Section 2 surveys the motivation and positions SPL 2.0 within the landscape of related work. Section 3 presents the language design — primitives, semantics, and formal grammar. Section 4 describes the implementation architecture. Section 5 discusses compilation and optimization opportunities. Section 6 evaluates expressiveness, performance, cost efficiency, and developer experience. Section 7

discusses strengths, limitations, and future directions. Section 8 concludes. Appendix E describes the Momagrid decentralized runtime architecture. Appendix F contains the complete cookbook and end-to-end experiment results. Appendix G presents the full batch run data.

# 2. Motivation and Related Work

## 2.1 The Evolution of Data Programming

The history of data programming offers a clear precedent for the transition SPL 2.0 proposes. In the 1970s, IBM's System R project demonstrated that a declarative query language (SQL) could match or exceed the performance of hand-coded navigational database access, because the declarative form exposed optimization opportunities invisible in imperative code [5]. By the 1980s, SQL had become the industry standard, and Oracle's PL/SQL extension demonstrated that procedural control flow (loops, variables, exception handling) could be layered on top of declarative queries without sacrificing the optimizer's ability to reason about individual SELECT statements [6].

The parallel to LLM programming is direct. Today's imperative frameworks — LangGraph's state graphs, AutoGen's conversation loops, CrewAI's role-based agents — are the equivalent of pre-SQL navigational code: powerful, flexible, and impossibly fragmented. SPL 2.0 proposes that the same declarative approach can occur for LLM orchestration, following the same evolutionary pattern:

1. **Atomic declarative queries** (SQL SELECT → SPL PROMPT/SELECT): express what you want from a single LLM call.
2. **Procedural extensions** (PL/SQL → SPL WORKFLOW): add variables, loops, branching, and error handling around those queries.
3. **Optimizer infrastructure** (Catalyst → SPL IR): enable the system to rewrite and optimize execution plans.

## 2.2 Current Approaches to Agentic Systems

The landscape of LLM orchestration systems can be organized along two axes: **declarative vs. imperative** paradigm and **query-level vs. workflow-level** scope.

### Imperative Frameworks

**LangGraph** [1] models agent workflows as directed graphs with explicit state management. Nodes represent computation steps, edges represent transitions, and state is passed through a mutable dictionary. The framework provides fine-grained control but requires significant Python expertise and offers no formal grammar or optimizer.

**AutoGen** [2] focuses on multi-agent conversation patterns where agents exchange messages. Its strength lies in natural dialogue flows, but orchestration logic is embedded in Python code and tightly coupled to the runtime.

**CrewAI** [3] introduces a role-based metaphor where specialized agents collaborate. While intuitive for team-like workflows, it remains Python-first and provides no path to compilation or optimization.

**PydanticAI** takes a code-first approach, rejecting complex abstractions in favor of standard Python with type hints. It prioritizes developer ergonomics over declarative specification.

### Declarative Systems

**LMQL** [7] (ETH Zurich, 2022) is the pioneer of SQL-like LLM interaction, combining declarative queries with output constraints (type, length, regex). LMQL achieves 26–85% cost savings through constrained decoding. However, it operates at the *query level* — individual prompts with output validation — and does not provide constructs for multi-step workflows, iterative refinement, or exception handling.

**DSPy** [8] (Stanford, 2023) introduces declarative modules (Predict, ChainOfThought) and a compiler that automatically optimizes prompts via few-shot bootstrapping. DSPy achieves remarkable results — GPT-3.5 outperforming expert-crafted demonstrations by up to 46%. However, DSPy's "declarative" is at the *module interface level*: it optimizes what prompts to use, not how to orchestrate multi-step workflows. It provides no loops, no exception handling, and no formal grammar.

**SGLang** [9] (UC Berkeley, 2023) provides a structured generation frontend with runtime optimizations (RadixAttention for KV cache reuse), achieving up to 6.4x throughput improvement. SGLang optimizes *inference execution*, not workflow orchestration — it operates below the abstraction level of agentic patterns.

### Specification Schemas

**Agent Spec** [10] (Oracle Labs, 2025) defines a YAML/JSON schema for specifying agents and workflows in a framework-agnostic way, with a standardized evaluation harness. Agent Spec solves *portability* (define once, run anywhere) but is a configuration format, not an executable language — it lacks control flow, iteration, and exception handling within the specification itself.

**SLANG** [11] (2025) proposes a minimal meta-language with three primitives (stake, await, commit) for formalizing multi-agent routing. SLANG's contribution is formalization of *coordination*, not computation — it does not address generation, evaluation, or error recovery.

### Concurrent Work

**PayPal** [12] (December 2025) presents a declarative system separating agent workflow specification from implementation, achieving 60% reduction in development time and 3x deployment velocity. The PayPal DSL focuses on enterprise pipeline configuration (data serialization, filtering, API orchestration) across multiple backend languages. It does not formalize LLM-specific primitives like semantic evaluation or hallucination handling.

**Auton** [13] (February 2026) formalizes agent execution as an augmented POMDP with a declarative "Cognitive Blueprint" schema. Auton's contribution is governance and safety constraint enforcement via policy projection. It specifies agents as configuration artifacts (YAML/JSON) rather than providing an executable programming language.

**Agentics 2.0** [14] (IBM, March 2026) introduces a logical transduction algebra where LLM calls are typed semantic transformations composable via algebraically grounded operators in asynchronous Map-Reduce programs. This is the closest concurrent work in terms of formalism, but takes an algebraic rather than SQL-like approach.

### Workflow Orchestration and Adjacent Systems

Several established systems influence the workflow orchestration landscape but do not address LLM-specific concerns:

**LangChain LCEL** provides a declarative chain composition model via a Runnable interface, supporting piped composition of prompt templates, models, and output parsers. LCEL's composability is powerful for linear chains but does not provide iteration primitives, semantic condition evaluation, or a formal exception hierarchy for LLM-specific failures.

**Haystack Pipelines** (deepset) offers YAML-declarative pipeline definitions for document retrieval and question-answering workflows. Like Agent Spec, Haystack is a configuration format: it declares components and connections but delegates control flow and error recovery to the host application.

**Prefect / Airflow** are mature declarative workflow orchestrators for data engineering pipelines. Their abstractions (DAGs, tasks, flows) were designed for deterministic, side-effectful operations — not for semantic evaluation of LLM outputs, quality-based iteration, or hallucination handling. SPL 2.0 extends these orchestration patterns into the LLM domain.

**OpenAI Assistants API** provides built-in iteration and tool-use loops via a managed stateful conversation runtime. It abstracts control flow entirely from the developer, offering convenience at the cost of programmability: developers cannot express custom branching conditions, quality thresholds, or exception recovery policies.

None of these systems provide a programmable language with formal semantics for LLM-specific concerns. SPL 2.0 is distinguished by its formal EBNF grammar, explicit semantic evaluation primitive (EVALUATE), and structured LLM exception taxonomy (EXCEPTION).

**Positioning Summary**

**Table 1: Landscape Comparison**

| System | Paradigm | Scope | Has Grammar | Semantic Eval | LLM Exceptions | Multi-Runt |
|---|---|---|---|---|---|---|
| LangGraph | Imperative | Workflow | No | No | No | No |
| AutoGen | Imperative | Conversation | No | No | No | No |
| CrewAI | Imperative | Role-based | No | No | No | No |
| LMQL | Declarative | Query | Yes | No (structural) | No | No |
| DSPy | Declarative | Module | No | No (compile-time) | No | No |
| SGLang | Declarative | Inference | Yes | No | No | No |
| Agent Spec | Declarative | Configuration | Yes (schema) | No | No | Yes (portab |
| PayPal | Declarative | Pipeline | Yes | No | No | Yes |
| Agentics 2.0 | Algebraic | Data workflow | Yes (algebra) | No | No | No |
| **SPL 2.0** | **Declarative** | **Query + Workflow** | **Yes (EBNF)** | **Yes** | **Yes** | **Yes (5 ada** |

SPL 2.0 is the only system that combines a formal EBNF grammar, LLM-as-judge primitives (EVALUATE), an LLM-specific exception hierarchy (EXCEPTION), and a unified adapter layer spanning local, cloud, and decentralized execution — all within a single coherent language that spans both atomic queries and procedural workflows.

## 2.3 SQL Extensions and Their Limitations

A natural question is whether existing SQL extensions could serve the purpose SPL 2.0 addresses. We analyze three candidates:

**Common Table Expressions (CTEs)** provide sequential composition of queries but offer no iteration, no branching on computed results, and no mechanism for invoking external computation

(like an LLM).

**Recursive CTEs** add iteration but with strictly deterministic termination conditions — the recursion must converge on a fixed point defined by the query algebra. Semantic termination ("continue until the output is coherent") is impossible.

**Stored Procedures (PL/SQL, T-SQL)** provide full procedural control flow but are designed for database operations, not LLM invocation. They have no concept of probabilistic output, semantic evaluation, or model routing.

**Table 2: SQL Extension Capability Matrix**

| Capability | SQL CTE | Recursive CTE | PL/SQL | **SPL 2.0** |
|---|---|---|---|---|
| Sequential composition | Yes | Yes | Yes | Yes |
| Iteration | No | Yes (fixed) | Yes | Yes (semantic) |
| LLM invocation | No | No | No | Yes |
| Semantic evaluation | No | No | No | Yes |
| Exception handling | No | No | Yes | Yes (LLM-aware) |
| Model routing | No | No | No | Yes |

### 2.4 Lessons from Proven Architectures

SPL 2.0's design draws deliberately from three proven systems:

**From Oracle PL/SQL**: block structure (`DO...END`), exception handling with typed exceptions (`WHEN HallucinationDetected THEN`), procedures and packages for code reuse, and the fundamental insight that procedural extensions should layer on top of declarative queries without disturbing the query optimizer.

**From Apache Spark**: the separation between logical and physical execution plans, the concept of a query optimizer (Catalyst) that rewrites execution strategies, and the idea that the same logical plan can execute on different backends (local, YARN, Kubernetes).

**From the JVM**: the stable intermediate representation (bytecode) that enables multiple frontend languages and multiple backend targets, and the principle that a well-defined IR is the foundation for an optimization ecosystem.

## 3. SPL 2.0 Language Design

### 3.1 Design Principles

SPL 2.0 is governed by five design principles:

**P1. Declarative intent over imperative implementation.** The programmer specifies *what* the workflow should accomplish — not *which model*, *which node*, or *which execution path* achieves it. This is the essential simplification: the zoo of models deployed around the world is the runtime's problem, not the programmer's. A student writing their first AI workflow and a senior engineer deploying a production pipeline write the same SELECT/GENERATE syntax; the runtime resolves the infrastructure details beneath them. This separation enables the compiler and runtime to make optimization decisions that are invisible to the user but structurally unavailable to an imperative framework, where execution order is hard-coded by the programmer rather than derived from declared intent.

**P2. LLM-aware primitives.** Rather than repurposing generic control flow constructs for LLM orchestration, SPL 2.0 provides primitives designed specifically for the unique characteristics of LLM computation: probabilistic output, semantic quality assessment, hallucination risk, context length limits, and cost sensitivity.

**P3. Compilation target.** The language is designed to produce an optimizable intermediate representation (IR), enabling future optimizer passes without changing the surface language.

**P4. Unified runtime — one specification, any model, any infrastructure.** The adapter pattern decouples the language from any specific model provider. The same `.spl` file executes on a local Ollama instance (a developer's laptop), OpenRouter (100+ cloud models), Claude CLI (Anthropic), or a Momagrid community node (a contributed GPU anywhere in the world) — without changing a single line. The programmer writes against the zoo; the adapter layer tames it. This is not merely an engineering convenience — it is the architectural prerequisite for broad accessibility.

**P5. Backward compatibility.** All SPL 1.0 [4] programs are valid SPL 2.0 programs. New constructs (WORKFLOW, EVALUATE, WHILE, EXCEPTION) are purely additive extensions.

**P6. Explicit probabilistic/deterministic boundary (CALL vs. GENERATE).** Every operation in an SPL workflow declares its computational character at the language level. `GENERATE` invokes an LLM — probabilistic, creative, slow, expensive. `CALL` invokes a Python function — deterministic, testable, instant, free. Most LLM frameworks blur this line; SPL draws it in ink. This forces workflow authors to be explicit about which operations can be cached, replayed, or optimized, and which cannot. As a corollary, the `@spl_tool` decorator serves as a Foreign Function Interface (FFI): any Python callable becomes a first-class `CALL` target with one decorator line, connecting SPL to the entire Python ecosystem without extending the language itself.

## 3.2 Core Primitives

The mental model of SPL 2.0 is grounded in two fundamental primitives that mirror the I/O contract with an LLM: **SELECT** assembles the context that flows *into* the model; **GENERATE** invokes the model and captures what flows *out*. Every other construct in the language is either a named container for SELECT and GENERATE (PROMPT, WORKFLOW) or a flow control mechanism that orchestrates when and how they are invoked (EVALUATE, WHILE, EXCEPTION, PROCEDURE).

This two-sided I/O model deliberately parallels SQL's interaction with a database — SELECT gathers data in, the result set flows out — while extending it to handle the probabilistic, semantic nature of LLM computation.

### 3.2.1 SELECT — Context Assembly (Input to LLM)

SELECT is the *input* primitive. It aggregates context from multiple sources into the prompt payload that will be sent to the LLM. Nothing executes until SELECT has finished assembling the input.

```
SELECT
    system_role('You are a research analyst. Be precise and cite evidence.'),
    context.topic AS topic,
    memory.get('prior_findings') AS background,
    rag.query('recent developments in ' + @topic, top_k=5) AS references
```

SPL 2.0 provides four context sources within SELECT:

| Source | Syntax | What it provides |
|---|---|---|
| System role | `system_role('...')` | The model's behavioral instruction |
| Runtime parameter | `context.field_name` | CLI-injected value at execution time |
| Persistent memory | `memory.get('key')` | Key-value store across workflow runs |
| Semantic retrieval | `rag.query('text', top_k=N)` | Top-N chunks from a vector index |

Token budgets on the input side are enforced with `WITH BUDGET n TOKENS`, preventing context window overflow before the LLM is ever invoked.

The analogy to SQL is precise: SQL's `SELECT ... FROM ... WHERE` gathers and filters rows from a database; SPL's `SELECT` gathers and filters context from heterogeneous sources (parameters, memory, retrieval indices) into a prompt. In both cases the reader specifies *what information to collect* — not how the storage layer retrieves it.

### 3.2.2 GENERATE — LLM Invocation (Output from LLM)

GENERATE is the *output* primitive. It names the generation task, invokes the LLM with the assembled SELECT context, and captures the result.

```
GENERATE answer(prompt)                          -- standalone: output to stdout
GENERATE analysis(@research) INTO @result        -- workflow: capture into variable
GENERATE summary(@analysis) INTO @summary
    WITH OUTPUT BUDGET 500 TOKENS                -- cap output length
    USING MODEL 'claude-opus-4-6'               -- model override per step
```

The argument passed to GENERATE — `answer`, `analysis`, `summary` — is a semantic label that tells the LLM what kind of output to produce. This is the declarative analog of a function name: the developer states *what* the model should generate, not *how*.

Key behaviors: - **Standalone** (`GENERATE x()`): output goes to stdout; used inside PROMPT statements. - **Capturing** (`GENERATE x()INTO @var`): output is stored in a workflow variable for downstream steps. - **Output budget** (`WITH OUTPUT BUDGET n TOKENS`): limits response length, enabling cost control. - **Model binding** (`USING MODEL '...'`): per-step model override without editing the file; the CLI `--model/-m` flag applies globally across all GENERATE calls at runtime.

GENERATE has no direct SQL equivalent — SQL does not produce probabilistic text — but the pattern it enables is familiar: `GENERATE` is to SPL what a function call is to PL/SQL, except the "function" is the LLM itself.

Together, SELECT and GENERATE form the complete I/O interface to the LLM. SELECT is the *input contract*; GENERATE is the *output contract*. PROMPT and WORKFLOW are simply named scopes that bind them together.

### 3.2.3 PROMPT — Atomic LLM Query

The PROMPT statement is the fundamental named unit of SPL 2.0: a SELECT clause followed by a GENERATE clause, given a name and made independently executable.

```
PROMPT hello_world
SELECT
    system_role('You are a friendly assistant. Introduce yourself and SPL 2.0 in two sentences.')
GENERATE greeting()
```

This four-line program is a complete, executable SPL file. SELECT assembles the system instruction; GENERATE invokes the LLM and streams the result to stdout. No imports, no configuration, no API keys in code.

Parametric context is supported via `context.param`:

```
PROMPT ollama_proxy
SELECT
    system_role('You are a helpful, knowledgeable assistant.'),
    context.prompt AS prompt
GENERATE answer(prompt)
```

Runtime parameter injection via the CLI: `spl2 run proxy.spl --adapter ollama prompt="Explain quantum computing"`. The model can be overridden at runtime with `--model/-m` without editing the `.spl` file — the CLI flag patches the GENERATE binding at execution time.

PROMPT is to SPL 2.0 what a single SQL SELECT statement is to SQL: the atomic, declarative query unit. WORKFLOW (Section 3.2.4) is what PL/SQL is to SQL: the procedural extension that adds state, loops, and branching around that query unit.

### 3.2.4 WORKFLOW — Procedural Orchestration

The WORKFLOW statement adds procedural control flow around PROMPT-level operations, analogous to how PL/SQL extends SQL with variables, loops, and exception handling.

```
WORKFLOW chain_of_thought
    INPUT: @topic TEXT
    OUTPUT: @summary TEXT
DO
    GENERATE research(@topic) INTO @research
    GENERATE analyze(@research) INTO @analysis
    GENERATE summarize(@analysis) INTO @summary
    COMMIT @summary WITH status = 'complete'
END
```

Key constructs: - **INPUT/OUTPUT** declarations define the workflow's typed interface. - **@variables** (prefixed with `@`) provide mutable state within the workflow scope. - **GENERATE...INTO** captures LLM output into a variable for use in subsequent steps. - **COMMIT** finalizes the workflow output with optional metadata. Execution halts after COMMIT.

### 3.2.5 EVALUATE — LLM-Based Condition Evaluation

EVALUATE is SPL 2.0's first novel primitive. It provides branching where the condition itself may be evaluated by an LLM — a capability no existing language formalizes.

SPL 2.0 supports three evaluation types, detected automatically from the condition syntax:

**Deterministic evaluation** (zero LLM cost):

```
EVALUATE @score
    WHEN > 0.8 THEN COMMIT @result WITH status = 'high_quality'
    WHEN > 0.5 THEN GENERATE improved(@result) INTO @result
    OTHERWISE RETRY WITH temperature = 0.1
END
```

When the condition is a numeric or boolean comparison, the runtime evaluates it directly with no LLM call.

**Semantic evaluation** (LLM as judge):

```
EVALUATE @summary
    WHEN 'coherent' THEN COMMIT @summary
    WHEN 'hallucination' THEN RETRY
    WHEN 'incomplete' THEN GENERATE expand(@summary) INTO @summary
END
```

When the condition is a quoted string literal, the runtime invokes the LLM as a judge: "Does this output satisfy the criterion 'coherent'?" This is a low-cost classification call (typically 50–200 tokens) that enables the workflow to make quality-based decisions without external tooling.

**Hybrid evaluation** (mixed):

```
EVALUATE @response
    WHEN token_count > 4000 THEN truncate          -- deterministic
    WHEN 'off_topic' THEN reframe                   -- semantic
    WHEN confidence < 0.7 THEN verify               -- deterministic
END
```

The detection rule is simple and unambiguous: if the condition is a quoted string, it is semantic; otherwise, it is deterministic. This design choice avoids the need for explicit annotations while maintaining clear semantics.

**Semantic Evaluation Protocol.** When a semantic condition is detected, the runtime sends a standardized two-message judge call to the LLM. The system prompt is: *"You are a quality judge. Answer ONLY 'true' or 'false' with no additional text."* The user prompt is: *"Does the following output satisfy the criterion '<condition>'?"* followed by the value being evaluated. The runtime extracts the first token of the response and maps it to a boolean; any non-boolean response is treated as `false` and may raise a `QualityBelowThreshold` exception. By default, the judge call uses the same adapter and model as the enclosing GENERATE statement, but a dedicated judge model can be specified per-EVALUATE via `USING MODEL`.

### 3.2.6 WHILE...DO — Iteration with Semantic Termination

WHILE is SPL 2.0's second novel primitive. It enables loops where the termination condition can be LLM-judged, with a built-in safety guard against infinite iteration.

```
WHILE @iteration < @max_iterations DO
    GENERATE critique(@current) INTO @feedback

    EVALUATE @feedback
        WHEN 'satisfactory' THEN
            COMMIT @current WITH status = 'complete'
        OTHERWISE
```

```
            GENERATE refined(@current, @feedback) INTO @current
            @iteration := @iteration + 1
    END
END
```

The construct supports the same deterministic/semantic/hybrid condition types as EVALUATE. A global iteration limit (default: 10) is enforced by the runtime, raising `MaxIterationsReached` if exceeded. This prevents runaway loops while allowing the workflow to iterate based on output quality rather than arbitrary counts.

Context accumulation across iterations is a natural pattern: each iteration's output becomes the input for the next, enabling progressive refinement without explicit history management.

### 3.2.7 EXCEPTION — Structured LLM Error Handling

EXCEPTION is SPL 2.0's third novel primitive. It defines a formal taxonomy of LLM-specific failure modes, each with clear semantics and recovery strategies.

**Table 3: LLM Exception Taxonomy**

| Exception | Trigger Condition | Typical Recovery |
|---|---|---|
| HallucinationDetected | Output fails fact-checking | Regenerate with lower temperature |
| RefusalToAnswer | Model declines the request | Log and COMMIT partial result |
| ContextLengthExceeded | Input exceeds model context window | Compress context and RETRY |
| ModelOverloaded | Provider rate limit or capacity error | Wait and RETRY |
| BudgetExceeded | Token budget exhausted | COMMIT best available result |
| QualityBelowThreshold | Output quality score too low | Regenerate or escalate to better model |
| MaxIterationsReached | WHILE loop hit iteration limit | COMMIT current best effort |

These exceptions are caught in structured handlers:

```
WORKFLOW safe_generation
    INPUT: @prompt text
    OUTPUT: @result text
DO
    GENERATE response(@prompt) INTO @result
    GENERATE quality_score(@result) INTO @score

    EVALUATE @score
        WHEN > 0.8 THEN COMMIT @result WITH status = 'high_quality'
        WHEN > 0.5 THEN
            GENERATE improved(@result, @prompt) INTO @result
            COMMIT @result WITH status = 'refined'
        OTHERWISE RETRY WITH temperature = 0.1
    END

EXCEPTION
```

```
    WHEN HallucinationDetected THEN
        GENERATE response(@prompt) INTO @result
            WITH TEMPERATURE 0.1
        COMMIT @result WITH status = 'conservative'
    WHEN ContextLengthExceeded THEN
        RETRY WITH temperature = 0.1
    WHEN RefusalToAnswer THEN
        COMMIT @prompt WITH status = 'refused'
    WHEN BudgetExceeded THEN
        COMMIT @result WITH status = 'truncated'
END
```

This 25-line workflow expresses a production-grade generation pattern — quality checking, iterative refinement, and graceful degradation across four failure modes — that would require 80–120 lines of Python in LangGraph with manual try/except blocks and custom error classes.

No existing language or framework formalizes LLM failure modes as a typed exception hierarchy. LangGraph and AutoGen rely on Python's generic exception mechanism; LMQL and DSPy have no error handling constructs at all.

### 3.2.8 PROCEDURE — Reusable Workflow Patterns

PROCEDUREs provide parameterized, reusable workflow components that can be composed via CALL:

```
PROCEDURE researcher(topic TEXT) RETURNS TEXT
DO
    GENERATE research_facts(topic) INTO @facts
    GENERATE identify_key_themes(@facts) INTO @themes
    @result := @facts + '\n\nKey Themes:\n' + @themes
    COMMIT @result
END
```

Procedures enable a standard library of agentic patterns: ReAct (reasoning + acting), Self-Refine (iterative improvement), Tree-of-Thought (branching exploration), and others can be implemented once and reused across workflows via `CALL researcher(@topic)INTO @result`.

### 3.2.9 Standard Library and User-Defined Functions (CALL Extension System)

**SPL Standard Library.** SPL 2.0 ships with `spl.stdlib`, a 55-function built-in library auto-loaded without any `--tools` flag or import statement. The functions mirror SQL built-ins, organized into 11 categories:

| Category | Representative functions |
|---|---|
| Type conversion | `to_float`, `to_int`, `to_text`, `to_bool` |
| String | `upper`, `lower`, `trim`, `substr`, `replace`, `concat` |
| Pattern matching | `like`, `startswith`, `endswith`, `regexp_match` |
| Numeric | `abs_val`, `round_val`, `ceil_val`, `clamp` |
| Conditional | `coalesce`, `iif`, `nullif` |
| Null/empty checks | `isnull`, `isblank`, `nvl` |
| Text aggregates | `word_count`, `char_count`, `line_count` |
| JSON | `json_get`, `json_set`, `json_keys`, `json_length` |

| Category | Representative functions |
|---|---|
| Date/time | `now_iso`, `date_format_val`, `date_diff_days` |
| Hashing | `md5_hash`, `sha256_hash` |
| List/array | `list_get`, `list_length`, `list_contains` |

The `to_float` function is representative of the stdlib's design philosophy: it extracts the first numeric token from any string, so LLM output like `"Quality score: 0.85\n"` returns `0.85` without a secondary LLM call. This is the CALL/GENERATE boundary in practice — the computation is deterministic and belongs in CALL, not GENERATE.

**User-Defined Functions (UDFs).** The `@spl_tool` decorator is SPL's Foreign Function Interface. Any Python callable decorated with `@spl_tool` becomes a first-class `CALL` target, connecting SPL to the entire Python ecosystem:

```
from spl.tools import spl_tool

@spl_tool
def lookup_customer(customer_id: str) -> str:
    row = db.query("SELECT name, tier FROM customers WHERE id = ?",
                   customer_id)
    return json.dumps(row) if row else ""

CALL lookup_customer(@id) INTO @customer_json   -- deterministic, zero LLM cost
GENERATE write_response(@customer_json) INTO @reply  -- LLM for the language part
```

Database queries, REST API calls, ML inference (scikit-learn, transformers), file I/O, PII detection, embedding computation — all become `CALL` targets with a single decorator line. The `--tools tools.py` CLI flag loads them at runtime; the language itself requires no extension.

The dispatch order for `CALL foo(@arg)` is: (1) registered UDF; (2) SPL PROCEDURE defined in the same file; (3) LLM fallback (treated as GENERATE). UDFs take priority, enabling domain-specific overrides of stdlib functions.

## 3.3 Cross-Cutting Concerns

### 3.3.1 Context Management

SPL 2.0 provides four context sources within the SELECT clause:

| Source | Syntax | Description |
|---|---|---|
| System role | `system_role('description')` | Sets the system prompt |
| Parameter | `context.field_name` | Runtime parameter via CLI |
| Memory | `memory.get('key')` | Persistent key-value store |
| RAG | `rag.query('text', top_k=5)` | Semantic search over vector store |

Token budgets can be specified at both input (`WITH BUDGET n TOKENS`) and output (`WITH OUTPUT BUDGET n TOKENS`) levels.

### 3.3.2 Security and Accounting

For deployments on untrusted infrastructure (e.g., Momagrid's decentralized grid), SPL 2.0 supports security classifications and accounting tags:

```
WORKFLOW sensitive_analysis
    INPUT: @data TEXT
    OUTPUT: @result TEXT
    SECURITY: CLASSIFICATION: confidential
    ACCOUNTING: BILLABLE_TO: 'project-alpha'
DO
    ...
END
```

These metadata blocks are preserved through the compilation pipeline and enforced by the runtime adapter.

## 3.4 Formal Grammar

The complete EBNF grammar is provided in Appendix A. The grammar defines 12 statement types, 3 condition types (deterministic, semantic, hybrid), 5 primitive types (text, number, boolean, chunk, json), composite types (list, map), and metadata blocks for security, accounting, and compression.

SPL 2.0 extends SPL 1.0 with 8 new statement types (WORKFLOW, EVALUATE, WHILE, COMMIT, RETRY, RAISE, assignment, PROCEDURE) and 1 new condition type (semantic). All SPL 1.0 productions remain unchanged.

# 4. Implementation

## 4.1 Architecture Overview

The SPL 2.0 reference implementation is organized as a four-layer stack:

```
┌──────────────────────────────────────────┐│

  Natural Language / .spl files │←
      Interface├──────────────────────────────────────────┤│

  spl2 CLI (Click) │←                            Tooling│├──
   run / execute / explain / parse │└├──
   --adapter / --model / KEY=VALUE │└──
   init / memory / rag │├──────────────────────────────────────────┤│

  Lexer → Parser → AST → IR (JSON) │←
      Compiler├──────────────────────────────────────────┤│

  Adapter Layer │←                                Runtime│├──
   echo     (testing, no LLM) │└├──
   ollama   (local GPU) │└├──
   openrouter (cloud, 100+ models) │└├──
   claude_cli (Anthropic Claude) │└──
   momagrid  (decentralized grid) │└──────────────────────────────────────────┘
```

**Figure 1**: SPL 2.0 system architecture. The four layers correspond to the conceptual stack described in Section 1.3.

### 4.2 The spl2 CLI

The CLI is built on Python's Click framework, providing a consistent interface for all operations:

```
# Zero-config execution
spl2 run cookbook/01_hello_world/hello.spl

# With adapter and model override
spl2 run cookbook/02_ollama_proxy/proxy.spl --adapter ollama -m gemma3 \
    prompt="Explain quantum computing"

# Parse and validate syntax (no LLM needed)
spl2 parse cookbook/05_self_refine/self_refine.spl --json

# Show execution plan
spl2 explain cookbook/09_chain_of_thought/chain.spl
```

Key design decisions: - **Trailing KEY=VALUE parameters** without a `-p` flag, using Click's UNPROCESSED argument handling. This makes the command line feel natural: `spl2 run file.spl topic="AI"` rather than `spl2 run file.spl -p topic="AI"`. - **Runtime model override** via `--model/-m` patches the model binding on all PROMPT statements at execution time, allowing model comparison without editing `.spl` files. - **Auto-initialization** creates workspace directories (`.spl2/`, memory store, RAG index) on first run.

### 4.3 Lexer and Parser

The compiler pipeline consists of a tokenizer and recursive-descent parser that produce an abstract syntax tree (AST).

The **tokenizer** handles SPL-specific challenges including keyword-as-identifier resolution (e.g., `context.prompt` where `prompt` is also a keyword), `@variable` prefixing, and the `:=` assignment operator. Comment support follows SQL convention (`--` for single-line).

The **parser** produces typed AST nodes: `PromptStatement`, `WorkflowStatement`, `GenerateClause`, `EvaluateClause`, `WhileStatement`, `ExceptionBlock`, `ProcedureDeclaration`, and others. Each node carries source location information for error reporting.

### 4.4 Adapter Architecture for LLM Model/Provider

The `--adapter` flag encompasses two conceptual dimensions that together define the complete execution environment for an SPL workflow:

- **LLM Model runtime** (two sub-dimensions):
  - `--adapter` (or `--llm-provider`): the backend infrastructure — `ollama` (local), `openrouter` (cloud), `anthropic`, `momagrid` (decentralized), etc.
  - `--model` (`-m`): the specific model within that provider — `gemma3`, `claude-sonnet-4-6`, `gpt-4o`, etc.
- **Integration patterns** (planned extensions):

- `--tools`: external tool connectors (web search, code execution, PDF ingestion, transcription)
- `--datasets`: data source bindings for RAG or batch processing workflows
- `--resources`: named resource slots (memory stores, vector indices, external APIs)

Together, the LLM runtime dimension and the integration pattern dimension provide comprehensive adapter coverage: the same `.spl` file can route inference to any provider and pull context from any data source, purely through CLI flags — without modifying the SPL source.

Runtime adapters implement a common interface:

```
class BaseAdapter:
    def generate(self, prompt, system_role, model, **kwargs) -> Result
    def list_models(self) -> list[str]
    def count_tokens(self, text) -> int
```

New adapters are registered via `register_adapter(name, cls)`, allowing third-party extensions without modifying core code.

**Table 4: Adapter Capabilities**

| Adapter | Generate | Streaming | List Models | Model Override | Auth |
|---|---|---|---|---|---|
| echo | Yes | No | Yes (static) | Yes | None |
| ollama | Yes | Planned | Yes (API) | Yes | None |
| openrouter | Yes | Planned | Yes (API) | Yes | API key |
| claude_cli | Yes | Planned | Yes (static) | Yes | CLI auth |
| momagrid | Yes | No | Yes (fallback) | Yes | API key |

The **echo adapter** deserves special mention: it returns deterministic responses without invoking any LLM, enabling the entire test suite (231 tests) to run without GPU hardware or API credentials. This is critical for CI/CD and for validating the parser and execution engine in isolation.

### 4.5 Momagrid Runtime Integration

The Momagrid [15] adapter connects SPL 2.0 to a decentralized compute grid where GPU nodes are contributed by independent operators. The adapter uses a hub-and-spoke protocol:

1. **Submit**: POST /tasks with the prompt, model, and node requirements (min_tier, min_vram_gb).
2. **Poll**: GET /tasks/{id} until the task reaches a terminal state.
3. **Retrieve**: Extract the generated text from the completed task.

The async httpx client supports configurable poll intervals and timeout limits. Health checks and agent listing are available via helper methods for monitoring grid availability.

### 4.6 The `spl-cookbook` Package and Recipe Kitchen UI

**Distribution.** The cookbook is distributed as the `spl-cookbook` pip package, providing two CLI entry points:

```
pip install spl-cookbook
```

```
recipe-kitchen   # launch the Streamlit browser UI
bake-all         # run all 40 recipes and save a timestamped log
```

Path resolution follows a three-level fallback: environment variable `SPL_KITCHEN_APP` override; bundled application module; dev layout walk-up (for editable installs). Optional extras install provider-specific SDKs: `spl-cookbook[anthropic]`, `spl-cookbook[openai]`, `spl-cookbook[all]`.

**Recipe Kitchen.** The Recipe Kitchen is a Streamlit application that makes all 40 cookbook recipes accessible from a browser — no terminal, no flags, no adapter syntax. It provides three tabs:

- **Playground** — filter by category, select a recipe, edit input parameters as text fields, press Run. The base command is shown read-only so the executed command is always visible.
- **Recipe Maker** — describe a workflow in plain English; the text2SPL compiler (Recipe 22) generates a `.spl` file; edit in the inline editor; run or save.
- **Bake-All** — run all active recipes via `run_all.py`, save a timestamped log, browse historical logs with pass/fail counts.

The Recipe Kitchen is  350 lines of Python. It is a local tool designed for developer and team use; deploying it on a shared machine with Streamlit's `--server.address` flag enables team access without public deployment.

The catalog (`cookbook_catalog.json`) is the single source of truth for the Playground and Bake-All. Adding a recipe to the catalog makes it appear in both views on the next page load — no UI changes required.

# 5. Compilation and Optimization

## 5.1 The text2SPL Compiler

The top layer of the SPL stack — translating natural language intent into executable SPL 2.0 code — has a working prototype accessible via `spl2 text2spl`. The current prototype handles simple PROMPT pattern generation: given a natural language description, it produces a syntactically correct PROMPT statement with appropriate SELECT and GENERATE clauses. For example:

```
$ spl2 text2spl "summarize a document with a 2000 token budget" \
    --adapter ollama -m gemma3 --mode prompt
```

produces a syntactically valid `.spl` file that `spl2 parse` accepts without error. The prototype also supports `--mode workflow` for WORKFLOW generation and `--mode auto` where the model decides the appropriate form; these modes are less mature and their output requires manual review before use.

The full compiler pipeline envisioned for production would:

1. Parse natural language descriptions of workflow intent.
2. Extract the agentic pattern (self-refine, ReAct, chain-of-thought, etc.).
3. Generate syntactically correct SPL 2.0 code, including full WORKFLOW structures.

Full WORKFLOW generation from natural language — "I want an agent that researches a topic, analyzes the findings, and writes a summary" — remains future work. The declarative form makes this tractable: because SPL 2.0 specifies intent rather than execution steps, the mapping from natural language to SPL is structurally simpler than generating imperative Python code.

### 5.2 Intermediate Representation (IR)

The current implementation compiles SPL 2.0 source to an AST, which is serialized as JSON via the `spl2 explain` command. The planned IR extends this with:

- **Execution cost annotations**: estimated token consumption per GENERATE node.
- **Dependency edges**: data flow between GENERATE steps for parallelism analysis.
- **Adapter hints**: preferred runtime for each node based on model requirements.

### 5.3 Future Optimization Opportunities

This section describes optimization strategies that are *structurally enabled* by the declarative form — not yet implemented, but architecturally possible in ways that are impossible in imperative frameworks where execution order is hard-coded by the programmer. The declarative form is the prerequisite; the optimizer is future work.

The key claim: an imperative framework cannot perform these optimizations because the programmer's control flow interleaves logic with execution sequence, eliminating the system's visibility into intent. SPL 2.0's separation of specification from execution preserves that visibility.

**GENERATE merging**: When consecutive GENERATE calls share the same context, they can be merged into a single LLM call with structured output parsing, reducing both latency and cost.

**Deterministic condition hoisting**: In a WHILE loop, deterministic conditions in EVALUATE blocks can be evaluated before invoking the LLM for semantic conditions, avoiding unnecessary LLM calls.

**SELECT caching**: When the same RAG query or memory lookup appears multiple times, the result can be cached and reused.

**Model selection optimization**: A cost-quality optimizer can route simple GENERATE calls to smaller, cheaper models while reserving large models for complex reasoning steps.

**Parallel GENERATE**: When consecutive GENERATE steps have no data dependencies (e.g., generating pro and con arguments in a debate), they can execute in parallel.

**Table 5: Optimization Opportunity Analysis**

| Optimization | Applicable Patterns | Estimated Impact | Status |
|---|---|---|---|
| GENERATE merging | Chain-of-Thought, Multi-Agent | Reduce LLM calls | Planned |
| Condition hoisting | Self-Refine, Reflection | Reduce judge calls | Planned |
| SELECT caching | RAG Query | Eliminate redundant retrieval | Planned |
| Model downgrade | Hello World, Proxy | Cost reduction | Planned |
| Parallel GENERATE | Debate, Tree-of-Thought, Ensemble | Latency reduction | Planned |

### 5.4 Logical Plan to Physical Plan

The `spl2 explain` command visualizes the execution plan for any `.spl` file, showing the sequence of operations, data dependencies, and estimated costs. This is directly analogous to SQL's EXPLAIN

ANALYZE and provides transparency into how the runtime will execute the workflow.

# 6. Evaluation

## 6.1 Expressiveness

We compare SPL 2.0's expressiveness against imperative frameworks by measuring lines of code (LOC) required to implement common agentic patterns.

**Table 6: Agentic Pattern Expressiveness (Lines of Code)**

| Pattern | SPL 2.0 | LangGraph (est.) | AutoGen (est.) | CrewAI (est.) |
|---|---|---|---|---|
| Hello World | 4 | ~15 | ~20 | ~25 |
| LLM Proxy | 5 | ~20 | ~25 | ~30 |
| Parametric Context | 6 | ~25 | ~30 | ~30 |
| Chain of Thought | 14 | ~50 | ~60 | ~55 |
| Self-Refine | 35 | ~100 | ~120 | ~110 |
| ReAct Agent | 43 | ~120 | ~100 | ~90 |
| Safe Generation | 37 | ~90 | ~110 | N/A |
| RAG Query | 7 | ~40 | ~45 | ~50 |

SPL 2.0 achieves 3–6x reduction in LOC compared to imperative alternatives. More significantly, the SPL code is *readable by non-programmers* — the SQL-like syntax is familiar to analysts, data scientists, and researchers who would not write Python framework code.

## 6.2 Performance

**Table 7: Implementation Metrics**

| Metric | Value | Notes |
|---|---|---|
| Total unit tests | 231 | All passing |
| Cookbook recipes | 40 | All 40 benchmarked including recipe 22 (Text2SPL Demo), which was previously planned/skipped; 5 new recipes (36–40) added in v0.7 |
| Recipe categories | 8 | basics, agentic, reasoning, safety, retrieval, multi-agent, application, benchmarking |
| Runtime adapters | 5 | echo, ollama, openrouter, claude_cli, momagrid |
| SPL 2.0 statement types | 12 | PROMPT, WORKFLOW, GENERATE, EVALUATE, WHILE, COMMIT, RETRY, RAISE, PROCEDURE, CALL, assignment, EXCEPTION |
| Cookbook success rate | 40/40 (100%) | gemma3 via Ollama, single GTX 1080 Ti; all 40 recipes pass; recipe 22 now fully operational (was SKIP in v0.6) |

| Metric | Value | Notes |
|---|---|---|
| PROMPT recipe avg latency | 2.3s | 7 PROMPT recipes; range 0.6s–6.1s depending on output length |
| WORKFLOW recipe avg latency | 44.8s | 33 WORKFLOW recipes; range 3.7s–132.3s |
| Fastest workflow | 1.7s | Document Q&A (recipe 39); CALL + GENERATE minimal pipeline |
| Slowest workflow | 132.3s | Reflection Agent (recipe 16); deep meta-cognition loop |
| CALL-based recipe latency | 1.7s–8.7s | Recipes 36–40 (Tool-Use, Headline News, Human-in-Loop, Doc Q&A, Portability); CALL overhead is negligible |
| Total cookbook wall time | 1,493.4s | 40/40 recipes; single GTX 1080 Ti, gemma3 via Ollama, sequential execution; 2026-03-23 |

### 6.3 Cost Efficiency

The declarative approach enables token savings through two mechanisms:

1. **Implicit context management**: SPL automatically constructs minimal prompts from SELECT clause elements, avoiding the verbose system prompts typical of manual prompt engineering.
2. **Budget enforcement**: `WITH BUDGET n TOKENS` and `WITH OUTPUT BUDGET n TOKENS` clauses enable explicit token accounting, preventing runaway costs.

Future optimizer passes (GENERATE merging, model downgrade) will provide additional cost reduction.

### 6.4 Developer Experience

**Table 8: Barrier to Entry**

| Skill Required | SPL 2.0 | LangGraph | AutoGen | CrewAI |
|---|---|---|---|---|
| Python programming | No | Yes | Yes | Yes |
| API knowledge | No | Yes | Yes | Yes |
| SQL familiarity | Helpful | not applicable | not applicable | not applicable |
| CLI usage | Basic | not applicable | not applicable | not applicable |
| Prompt engineering | Minimal | Full | Full | Full |

SPL 2.0's zero-config CLI (`spl2 run file.spl --adapter ollama`) eliminates the setup friction that characterizes imperative frameworks: no virtual environments to configure, no API clients to instantiate, no state management classes to define.

### 6.5 Case Study: Self-Refining Agent

The self-refining agent — a workflow that **drafts**, **critiques**, and **revises** its own output — is the canonical benchmark for agentic frameworks. It exercises every layer of SPL syntax: atomic LLM

queries (PROMPT/SELECT/GENERATE), multi-step orchestration (WORKFLOW), iterative refinement (WHILE), semantic quality gating (EVALUATE), and graceful degradation (EXCEPTION).

The complete SPL 2.0 implementation in 22 lines illustrates all seven core grammar constructs:

```
-- Illustrates: WORKFLOW, PROMPT, SELECT, GENERATE, WHILE,
--              EVALUATE, EXCEPTION
WORKFLOW self_refine
  INPUT:  @task text
  OUTPUT: @result text
DO
  -- Step 1: initial draft  (GENERATE)
  GENERATE draft(@task) INTO @current

  -- Step 2: iterative critique-and-revise loop  (WHILE + EVALUATE)
  WHILE @iteration < @max_iterations DO
    GENERATE critique(@current) INTO @feedback

    EVALUATE @feedback
      WHEN 'satisfactory' THEN
        COMMIT @current
      OTHERWISE
        GENERATE revised(@current, @feedback) INTO @current
    END
  END

  COMMIT @current

-- Step 3: handle runaway loops  (EXCEPTION)
EXCEPTION
  WHEN MaxIterationsReached THEN
    COMMIT @current WITH status = 'partial'
END
```

**Syntax grammar illustrated:**

- **PROMPT** — the named atomic LLM query unit (used inside WORKFLOW steps as GENERATE targets)
- **SELECT** — assembles context going *into* the model (implicit in each GENERATE call above)
- **GENERATE** — invokes the LLM and captures output into a variable (`INTO @var`)
- **WORKFLOW** — binds multiple GENERATE steps with shared variables and control flow
- **WHILE** — iterates until a counter or semantic condition terminates the loop
- **EVALUATE** — LLM-as-judge branch: the LLM evaluates `@feedback` and routes accordingly
- **EXCEPTION** — handles LLM-specific failure modes with structured recovery

For a detailed LOC comparison of this pattern implemented in SPL, LangGraph, AutoGen, and CrewAI, see Appendix H.

# 7. Discussion

## 7.1 Strengths

**Declarative form enables optimization.** Because SPL 2.0 specifies *what* should happen rather than *how*, the compiler can analyze data dependencies, identify parallelism opportunities, and select cost-optimal models — optimizations that are not structurally available in imperative code where the execution order is hard-coded by the programmer, leaving no intent visible to the system.

**Adapter pattern broadens reach.** The same `.spl` file runs on local GPU (Ollama), cloud APIs (OpenRouter, Claude), and decentralized compute (Momagrid). No other system spans all three deployment models from a single specification.

**SQL familiarity lowers the barrier.** SQL is the most widely known programming language in the world. SPL 2.0's syntax leverages this familiarity, making agentic workflows accessible to data analysts, researchers, and domain experts who are fluent in SQL but not in Python framework APIs.

**LLM-aware primitives address real failure modes.** The EXCEPTION taxonomy (HallucinationDetected, RefusalToAnswer, BudgetExceeded, etc.) codifies failure patterns that every LLM application must handle but that no existing language formalizes.

**Human reviewability enables non-technical governance.** With imperative frameworks, reviewing an agentic workflow requires tracing Python control flow, understanding framework-specific abstractions, and mentally simulating execution — skills that exclude the domain experts, analysts, and business stakeholders who are best positioned to validate whether the workflow is *correct for the problem*. An SPL script, by contrast, reads like a structured specification: SELECT declares what context flows in, GENERATE names the task, EVALUATE makes the branching condition explicit, and EXCEPTION names the failure modes. A business owner can read a 22-line SPL workflow and approve it without running it. This is a qualitatively different governance model — one that imperative code cannot provide regardless of documentation quality.

**Declarative operations simplify the deployment lifecycle.** SPL achieves not only declarative *implementation* but also declarative *operation*. The CLI flag surface — `--adapter` (LLM provider), `--model` (specific model), `--tools` (external connectors), `--datasets` (data sources) — acts as the environment configuration layer, analogous to environment variables in the 12-factor app methodology. The identical `.spl` file promotes unchanged through Dev, Test, and Production:

```
# Dev: local model, mock tools
spl run pipeline.spl --adapter ollama --model gemma3 \
    --tools search=mock --datasets corpus=sample.jsonl

# Test: cloud model, staging data
spl run pipeline.spl --adapter openrouter --model gpt-4o \
    --tools search=brave --datasets corpus=staging.jsonl

# Production: cost-optimised model, live data
spl run pipeline.spl --adapter anthropic --model claude-sonnet-4-6 \
    --tools search=brave --datasets corpus=prod
```

No code changes, no re-review, no merge. The workflow logic is environment-agnostic by construction; only the runtime configuration changes. Imperative frameworks require either separate code

paths per environment or fragile runtime injection — neither of which matches the simplicity of a flag swap.

### 7.2 Limitations

We acknowledge several limitations of the current work:

**No optimizer yet.** The optimization opportunities described in Section 5.3 are architectural possibilities enabled by the declarative form, not implemented features. The current system executes GENERATE statements sequentially without parallelism or model-routing optimizations.

**text2SPL (basic prototype).** The natural language to SPL compiler (Layer 2 of the stack) has a working prototype via `spl2 text2spl`, limited to simple PROMPT pattern generation. Full WORKFLOW generation from natural language — the complete vision for Layer 2 — remains future work.

**Text-only modality.** SPL 2.0 currently handles text input and output. Multi-modal support (images, audio, video) is planned for a future version.

**No streaming (deliberate for batch targets).** Adapters do not yet support token-by-token streaming of LLM output; results are returned in full at completion. For Momagrid's primary use case — batch inference on community-contributed GPU nodes — streaming adds protocol complexity without clear benefit, as batch jobs do not require real-time token delivery. Streaming support for latency-sensitive local adapters is planned for a future release.

**Single-model judge in EVALUATE.** By default, the same model that performs generation also acts as the quality judge in semantic EVALUATE conditions. If the generator is a weak model, the judge may be equally unreliable. A per-EVALUATE `USING MODEL` override mitigates this, but consistency of LLM-as-judge across repeated evaluations has not been formally measured.

**WHILE termination not formally guaranteed.** The WHILE loop relies on semantic conditions or iteration guards for termination. Non-termination is possible if the quality condition is never met and no explicit iteration cap is set. The runtime default of 10 iterations provides a practical safety net, but formal termination proofs are not provided.

**Single-node baseline.** The cookbook experiments reported here were run on a single GTX 1080 Ti GPU via Ollama with gemma3. Multi-node Momagrid execution and cross-model comparative benchmarks are planned for a follow-up evaluation.

**memory.get() is read-only in the current implementation.** The `memory.get('key')` context source is parsed, evaluated, and injected into the SELECT clause. However, write-back from within a WORKFLOW — persisting a computed value to the memory store for retrieval in a subsequent run — is not yet implemented. The `STORE` statement is defined in the grammar and accepted by the parser but is not executed by the runtime. Cross-run memory persistence is planned for a future release.

**rag.query() requires a pre-built index.** The `rag.query('text', top_k=N)` context source is functional only when a vector index has been pre-populated via `spl2 rag add <document>`. The 34-recipe cookbook benchmark was run without a pre-built index; recipe 08 (RAG Query) confirmed the execution path and correct syntax handling, but the LLM received empty retrieval context because no documents were indexed. End-to-end retrieval validation over populated indices is deferred to a follow-up evaluation.

### 7.3 Comparison with Imperative Frameworks

The contrast between imperative and declarative approaches is most visible in the ReAct agent pattern. A LangGraph implementation requires approximately 100+ lines of Python: defining state classes, creating graph nodes, specifying edge conditions, configuring the state machine, and handling errors through try/except blocks. The equivalent SPL 2.0 implementation is 43 lines with the workflow logic, branching conditions, and error handling all expressed declaratively.

More importantly, the declarative form makes the *intent* immediately visible. A reader of the SPL code can see at a glance: "this agent reasons, decides on an action, executes it, and accumulates history — with handlers for hallucination and max iterations." The same understanding from the LangGraph implementation requires tracing through multiple function definitions, graph edge conditions, and state class hierarchies.

### 7.4 Design Intent: The PL/SQL Transfer Path

A central motivation behind SPL 2.0's syntax is not merely aesthetic coherence with SQL — it is a deliberate strategy for workforce transfer. Relational databases have produced, over five decades, a global community of practitioners who think fluently in the stored-procedure mental model: declare inputs and outputs, assemble data with queries, branch on conditions, loop until a criterion is met, and handle named exceptions. These engineers built the data pipelines that power the world's enterprises. Their instincts — separate *what* from *how*, compose transformations, handle failure explicitly — are precisely the instincts an AI workflow engineer needs.

SPL 2.0's **WORKFLOW** is a stored procedure for LLM orchestration. The structural parallel is exact, not approximate:

| **Oracle PL/SQL** | **SPL 2.0** |
|---|---|
| `CREATE PROCEDURE p(x IN VARCHAR2)` | `WORKFLOW p INPUT @x TEXT` |
| `v_result VARCHAR2(200);` | `@result TEXT` |
| `SELECT ... INTO v_result FROM ...;` | `GENERATE f(@x) INTO @result;` |
| `WHILE condition LOOP ... END LOOP;` | `WHILE condition DO ... END;` |
| `IF cond THEN ... ELSIF ... END IF;` | `EVALUATE WHEN ... THEN ... END;` |
| `EXCEPTION WHEN DUP_VAL_ON_INDEX THEN` | `EXCEPTION WHEN HallucinationDetected THEN` |
| `COMMIT;` | `COMMIT @result;` |

A data engineer who has written Oracle PL/SQL, PostgreSQL PL/pgSQL, or SQL Server T-SQL stored procedures can read and write SPL 2.0 WORKFLOW blocks with no tutorial. The block structure, exception taxonomy, variable lifecycle, and loop semantics are already in their muscle memory. What is new is only the backend: instead of `SELECT ... FROM orders` querying a relational table, `GENERATE answer(question)` queries a large language model. The declarative SELECT that assembles context mirrors the SELECT that assembles data — the difference is that the "rows" are prompt fragments, memories, and retrieved documents rather than database records.

This is the core accessibility claim of SPL 2.0: the millions of data engineers who currently maintain stored procedures, ETL pipelines, and reporting workflows are not learning a new paradigm to become AI engineers. They are applying a paradigm they have used for years to a new backend. The conceptual distance from PL/SQL to SPL 2.0 is a single substitution — *language model* for *relational engine* — while the surrounding architecture, discipline, and vocabulary remain unchanged.

The practical implication is day-one productivity. In the experiments presented in Section 6, the

cookbook of 40 recipes spans patterns from simple single-step inference (Recipe 01) through multi-agent debate (Recipe 11), MapReduce over documents (Recipe 13), and CALL-based tool use with Python UDFs (Recipe 36). Every one of these recipes could be written, read, and maintained by a senior PL/SQL developer encountering SPL 2.0 for the first time. No Python runtime knowledge, no framework-specific graph API, no imperative state machine design is required. The expressiveness of the language follows directly from this design choice: a language built on a fifty-year-old abstraction that has already proven it can express arbitrary data transformations at scales. Now that same mental power can be applied to agentic workflow orchestration over LLM.

### 7.5 Future Work

**text2SPL compiler.** Natural language to SPL 2.0 translation would complete the four-layer stack, enabling zero-syntax access to agentic workflows.

**Query optimizer.** A cost-based optimizer that rewrites execution plans — merging GENERATE calls, parallelizing independent steps, routing to optimal models — is the highest-impact future contribution.

**Visual workflow builder.** A graphical interface for composing WORKFLOW blocks would complement the text-based CLI for users who prefer visual programming.

**Multi-modal I/O.** Extending GENERATE to handle images, audio, and video inputs/outputs, with corresponding context source types.

**Extended procedure library.** A standard library of validated PROCEDURE implementations for common patterns (ReAct, Tree-of-Thought, Self-Refine, Debate, Map-Reduce, etc.) would accelerate adoption.

**Formal verification.** Proving properties of workflow specifications — termination guarantees, budget bounds, safety constraints — through static analysis.

**Momagrid production integration.** End-to-end testing of the decentralized runtime with real GPU node operators.

## 8. Conclusion

SPL 2.0 represents a paradigm shift from imperative to declarative orchestration of LLM-powered agentic workflows. The language addresses a fundamental fragmentation in the current ecosystem — where building agents requires Python programming, API expertise, and framework-specific knowledge — by providing a SQL-like specification layer that separates *what* a workflow should accomplish from *how* it executes.

The language is built on six core primitives. Three are inherited from SPL 1.0 [5] (the declarative context management foundation) and remain fully compatible:

- **PROMPT** — the named, executable unit binding a SELECT and a GENERATE into a single atomic LLM query
- **SELECT** prepares and assembles input context for the LLM
- **GENERATE** invokes the LLM and captures its output

Three are novel SPL 2.0 primitives that no existing language formalizes:

- **EVALUATE** enables hybrid deterministic-semantic branching, where the LLM itself serves as a runtime judge.
- **WHILE** with semantic termination enables iteration based on output quality rather than arbitrary counters.
- **EXCEPTION** with a formal taxonomy of LLM failure modes enables structured error handling for hallucination, refusal, context overflow, and budget exhaustion.

Readers seeking the complete picture of the SPL 1.0 foundation (PROMPT, SELECT, GENERATE) are directed to the SPL 1.0 paper [5].

This paper (v0.7) adds three further contributions beyond the SPL 2.0 language design: **(1) the CALL/GENERATE design principle** — an explicit language-level boundary between deterministic and probabilistic operations, with the `@spl_tool` decorator as a Foreign Function Interface to the Python ecosystem; **(2) the SPL Standard Library** — 55 SQL-equivalent built-in functions auto-loaded in every workflow; and **(3) the `spl-cookbook` package and Recipe Kitchen UI** — a pip-installable distribution with a Streamlit browser interface for all 40 recipes.

These primitives and extensions, combined with the PROMPT/WORKFLOW layering (modeled after the SQL/PL/SQL architectural pattern), a formal EBNF grammar, and a pluggable adapter layer spanning local, cloud, and decentralized execution, form a coherent language design validated through 231 automated tests and a cookbook of 40 recipes covering 8 categories of agentic workflow patterns, with 40/40 recipes passing on a single local GPU node.

The declarative form of SPL 2.0 is not merely a notational convenience — it is the structural prerequisite for a future ecosystem of compilers, optimizers, and visual tools that can reason about, transform, and improve agentic workflows in ways that imperative code fundamentally cannot support.

Our reference implementation in python is open-source, released under Apache 2.0 license, available at GitHub (https://github.com/digital-duck/SPL20).

## Acknowledgements

The author used Claude (Anthropic), Gemini (Google), and Z.ai (ZhiPu AI) as AI assistants during the preparation of this work: for literature search and related-work synthesis, for implementation and debugging of the reference codebase, and for editing portions of the manuscript. All technical claims, experimental results, and final text were reviewed and approved by the author, who takes full responsibility for the content.

## References

[1] LangGraph. LangChain, Inc. https://github.com/langchain-ai/langgraph

[2] Wu, Q., et al. "AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation." arXiv:2308.08155, 2023.

[3] CrewAI. https://github.com/crewAIInc/crewAI

[4] Chamberlin, D. D. and Boyce, R. F. "SEQUEL: A Structured English Query Language." SIGFIDET Workshop, 1974.

[5] Gong, Wen G. "Structured Prompt Language: Declarative Context Management for LLMs." arXiv:2602.21257, 2026.

[6] Oracle Corporation. "PL/SQL User's Guide and Reference." Oracle Database Documentation.

[7] Beurer-Kellner, L., Fischer, M., and Vechev, M. "Prompting Is Programming: A Query Language for Large Language Models." arXiv:2212.06094, 2022.

[8] Khattab, O., et al. "DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines." arXiv:2310.03714, 2023.

[9] Zheng, L., et al. "SGLang: Efficient Execution of Structured Language Model Programs." arXiv:2312.07104, 2023.

[10] Agent Spec: Open Agent Specification. Oracle Labs. arXiv:2510.04173, 2025.

[11] SLANG: Standard Language for Agentic Workflows. https://dev.to/argonauta/why-we-need-a-standard-language-for-agentic-workflows-and-why-i-built-one-138

[12] "A Declarative Language for Building And Orchestrating LLM-Powered Agent Workflows." arXiv:2512.19769, 2025.

[13] "The Auton Agentic AI Framework: A Declarative Architecture for Specification, Governance, and Runtime Execution of Autonomous Agent Systems." arXiv:2602.23720, 2026.

[14] Gliozzo, A. M., Lee, J., and Defosse, N. "Agentics 2.0: Logical Transduction Algebra for Agentic Data Workflows." arXiv:2603.04241, 2026.

[15] Gong, Wen G. "Momagrid: A Decentralized Inference Runtime with Linear Complexity via Semantic Chunking." arXiv:submit/7373026, 2026.

[16] Yang, J., et al. "Auto-GPT: An Autonomous GPT-4 Experiment." arXiv:2306.02224, 2023.

---

## Appendix A: Complete SPL 2.0 Grammar (EBNF)

Productions are listed top-down. SPL 1.0 productions (unchanged) are marked †. New SPL 2.0 productions are unmarked. Keywords are case-insensitive. Comments use -- (end-of-line).

```
(* Top-level *)
program            := statement+

statement          := prompt_stmt
                    | workflow_stmt
                    | procedure_stmt
                    | create_function_stmt
                    | explain_stmt †
                    | execute_stmt †

(* PROMPT — atomic LLM query †  extended with CTE block *)
prompt_stmt        := 'PROMPT' identifier
                       cte_block?
                       select_clause
                       generate_clause
                       store_clause?

cte_block          := 'WITH' cte_def (',' cte_def)*
cte_def            := identifier 'AS' '(' prompt_stmt ')'

select_clause      := 'SELECT' select_item (',' select_item)*
select_item        := expression ('AS' identifier)?
                       ('LIMIT' integer 'TOKENS')?

generate_clause    := 'GENERATE' identifier '(' argument_list? ')'
                       ('INTO' '@' identifier)?
                       ('WITH' generate_option (',' generate_option)*)?
                       ('USING' 'MODEL' string)?

generate_option    := 'OUTPUT' 'BUDGET' integer 'TOKENS'
                    | 'TEMPERATURE' float

store_clause       := 'STORE' 'RESULT' 'AS' string
                       ('CACHE' 'FOR' integer)? †

(* WORKFLOW — procedural orchestration *)
workflow_stmt      := 'WORKFLOW' identifier
                       ('INPUT'  ':' param_list)?
                       ('OUTPUT' ':' param_list)?
                       security_block?
                       accounting_block?
                       do_block
                       'END'

(* PROCEDURE — reusable parameterized sub-workflow *)
procedure_stmt     := 'PROCEDURE' identifier
                       '(' param_list? ')'
                       ('RETURNS' identifier)?
                       security_block?
```

```
                        accounting_block?
                        do_block
                        'END'

(* DO block — shared by WORKFLOW and PROCEDURE *)
do_block           := 'DO' body_stmt+ exception_block?

body_stmt          := generate_into_stmt
                    | assignment_stmt
                    | evaluate_stmt
                    | while_stmt
                    | commit_stmt
                    | retry_stmt
                    | raise_stmt
                    | call_stmt
                    | select_into_stmt

generate_into_stmt := 'GENERATE' identifier '(' argument_list? ')'
                        'INTO' '@' identifier
                        ('WITH' generate_option (',' generate_option)*)?

assignment_stmt    := '@' identifier ':=' expression

evaluate_stmt      := 'EVALUATE' expression
                        when_clause+
                        ('OTHERWISE' body_stmt+)?
                        'END'

when_clause        := 'WHEN' condition 'THEN' body_stmt+

condition          := comparison_condition      (* deterministic *)
                    | semantic_condition         (* LLM-judged    *)
                    | expression

comparison_condition := comparison_op expression
comparison_op      := '>' | '<' | '>=' | '<=' | '=' | '!='

semantic_condition := string_literal            (* quoted string triggers LLM judge *)

while_stmt         := 'WHILE' condition 'DO' body_stmt+ 'END'

commit_stmt        := 'COMMIT' expression ('WITH' option_list)?

retry_stmt         := 'RETRY' ('WITH' option_list)?

raise_stmt         := 'RAISE' exception_type string?

call_stmt          := 'CALL' identifier '(' named_arg_list? ')'
                        ('INTO' '@' identifier)?

select_into_stmt   := 'SELECT' select_item (',' select_item)*
                        'INTO' '@' identifier

exception_block    := 'EXCEPTION' exception_handler+
```

```
exception_handler  := 'WHEN' exception_type 'THEN' body_stmt+

exception_type     := 'HallucinationDetected'  | 'RefusalToAnswer'
                    | 'ContextLengthExceeded'   | 'ModelOverloaded'
                    | 'BudgetExceeded'          | 'QualityBelowThreshold'
                    | 'MaxIterationsReached'    | 'NodeUnavailable'
                    | 'GenerationError'         | 'GenerationFailed'

(* CREATE FUNCTION — schema / constant definition *)
create_function_stmt := 'CREATE' 'FUNCTION' identifier
                        '(' param_list? ')'
                        'RETURNS' identifier
                        'AS' '$$' text '$$'

(* Expressions *)
expression         := literal
                    | param_ref
                    | dotted_name
                    | function_call
                    | binary_op
                    | select_source

select_source      := system_role | context_ref | rag_query | memory_get

system_role        := 'system_role' '(' string ')'
context_ref        := 'context' '.' identifier
rag_query          := 'rag' '.' 'query' '(' string (',' named_arg)* ')'
memory_get         := 'memory' '.' 'get' '(' string ')'

literal            := string_literal | integer_literal | float_literal
param_ref          := '@' identifier
dotted_name        := identifier ('.' identifier)+
function_call      := identifier '(' argument_list? ')'
binary_op          := expression ('+' | '-' | '*' | '/') expression

(* Parameter and argument lists *)
param_list         := param_def (',' param_def)*
param_def          := '@'? identifier identifier? ('DEFAULT' expression)?

argument_list      := argument (',' argument)*
argument           := named_arg | expression
named_arg          := identifier '=' expression

named_arg_list     := named_arg (',' named_arg)*

option_list        := option (',' option)*
option             := identifier '=' expression

(* Metadata blocks *)
security_block     := 'SECURITY' ':' 'CLASSIFICATION' ':' identifier
accounting_block   := 'ACCOUNTING' ':' 'BILLABLE_TO' ':' string
                       ('BUDGET_LIMIT' ':' integer identifier)?
```

```
(* Lexical *)
identifier         := [a-zA-Z_][a-zA-Z0-9_]*
string_literal     := '"' .* '"' | "'" .* "'"
integer_literal    := [0-9]+
float_literal      := [0-9]+ '.' [0-9]+
```

**Type system.** SPL 2.0 supports five primitive types: `TEXT`, `NUMBER` (integer or float), `BOOLEAN`, `JSON`, `CHUNK` (token-bounded text segment). Composite types: `LIST` (ordered sequence), `MAP` (key-value). Type annotations are optional; untyped parameters default to `TEXT`.

**Condition type detection rule.** If a `WHEN` condition is a quoted string literal, the runtime treats it as a semantic condition and invokes the LLM as a judge. Otherwise (numeric comparison, boolean expression) it is evaluated deterministically with no LLM call. This detection is unambiguous and requires no explicit annotation.

## Appendix B: IR Specification (JSON Schema)

The SPL 2.0 Intermediate Representation (IR) is a JSON-serialized AST produced by `spl2 explain`. It is the portable format consumed by the Momagrid runtime, plan visualization tools, and future optimizer passes.

### Program envelope

```
{
  "type": "Program",
  "statements": [ <statement>+ ]
}
```

**Prompt plan** (produced by `spl2 explain` for PROMPT statements)

```
{
  "type": "PromptPlan",
  "prompt_name": "string",
  "model": "string | null",
  "total_budget": "integer | null",
  "output_budget": "integer | null",
  "total_input_tokens": "integer",
  "buffer_tokens": "integer",
  "estimated_cost": "float",
  "optimizations": [ "string" ],
  "warnings": [ "string" ],
  "steps": [
    {
      "operation": "string",
      "source": "string",
      "alias": "string | null",
      "estimated_tokens": "integer",
      "limit_tokens": "integer | null",
      "allocated_tokens": "integer",
      "compressed": "boolean",
      "compression_ratio": "float",
      "cache_status": "string",
      "priority": "integer"
```

```
    }
  ]
}
```

**Workflow plan** (produced by `spl2 explain` for WORKFLOW statements)

```
{
  "type": "WorkflowPlan",
  "workflow_name": "string",
  "inputs": [ { "name": "string", "type": "string", "default": "<expr> | null" } ],
  "outputs": [ { "name": "string", "type": "string" } ],
  "security": { "classification": "string" } | null,
  "accounting": { "billable_to": "string" } | null,
  "total_estimated_llm_calls": "integer",
  "total_estimated_tokens": "integer",
  "steps": [ <workflow_step> ],
  "exception_handlers": [ <workflow_step> ]
}
```

**Workflow step** (recursive — branches and substeps are the same type)

```
{
  "step_type": "GenerateInto | Evaluate | While | Commit | Retry | Call | Assignment | ...",
  "description": "string",
  "estimated_llm_calls": "integer",
  "estimated_tokens": "integer",
  "substeps": [ <workflow_step> ] | null,
  "branches": [
    {
      "condition": "string",
      "condition_type": "semantic | deterministic | hybrid",
      "steps": [ <workflow_step> ]
    }
  ] | null,
  "exception_handlers": [ <workflow_step> ] | null
}
```

**Condition node**

```
{ "type": "SemanticCondition",   "value": "string" }
{ "type": "ComparisonCondition", "operator": "> | < | >= | <= | = | !=", "right": <expr> }
{ "type": "Condition",           "left": <expr>, "operator": "string", "right": <expr> }
```

**Expression nodes** (used in step arguments, conditions, assignments)

```
{ "type": "Literal",    "value": "string | integer | float" }
{ "type": "ParamRef",   "name": "string" }
{ "type": "ContextRef", "field_name": "string" }
{ "type": "DottedName", "parts": ["string"], "full_name": "string" }
{ "type": "SystemRole", "description": "string" }
{ "type": "RagQuery",   "query": "string", "source": "string | null", "top_k": "integer" }
{ "type": "MemoryGet",  "key": "string" }
{ "type": "FunctionCall","name": "string", "arguments": [ <expr> ] }
{ "type": "BinaryOp",   "op": "+ | - | * | /", "left": <expr>, "right": <expr> }
```

The IR is produced by the `spl2.ir` module and consumed by the executor (`spl2.executor`) and Momagrid adapter (`spl2.adapters.momagrid`). All field names are stable across SPL 2.0 patch releases; additions are backward-compatible.

## Appendix C: Hub ↔ Node Protocol Specification

The Momagrid Hub-Node protocol is a REST-over-HTTPS request-response protocol. The SPL 2.0 Momagrid adapter (`spl2.adapters.momagrid`) implements the client (Hub) side using async `httpx`.

### Task lifecycle

```
Client (SPL runtime)          Hub                      GPU Node||||—

         POST /tasks ———————|> |||—
                              dispatch ———————————————|>|-
        < 202 { task_id } ——||||
                                                 (inference)|—
         GET /tasks/{id} —|> ||-
        < 200 { pending } ——||||—
                             < result ———————————————||—
         GET /tasks/{id} —|> ||-
        < 200 { complete } —||
```

### POST /tasks — Submit inference request

```
Request body:
{
  "model":         "string",
  "prompt":        "string (plaintext or encrypted ciphertext)",
  "system_role":   "string | null",
  "temperature":   "float | null",
  "max_tokens":    "integer | null",
  "node_requirements": {
    "min_tier":    "integer –(05)",
    "min_vram_gb": "float"
  },
  "session_key_id": "string | null"
}

Response 202:
{
  "task_id": "uuid",
  "status":  "pending",
  "eta_ms":  "integer | null"
}
```

### GET /tasks/{task_id} — Poll task status

```
Response 200 (pending):
{ "task_id": "uuid", "status": "pending | running" }

Response 200 (complete):
{
  "task_id":    "uuid",
```

```
  "status":     "complete",
  "result":     "string (plaintext or encrypted ciphertext)",
  "tokens_in":  "integer",
  "tokens_out": "integer",
  "latency_ms": "integer",
  "node_id":    "string"
}

Response 200 (failed):
{
  "task_id": "uuid",
  "status":  "failed",
  "error":   "string",
  "code":    "node_unavailable | model_not_found | overloaded | timeout"
}
```

### Error code → SPL exception mapping

| HTTP / error code | SPL Exception raised |
|---|---|
| `overloaded` or HTTP 429 | `ModelOverloaded` |
| `node_unavailable` or HTTP 503 | `NodeUnavailable` |
| `timeout` | `NodeUnavailable` |
| `context_exceeded` | `ContextLengthExceeded` |
| any other `failed` status | `GenerationError` |

### Encryption handshake (zero-trust mode)

When `SECURITY: CLASSIFICATION: confidential` is set in the SPL workflow, the adapter negotiates an ephemeral session key before task submission:

1. Client generates an ephemeral key pair (X25519).
2. Client sends the public key to `POST /session/init`.
3. Hub returns an encrypted session key, signed by the Hub's master key.
4. Subsequent task payloads (`prompt`, `result`) are encrypted with AES-256-GCM under the session key.
5. Session keys expire after 1 hour or 1000 tasks, whichever comes first.

In plaintext mode (default for development), the `session_key_id` field is omitted and `prompt`/`result` are transmitted unencrypted.

### Accounting integration

When `ACCOUNTING: BILLABLE_TO: 'project-id'` is set, the Hub records a Moma Points debit against the project account on task completion. If `BUDGET_LIMIT` is exceeded, the Hub returns HTTP 402 and the adapter raises `BudgetExceeded`.

## Appendix D: Exception Type Hierarchy

| Exception | Trigger | Default Recovery |
|---|---|---|
| HallucinationDetected | Output fails fact-checking or contains fabricated information | Regenerate with temperature = 0.1 |
| RefusalToAnswer | Model declines due to safety filters or content policy | Log refusal, COMMIT with status = 'refused' |
| ContextLengthExceeded | Combined input exceeds model's context window | Compress context, RETRY |
| ModelOverloaded | Provider returns 429 or capacity error | Exponential backoff, RETRY |
| BudgetExceeded | Token budget (input or output) exhausted | COMMIT best available result |
| QualityBelowThreshold | Generated output scores below configured minimum | Regenerate or escalate to larger model |
| MaxIterationsReached | WHILE loop reached iteration guard limit | COMMIT current result with status = 'partial' |
| NodeUnavailable | Momagrid GPU node unreachable or timed out *(Momagrid-specific)* | Route to alternative node, RETRY |

## Appendix E: Momagrid Runtime Architecture

SPL 2.0's adapter pattern supports a unique deployment model: decentralized execution on independently operated GPU nodes. Momagrid targets a fundamentally different niche from commercial cloud APIs: **batch inference for users who prioritize accessibility and cost over latency.** The design goal is low- or no-cost agentic AI inference for researchers, students, and developers in resource-constrained environments — complementary to, not competitive with, services like Anthropic or Google.

**Current status:** Research prototype. LAN grid testing is complete. WAN deployment simulation is planned. A research grant proposal has been submitted to Mozilla Foundation. Latency variability is an acknowledged trade-off for the accessibility target.

The final arXiv submission will include a multi-configuration LAN scaling experiment running the full cookbook under three grid sizes on a home LAN lab:

| Configuration | Nodes | Expected Result |
|---|---|---|
| 1-GPU | 1 node (GTX 1080 Ti) | Baseline (matches Appendix G) |
| 2-GPU | 2 nodes | Throughput scaling, distributed overhead |
| 3-GPU | 3 nodes | Full LAN grid, throughput ceiling |

Results (total wall time, per-recipe latency, scaling factor vs. single-node) will be reported in Appendices F and G upon completion.

| Adapter | Target Use Case | Latency Profile |
|---|---|---|
| Ollama | Local dev, single GPU | Low (local) |
| OpenRouter | Cloud, 100+ models | Medium (API) |
| Claude CLI | Anthropic direct | Medium (API) |
| **Momagrid** | **Batch jobs, low-cost inference** | **Variable (decentralized)** |

### E.1 The Hub

The Momagrid Hub serves as the coordination layer for decentralized workflows:

- **Task submission**: Receives SPL-generated inference requests with model and hardware requirements.
- **Node selection**: Matches tasks to available GPU nodes based on min_tier, min_vram_gb, and geographic constraints.
- **State management**: Maintains global workflow context, enabling checkpointing and recovery for long-running workflows.

### E.2 GPU Nodes

Each GPU node in the Momagrid network operates as a stateless executor:

- Receives encrypted task payloads from the Hub.
- Decrypts, performs inference, and encrypts the result.
- Reports completion status and resource consumption metrics.

Node selection criteria (min_tier, min_vram_gb) are specified in SPL 2.0's SECURITY metadata block, allowing workflows to express hardware requirements declaratively.

### E.3 Compression, Encryption, Accounting

Three cross-cutting concerns are built into the Momagrid protocol:

**Compression**: Two levels — *semantic compression* reduces context length before transmission (preserving meaning while reducing tokens), and *transmission compression* (zstd, gzip) reduces network bandwidth.

**Encryption**: A zero-trust model with a three-level key hierarchy (Master → Session → Data) ensures that no single party — including the Hub — has access to both the prompt and the response in plaintext.

**Accounting**: The Moma Points system tracks resource consumption, with event logging and a reputation system for node operators. SPL 2.0's ACCOUNTING metadata block integrates with this system: `BUDGET_LIMIT: 1000 moma_points` enforces cost caps at the workflow level.

## Appendix F: Cookbook — Practical Validation

### F.1 Recipe Design Philosophy

Each cookbook recipe is designed as a self-contained demonstration: - A single `.spl` file with inline comments explaining the pattern. - Usage examples as CLI commands in the file header. - Progressive complexity from atomic queries to multi-agent workflows. - Catalog-driven organization via `cookbook_catalog.json` and `run_all.py` batch runner.

### F.2 Recipe Catalog

**Table E-1: SPL 2.0 Cookbook — 40 Recipes**

| # | Recipe | Category | Pattern | LOC | Key SPL Features |
|---|---|---|---|---|---|
| 01 | Hello World | basics | PROMPT (atomic) | 4 | system_role, GENERATE |
| 02 | Ollama Proxy | basics | PROMPT + params | 5 | context.param, model override |
| 03 | Multilingual | basics | PROMPT + 2 params | 6 | Multiple context bindings |
| 04 | Model Showdown | benchmarking | WORKFLOW + multi-model | 30 | Cross-model comparison, 4 calls |
| 05 | Self-Refine | agentic | WORKFLOW + WHILE | 35 | Iterative improvement loop |
| 06 | ReAct Agent | agentic | WORKFLOW + EVALUATE | 43 | Tool-call branching, history |
| 07 | Safe Generation | safety | WORKFLOW + EXCEPTION | 37 | Quality scoring, error recovery |
| 08 | RAG Query | retrieval | PROMPT + rag.query | 7 | Vector search context |
| 09 | Chain of Thought | reasoning | WORKFLOW (3-step) | 14 | Sequential GENERATE pipeline |
| 10 | Batch Test | benchmarking | WORKFLOW + parallel CTE | 50 | Multi-recipe, multi-model regression |
| 11 | Debate Arena | multi-agent | WORKFLOW + adversarial | 45 | Multi-perspective, judge eval |
| 12 | Plan and Execute | agentic | WORKFLOW + re-planning | 50 | Dynamic plan adjustment |
| 13 | Map-Reduce Summarizer | application | WORKFLOW + chunking | 40 | Chunked GENERATE, reduce step |
| 14 | Multi-Agent Collab. | multi-agent | PROCEDURE + CALL | 45 | Role specialization, delegation |
| 15 | Code Review | application | WORKFLOW + multi-pass | 50 | Severity scoring, synthesis |
| 16 | Reflection Agent | agentic | WORKFLOW + meta-cognition | 45 | Confidence scoring, correction |
| 17 | Tree of Thought | reasoning | WORKFLOW + branching | 50 | Path scoring, verification |
| 18 | Guardrails Pipeline | safety | WORKFLOW + multi-gate | 55 | PII detection, content filtering |
| 19 | Memory Conversation | application | WORKFLOW + memory | 40 | memory.get, fact extraction |
| 20 | Ensemble Voting | multi-agent | WORKFLOW + consensus | 45 | Multi-candidate scoring, voting |
| 21 | Multi-Model Pipeline | agentic | WORKFLOW + adapters | 50 | Cross-adapter routing, fallback |
| 22 | Text2SPL Demo | basics | *(NEW — planned, not yet released)* | — | Natural language → SPL compilation |
| 23 | Structured Output | basics | PROMPT + CREATE FUNCTION | 36 | JSON schema-constrained generation |

| # | Recipe | Category | Pattern | LOC | Key SPL Features |
|---|---|---|---|---|---|
| 24 | Few-Shot Prompting | basics | PROMPT + CREATE FUNCTION | 47 | In-context learning, domain examples |
| 25 | Nested Procedures | agentic | WORKFLOW + PROCEDURE + CALL | 55 | PROCEDURE composability, delegation |
| 26 | Prompt A/B Test | benchmarking | WORKFLOW + comparison | 40 | Multi-prompt comparison, scoring |
| 27 | Data Extraction | application | PROMPT + CREATE FUNCTION | 32 | Schema-constrained field extraction |
| 28 | Customer Support Triage | application | WORKFLOW + EVALUATE + EXCEPTION | 60 | Classify → route → draft pipeline |
| 29 | Meeting Notes to Actions | application | WORKFLOW + EXCEPTION | 55 | Transcript → structured TODO list |
| 30 | Code Generator + Tests | application | WORKFLOW + EVALUATE + EXCEPTION | 65 | Generate → review → test pipeline |
| 31 | Sentiment Pipeline | application | WORKFLOW + multi-step | 50 | Multi-label sentiment + reporting |
| 32 | Socratic Tutor | application | WORKFLOW + iterative | 55 | Adaptive question-driven learning |
| 33 | Interview Simulator | application | WORKFLOW + PROCEDURE | 60 | Role-play, scoring, feedback |
| 34 | Progressive Summarizer | application | WORKFLOW + chunking | 45 | Hierarchical summarization |
| 35 | Hypothesis Tester | reasoning | WORKFLOW + EVALUATE | 55 | Hypothesis → test plan → evidence |
| 36 | Tool Use | agentic | WORKFLOW + CALL | 32 | CALL UDFs for math; GENERATE for narrative |
| 37 | Headline News | application | WORKFLOW + CALL | 28 | Fetch RSS, summarize, format report |
| 38 | Human in the Loop | agentic | WORKFLOW + CALL + EVALUATE | 38 | Human approval gate before COMMIT |
| 39 | Doc Q&A | retrieval | WORKFLOW + CALL | 30 | Load document, extract answers via GENERATE |
| 40 | Portability Test | benchmarking | WORKFLOW + stdlib | 25 | to_float from stdlib; runs on any adapter |

### F.3 End-to-End Demonstrations

**Experimental setup.** All 40 recipes were executed via the batch runner (`python cookbook/run_all.py --adapter ollama --model gemma3`) on a single workstation equipped with a GTX 1080 Ti GPU running the Ollama inference server. No cloud API credits were consumed. Recipes 36–40 are

new in v0.7 and include the Tool-Use (CALL UDFs), Human-in-Loop, Doc Q&A, and Portability Test (stdlib) patterns. The full v0.7 run completed in **1,493.4 seconds** — 40/40 recipes passing, including recipe 22 (Text2SPL Demo) which now runs successfully for the first time.

**Table E-2: Batch Run Results — gemma3 / Ollama / GTX 1080 Ti (2026-03-23, v0.7)**

| # | Recipe | Category | Type | Status | Tokens in/out | Latency | LLM Calls |
|---|---|---|---|---|---|---|---|
| 01 | Hello World | basics | PROMPT | **OK** | — | 5.3s | 1 |
| 02 | Ollama Proxy | basics | PROMPT | **OK** | — | 0.7s | 1 |
| 03 | Multilingual Greeting | basics | PROMPT | **OK** | — | 0.6s | 1 |
| 04 | Model Showdown | benchmarking | WORKFLOW | **OK** | — | 26.2s | — |
| 05 | Self-Refine | agentic | WORKFLOW | **OK** | — | 3.9s | — |
| 06 | ReAct Agent | agentic | WORKFLOW | **OK** | — | 27.3s | — |
| 07 | Safe Generation | safety | WORKFLOW | **OK** | — | 10.3s | — |
| 08 | RAG Query | retrieval | PROMPT | **OK** | — | 6.1s | 1 |
| 09 | Chain of Thought | reasoning | WORKFLOW | **OK** | — | 23.5s | — |
| 10 | Batch Test | benchmarking | WORKFLOW | **OK** | — | 5.1s | — |
| 11 | Debate Arena | multi-agent | WORKFLOW | **OK** | — | 91.8s | — |
| 12 | Plan and Execute | agentic | WORKFLOW | **OK** | — | 89.9s | — |
| 13 | Map-Reduce Summarizer | application | WORKFLOW | **OK** | — | 10.6s | — |
| 14 | Multi-Agent Collab. | multi-agent | WORKFLOW | **OK** | — | 76.0s | — |
| 15 | Code Review | application | WORKFLOW | **OK** | — | 47.5s | — |
| 16 | Reflection Agent | agentic | WORKFLOW | **OK** | — | 132.3s | — |
| 17 | Tree of Thought | reasoning | WORKFLOW | **OK** | — | 80.6s | — |
| 18 | Guardrails Pipeline | safety | WORKFLOW | **OK** | — | 8.2s | — |
| 19 | Memory Conversation | application | WORKFLOW | **OK** | — | 3.9s | — |
| 20 | Ensemble Voting | multi-agent | WORKFLOW | **OK** | — | 61.2s | — |
| 21 | Multi-Model Pipeline | agentic | WORKFLOW | **OK** | — | 32.5s | — |
| 22 | Text2SPL Demo | basics | WORKFLOW | **OK** | — | 13.8s | — |
| 23 | Structured Output | basics | PROMPT | **OK** | — | 1.4s | 1 |
| 24 | Few-Shot Prompting | basics | PROMPT | **OK** | — | 0.8s | 1 |
| 25 | Nested Procedures | agentic | WORKFLOW | **OK** | — | 32.6s | — |
| 26 | Prompt A/B Test | benchmarking | WORKFLOW | **OK** | — | 75.3s | — |
| 27 | Data Extraction | application | PROMPT | **OK** | — | 1.3s | 1 |
| 28 | Customer Support Triage | application | WORKFLOW | **OK** | — | 55.8s | — |
| 29 | Meeting Notes to Actions | application | WORKFLOW | **OK** | — | 36.2s | — |
| 30 | Code Generator + Tests | application | WORKFLOW | **OK** | — | 34.4s | — |
| 31 | Sentiment Pipeline | application | WORKFLOW | **OK** | — | 83.1s | — |
| 32 | Socratic Tutor | application | WORKFLOW | **OK** | — | 90.6s | — |
| 33 | Interview Simulator | application | WORKFLOW | **OK** | — | 126.7s | — |
| 34 | Progressive Summarizer | application | WORKFLOW | **OK** | — | 10.2s | — |
| 35 | Hypothesis Tester | reasoning | WORKFLOW | **OK** | — | 36.9s | — |
| 36 | Tool-Use / Function-Call | agentic | WORKFLOW | **OK** | — | 8.7s | — |
| 37 | Headline News Aggregator | application | WORKFLOW | **OK** | — | 131.6s | — |
| 38 | Human-in-the-Loop Appr. Gate | agentic | WORKFLOW | **OK** | — | 3.7s | — |
| 39 | Document Q&A | retrieval | WORKFLOW | **OK** | — | 1.7s | — |
| 40 | Portability Test | benchmarking | WORKFLOW | **OK** | — | 5.0s | — |

**Summary:** 40/40 recipes passed (100%). Total wall time: **1,493.4 seconds**. This is the first run in which recipe 22 (Text2SPL Demo) passes — previously always skipped as "planned/not yet released." Five new recipes (36–40) all pass, with the CALL-based recipes completing in 1.7s–8.7s, demonstrating that delegating computation to deterministic Python UDFs adds negligible

overhead. The fastest workflow is Recipe 39 (Doc Q&A, 1.7s); the slowest is Recipe 16 (Reflection Agent, 132.3s).

**Aggregate statistics:**

| Segment | Count | Avg Latency | Avg Tokens (in+out) | Avg LLM Calls |
|---|---|---|---|---|
| PROMPT recipes (7) | 7 | 2.3s | — | 1 |
| WORKFLOW recipes (33) | 33 | 44.8s | — | — |
| CALL-based new recipes (36–40) | 5 | 9.9s | — | — |
| ReAct Agent (recipe 06, peak) | 1 | 27.3s | — | — |
| **Total** | **40** | — | — | **1,493.4s total** |

**Notable observations:** - **Recipe 22 (Text2SPL Demo)** now passes at 13.8s — this is the first successful run of the text-to-SPL compiler in a batch context, completing the "natural language → executable SPL" demonstration layer. Previously marked SKIP in all prior runs. - **Tool Use (recipe 36, 8.7s)** demonstrates the CALL/GENERATE boundary in its purest form: all arithmetic (`sum_values`, `percentage_change`, `format_currency`) is delegated to deterministic Python UDFs via CALL; the LLM is invoked exactly once to write the narrative summary. Zero hallucination risk on numerical values. - **Human-in-the-Loop (recipe 38, 3.7s)** is the fastest WORKFLOW recipe — the approval gate operates via simulated CALL, with no LLM involvement in the decision itself, demonstrating the CALL/GENERATE boundary for human-in-the-loop patterns. - **Doc Q&A (recipe 39, 1.7s)** is the fastest recipe overall — document loading via CALL is deterministic; only the answer extraction step calls the LLM. - **Portability Test (recipe 40)** uses only stdlib functions (`to_float`) with no `--tools` flag, demonstrating that the Standard Library eliminates common boilerplate tools. - **ReAct Agent (recipe 06, 27.3s)** executed multiple LLM calls autonomously — demonstrating the WHILE loop driving a genuine tool-reasoning loop without programmer intervention. - **Safe Generation (recipe 07, 10.3s)** demonstrates EVALUATE with semantic conditions: `WHEN 'high_quality'`, `WHEN 'acceptable'` — the LLM serves as its own quality judge, with no numeric parsing required. - **Reflection Agent (recipe 16, 132.3s)** is the slowest workflow, reflecting its deep meta-cognition loop. - **Interview Simulator (recipe 33, 126.7s)** demonstrates deep multi-turn role-play simulation within a single `spl2 run` command. - **Headline News Aggregator (recipe 37, 131.6s)** is slow because it fetches and summarizes multiple live news sources via CALL UDFs before generating the final report. - **Debate Arena (recipe 11, 91.8s)** and **Ensemble Voting (recipe 20, 61.2s)** both generated rich multi-perspective outputs within a single `spl2 run` command, with no Python orchestration code. - **Nested Procedures (recipe 25, 32.6s)** demonstrates three-level PROCEDURE composability — WORKFLOW calling PROCEDURE calling PROCEDURE — without any imperative orchestration code.

## Appendix G: Full Experiment Results

**Run configuration (v0.7, 2026-03-23):** - Date: 2026-03-23 - Model: gemma3 (via Ollama) - Hardware: GTX 1080 Ti (11 GB VRAM), single node - Runner: `python cookbook/run_all.py --adapter ollama --model gemma3` - Total wall time: 1,493.4s - Result: 40/40 SUCCESS

**Note:** A multi-GPU experiment on the Momagrid 2-node and 3-node configurations is planned for a follow-up evaluation (see Appendix E for Momagrid architecture). All 40 recipes now pass including recipe 22 (Text2SPL Demo), which ran successfully for the first time in v0.7.

**Pass/fail breakdown by category:**

| Category | Recipes | Passed | Failed | Notes |
|---|---|---|---|---|
| basics | 6 | 6 | 0 | Recipes 01, 02, 03, 22 (*first pass*), 23, 24 |
| agentic | 8 | 8 | 0 | Recipes 05, 06, 12, 16, 21, 25, 36 (Tool Use), 38 (Human-in-Loop) |
| reasoning | 3 | 3 | 0 | Recipes 09, 17, 35 |
| safety | 2 | 2 | 0 | Recipes 07, 18 |
| retrieval | 2 | 2 | 0 | Recipes 08, 39 (Doc Q&A, 1.7s) |
| multi-agent | 3 | 3 | 0 | Recipes 11, 14, 20 |
| application | 12 | 12 | 0 | Recipes 13, 15, 19, 27, 28, 29, 30, 31, 32, 33, 34, 37 (Headline News, 131.6s) |
| benchmarking | 4 | 4 | 0 | Recipes 04, 10, 26, 40 (Portability Test, 5.0s) |
| **Total** | **40** | **40** | **0** | **100% — 1,493.4s total** |

**Recipe 07 fix note:** The original recipe used numeric EVALUATE (`WHEN > 0.8 THEN`) to branch on a quality score, but `GENERATE quality_score()` returned a natural-language assessment string, not a bare float. The fix switches to semantic EVALUATE: `WHEN 'high_quality' THEN` / `WHEN 'acceptable' THEN`. The LLM now serves as its own quality judge via the semantic condition — which is exactly what EVALUATE's semantic mode is designed for. This is a better demonstration of the primitive than numeric scoring.

**Planned follow-up experiments:**

1. Run with `llama3.2` and `mistral` to measure cross-model variance.
2. Run full cookbook on a LAN Momagrid with 1-GPU, 2-GPU, or 3-GPU configurations to measure distributed throughput scaling and overhead (see Appendix E).
3. Measure energy consumption (joules/token) via differential method for $\eta_{\mathrm{AI}}$ baseline.
4. Quantify token savings of SPL implicit context management vs. manual prompting.

## Appendix H: Lines-of-Code Comparison — Self-Refining Agent

The self-refining agent (draft → critique → revise) is a widely used benchmark for agentic frameworks (cf. Auto-GPT [16]). Table H.1 compares implementations across SPL, LangGraph, AutoGen, and CrewAI. Counts include only logic lines (blank lines and comments excluded); boilerplate imports are counted once.

**Table H.1: LOC Comparison — Self-Refining Agent**

| Framework | LOC (approx.) | Notes |
|---|---|---|
| **SPL 2.0** | **22** | Full WORKFLOW with WHILE loop, EVALUATE gate, and EXCEPTION handler. No Python required. |
| LangGraph | 120–150 | StateGraph definition, node functions, edge routing, conditional edges, compile + invoke boilerplate. |
| AutoGen | 90–120 | AssistantAgent + UserProxyAgent setup, message-passing loop, termination condition, result extraction. |
| CrewAI | 80–110 | Agent + Task definitions, Crew assembly, process config, kickoff invocation. |

**SPL 2.0 implementation** (22 lines, reproduced from Section 6.5):

```
WORKFLOW self_refine
  INPUT:  @task text, @max_iterations int DEFAULT 5
  OUTPUT: @result text
DO
  GENERATE draft(@task) INTO @current

  WHILE @iteration < @max_iterations DO
    GENERATE critique(@current) INTO @feedback
    EVALUATE @feedback
      WHEN 'satisfactory' THEN
        COMMIT @current
      OTHERWISE
        GENERATE revised(@current, @feedback) INTO @current
    END
  END

  COMMIT @current
EXCEPTION
  WHEN MaxIterationsReached THEN
    COMMIT @current WITH status = 'partial'
END
```

**Key observations:**

- SPL 2.0 requires **5–7× fewer lines** than imperative frameworks for the same pattern.
- The EVALUATE primitive eliminates the need for regex parsing or numeric scoring logic to detect loop termination — a significant source of complexity in imperative implementations.

- The EXCEPTION block captures MaxIterationsReached with a single clause; imperative frameworks require explicit counter checks and try/except blocks scattered across multiple functions.
- The same `self_refine.spl` file runs on Ollama, Anthropic, OpenRouter, or Momagrid by changing a single CLI flag; imperative implementations are typically coupled to one provider SDK.

Beyond LOC, two structural advantages do not appear in the table but are significant in practice:

**Reviewability.** An imperative implementation (120–150 lines of Python across multiple classes and callbacks) requires execution to verify correctness. The 22-line SPL workflow is self-describing: a domain expert with no Python knowledge can read it, understand every step, and sign off. This enables a governance model where business stakeholders review and approve agentic workflows before deployment — structurally impossible with imperative code.

**Declarative operations.** The same `self_refine.spl` runs in any environment by swapping CLI flags:

```
# Dev        --adapter ollama  --model gemma3
# Test       --adapter openrouter --model gpt-4o
# Production --adapter anthropic --model claude-sonnet-4-6
```

No code changes across environments. Imperative implementations require environment-specific initialization, conditional adapter logic, or separate configuration files that must be kept in sync with the code.

**Note:** LOC counts for LangGraph, AutoGen, and CrewAI are estimated from public examples and documentation as of 2026-03. Exact counts vary with coding style; the order-of-magnitude difference is robust across stylistic variations.