arXiv:submit/7420152 [cs.CL] 29 Mar 2026

SPL 2.0: A Declarative Language for Agentic Workflow Orchestration

Wen G. Gong

2026-03-29

Abstract

We present SPL 2.0, a declarative language for orchestrating LLM-powered agentic workflows. Unlike imperative frameworks (e.g., LangGraph, AutoGen, CrewAI), SPL 2.0 expresses workflow intent in SQL-like syntax, separating **GENERATE** (LLM invocations) from **CALL** (deterministic functions) as first-class language constructs. Three novel primitives drive the design: **EVALUATE** for hybrid semantic/deterministic branching, **WHILE** for flow control, and **EXCEPTION** for LLM-aware error handling. A pluggable runtime supports various execution modes (via a single CLI flag): local (Ollama), cloud (OpenRouter, Anthropic), and decentralized inference network (Momagrid).

Our declarative design delivers compounding benefits: (1) workflows are readable and reviewable across engineering, governance, and business teams; (2) LOC (Line-of-Code) is reduced by 2 to 10x versus imperative frameworks; and (3) the same `.spl` script runs unchanged across Dev, Test, and Production environments, or from a single GPU to a decentralized GPU network. SPL 2.0 is validated through a cookbook of 30+ recipes spanning common agentic workflow patterns, achieving 100% success rate on a 4-node LAN GPU grid (3× throughput speedup over single-node sequential execution), with portability validated across 5 runtime adapters. The reference implementation is publicly available under Apache 2.0.

Keywords: declarative language, agentic workflows, LLM orchestration, structured prompt language, semantic evaluation

---

1. Introduction

1.1 The Problem: Fragmentation in LLM Programming

Building agentic AI systems today demands a daunting intersection of skills. A practitioner must master prompt engineering for effective LLM interaction, Python programming for orchestration logic, API wrangling across disparate model providers, and manual state management for multi-step reasoning. Each popular framework — LangGraph [1], AutoGen [2], CrewAI [3], Auto-GPT [16] — introduces its own abstractions, API surfaces, and execution models, creating a fragmented ecosystem where knowledge transfers poorly and vendor lock-in is the norm.

This situation echoes a pattern the data industry has seen before. In the early days of relational

databases, programmers wrote imperative C and COBOL code to traverse data structures, manage cursors, and handle errors — all in application-specific ways. The introduction of SQL [4] in the 1970s transformed this landscape by providing a declarative specification layer: users described *what* data they wanted, not *how* to retrieve it. This separation enabled decades of optimizer innovation underneath a stable surface language. The current state of LLM programming is analogous to the pre-SQL era: powerful capabilities buried under layers of imperative glue code.

The barrier to entry compounds the problem. To build even a simple self-refining agent — one that drafts, critiques, and revises its own output — a developer must write code across multiple framework abstractions. This excludes the vast population of domain experts, analysts, and researchers who understand their workflows but lack the software engineering skills to implement them as agentic systems.

1.2 The Opportunity: Declarative Agentic Workflows

Consider the world an AI application runs in today: a zoo of models — Llama, Gemma, Mistral, Claude, GPT, Qwen, Phi — deployed across a zoo of infrastructure: local GPUs in a developer's office, cloud APIs on three continents, and many data centers. Each model has different capabilities, different context limits, different cost profiles. Each deployment has different latency, availability, and authentication requirements.

A student new to AI who wants to build a simple self-refining essay agent faces an immediate question that has nothing to do with their goal: which model, on which infrastructure, accessed through which API, orchestrated with which Python framework? Before writing a single line of logic, they must make a dozen infrastructure decisions. Most give up. Those who persist write code that is tightly coupled to one model, one provider, one framework — and must be rewritten when any of those change.

The opportunity SPL 2.0 addresses is this: what if you could program against the entire zoo declaratively?

```
SELECT                          -- assemble what goes into the model
    system_role('You␣are␣a␣helpful␣assistant.'),
    context.question AS question
GENERATE answer(question)       -- invoke the model, capture what comes out
```

This is the complete I/O interface to any model, on any infrastructure. SELECT assembles the input. GENERATE invokes the model and captures the output. The runtime — not the programmer — decides which model in the zoo executes the inference, and where. The same two lines run on a local Ollama instance, an OpenRouter cloud API, or a Momagrid decentralized inference node on consumer-grade GPU, without modification.

A declarative approach to this zoo offers three structural advantages: (1) it enforces **separation of concerns** between what the workflow accomplishes and how it executes, allowing the same specification to run on local hardware, cloud APIs, or decentralized compute grids without modification. (2) it achieves **broad accessibility**: the SELECT/GENERATE mental model is simple enough for a student encountering AI programming for the first time, and powerful enough for production agentic workflows. (3) it enables **compiler-level optimization**: when the system understands intent rather than following imperative steps, it can merge redundant LLM calls, cache repeated contexts, and route queries to cost-optimal models — optimizations invisible to the user but not

achievable by an imperative framework, which by construction interleaves logic with execution order and thereby eliminates the system's visibility into intent.

This is no coincidence: SPL 2.0's design is a deliberate application of the SQL evolutionary pattern to LLM programming. We develop this parallel further in the next section.

1.3 Our Contribution: The SPL Stack

Figure 1 illustrates the **three-tier runtime stack**: Momagrid as the decentralized inference layer, SPL 1.0 as the declarative query layer, and SPL 2.0 as the agentic workflow orchestration layer. SPL 2.0 is an **additive extension** of SPL 1.0: every valid SPL 1.0 program runs unchanged inside the SPL 2.0 runtime, and the new WORKFLOW constructs are purely additive.

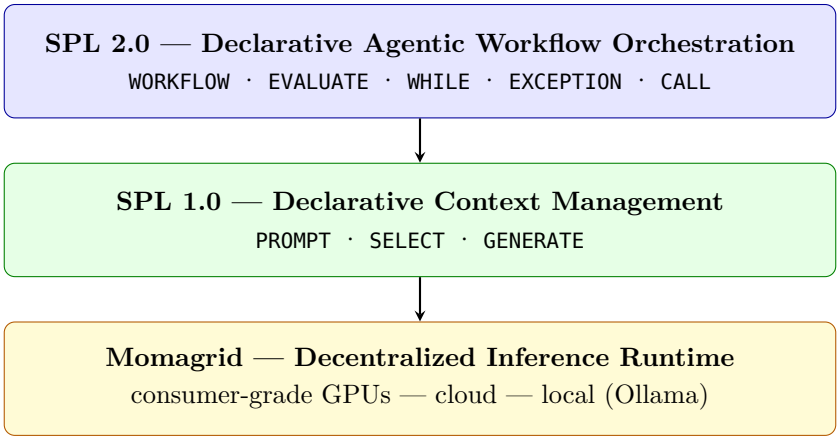

Figure 1: The SPL stack. SPL 2.0 contains SPL 1.0 as its foundation layer. Momagrid provides a decentralized inference runtime that SPL 2.0 targets via the `--adapter momagrid` flag, alongside local (Ollama) and cloud (OpenRouter, Anthropic) runtimes.

We make the following contributions:

1. **A four-layer execution architecture** (orthogonal to the runtime stack above) spanning:
   - *Natural Language* (interface) — the natural input modality
   - *Text2SPL* (compiler) — translates natural language intent into executable specification
   - *SPL 2.0* (specification) — a declarative language with a formal EBNF grammar
   - *Pluggable runtime* — adapters for Ollama (local), OpenRouter (cloud), Claude CLI (Anthropic), and Momagrid (decentralized)
2. **Three novel language primitives** for agentic workflows:
   - **EVALUATE** with hybrid deterministic-semantic condition evaluation, where the LLM itself serves as a runtime judge
   - **WHILE** with semantic termination conditions, enabling loops that iterate until output quality is "satisfactory" rather than until a counter reaches N
   - **EXCEPTION** for error-handling to address LLM-specific failure modes (HallucinationDetected, RefusalToAnswer, ContextLengthExceeded, BudgetExceeded, etc.)
3. **A reference implementation** consisting of a Python Click-based CLI, a recursive-descent parser, 5 runtime adapters, 231 automated tests, and a cookbook of 30+ recipes spanning 8 categories of agentic workflow patterns.

4. **The CALL/GENERATE design principle**: an explicit language-level separation between deterministic Python operations (CALL) and probabilistic LLM invocations (GENERATE), with the `@spl_tool` decorator as a Foreign Function Interface to the Python ecosystem.
5. **SPL Standard Library** (`spl.stdlib`): 55 SQL-equivalent built-in functions (type conversion, string, numeric, conditional, JSON, date, hashing, list) auto-loaded in every workflow without configuration.
6. **The `spl-cookbook` package** and **Recipe Kitchen** Streamlit UI: a pip-installable distribution with `recipe-kitchen` and `bake-all` CLI entry points, making the full cookbook accessible without manual path configuration.
7. **Backward compatibility** with SPL 1.0 [5]: all existing SPL 1.0 PROMPT statements execute unchanged in SPL 2.0, and the new WORKFLOW constructs are purely additive extensions.

1.4 Paper Organization

Section 2 surveys the motivation and positions SPL 2.0 within the landscape of related work. Section 3 presents the language design — primitives, semantics, and formal grammar. Section 4 describes the implementation architecture. Section 5 discusses compilation and optimization opportunities. Section 6 evaluates expressiveness, performance, cost efficiency, and developer experience. Section 7 discusses strengths, limitations, and future directions. Section 8 concludes. Appendix E describes the Momagrid decentralized runtime architecture. Appendix F contains the complete cookbook and end-to-end experiment results. Appendix G presents the full batch run data. Appendix H presents a case study comparing declarative SPL 2.0 with equivalent imperative implementations across three frameworks.

2. Motivation and Related Work

2.1 The Evolution of Data Programming

The history of data programming offers a clear precedent for the transition SPL 2.0 proposes. In the 1970s, IBM's System R project demonstrated that a declarative query language (SQL) could match or exceed the performance of hand-coded navigational database access, because the declarative form exposed optimization opportunities invisible in imperative code [4]. By the 1980s, SQL had become the industry standard, and SQL stored procedures [6] demonstrated that procedural control flow — loops, variables, exception handling — could be layered on top of declarative queries without sacrificing the optimizer's ability to reason about individual SELECT statements.

The parallel to LLM programming is direct. Today's imperative frameworks — LangGraph's state graphs, AutoGen's conversation loops, CrewAI's role-based agents — are the equivalent of pre-SQL navigational code: powerful, flexible, and impossibly fragmented. SPL 2.0 proposes that the same declarative approach can occur for LLM orchestration, following the same evolutionary pattern:

1. **Atomic declarative queries** (SQL SELECT → SPL PROMPT/SELECT): express what you want from a single LLM call.
2. **Procedural extensions** (SQL stored procedures → SPL WORKFLOW): add variables, loops, branching, and error handling around those queries.
3. **Optimizer infrastructure** (Catalyst → SPL IR): enable the system to rewrite and optimize execution plans.

2.2 Lessons from Proven Architectures

SPL 2.0's design draws deliberately from three proven systems:

From SQL stored procedures: block structure (`DO...END`), exception handling with typed exceptions (`WHEN HallucinationDetected THEN`), procedures and packages for code reuse, and the fundamental insight that procedural extensions should layer on top of declarative queries without disturbing the query optimizer.

From Apache Spark: the separation between logical and physical execution plans, the concept of a query optimizer (Catalyst) that rewrites execution strategies, and the idea that the same logical plan can execute on different backends (local, YARN, Kubernetes).

From the JVM: the stable intermediate representation (bytecode) that enables multiple frontend languages and multiple backend targets, and the principle that a well-defined IR is the foundation for an optimization ecosystem.

2.3 Current Approaches to Agentic Systems

The landscape of LLM orchestration systems can be organized along two axes: **declarative vs. imperative** paradigm and **query-level vs. workflow-level** scope.

Imperative Frameworks

LangGraph [1] models agent workflows as directed graphs with explicit state management. Nodes represent computation steps, edges represent transitions, and state is passed through a mutable dictionary. The framework provides fine-grained control but requires significant Python expertise and offers no formal grammar or optimizer.

AutoGen [2] focuses on multi-agent conversation patterns where agents exchange messages. Its strength lies in natural dialogue flows, but orchestration logic is embedded in Python code and tightly coupled to the runtime.

CrewAI [3] introduces a role-based metaphor where specialized agents collaborate. While intuitive for team-like workflows, it remains Python-first and provides no path to compilation or optimization.

PydanticAI [17] takes a code-first approach, rejecting complex abstractions in favor of standard Python with type hints. It prioritizes developer ergonomics over declarative specification.

Declarative Systems

LMQL [7] (ETH Zurich, 2022) is the pioneer of SQL-like LLM interaction, combining declarative queries with output constraints (type, length, regex). LMQL achieves 26–85% cost savings through constrained decoding. However, it operates at the *query level* — individual prompts with output validation — and does not provide constructs for multi-step workflows, iterative refinement, or exception handling.

DSPy [8] (Stanford, 2023) introduces declarative modules (Predict, ChainOfThought) and a compiler that automatically optimizes prompts via few-shot bootstrapping. DSPy achieves remarkable results — GPT-3.5 outperforming expert-crafted demonstrations by up to 46%. However, DSPy's "declarative" is at the *module interface level*: it optimizes what prompts to use, not how to orchestrate multi-step workflows. It provides no loops, no exception handling, and no formal grammar.

SGLang [9] (UC Berkeley, 2023) provides a structured generation frontend with runtime optimizations (RadixAttention for KV cache reuse), achieving up to 6.4x throughput improvement. SGLang optimizes *inference execution*, not workflow orchestration — it operates below the abstraction level of agentic patterns.

Specification Schemas

Agent Spec [10] (Oracle Labs, 2025) defines a YAML/JSON schema for specifying agents and workflows in a framework-agnostic way, with a standardized evaluation harness. Agent Spec solves *portability* (define once, run anywhere) but is a configuration format, not an executable language — it lacks control flow, iteration, and exception handling within the specification itself.

SLANG [11] (2025) proposes a minimal meta-language with three primitives (stake, await, commit) for formalizing multi-agent routing. SLANG's contribution is formalization of *coordination*, not computation — it does not address generation, evaluation, or error recovery.

Concurrent Work

PayPal [12] (December 2025) presents a declarative system separating agent workflow specification from implementation, achieving 60% reduction in development time and 3x deployment velocity. The PayPal DSL focuses on enterprise pipeline configuration (data serialization, filtering, API orchestration) across multiple backend languages. It does not formalize LLM-specific primitives like semantic evaluation or hallucination handling.

Auton [13] (February 2026) formalizes agent execution as an augmented POMDP with a declarative "Cognitive Blueprint" schema. Auton's contribution is governance and safety constraint enforcement via policy projection. It specifies agents as configuration artifacts (YAML/JSON) rather than providing an executable programming language.

Agentics 2.0 [14] (IBM, March 2026) introduces a logical transduction algebra where LLM calls are typed semantic transformations composable via algebraically grounded operators in asynchronous Map-Reduce programs. This is the closest concurrent work in terms of formalism, but takes an algebraic rather than SQL-like approach.

Workflow Orchestration and Adjacent Systems

Several established systems influence the workflow orchestration landscape but do not address LLM-specific concerns:

LangChain LCEL provides a declarative chain composition model via a Runnable interface, supporting piped composition of prompt templates, models, and output parsers. LCEL's composability is powerful for linear chains but does not provide iteration primitives, semantic condition evaluation, or a formal exception hierarchy for LLM-specific failures.

Haystack Pipelines (deepset) offers YAML-declarative pipeline definitions for document retrieval and question-answering workflows. Like Agent Spec, Haystack is a configuration format: it declares components and connections but delegates control flow and error recovery to the host application.

Prefect / Airflow are mature declarative workflow orchestrators for data engineering pipelines. Their abstractions (DAGs, tasks, flows) were designed for deterministic, side-effectful operations — not for semantic evaluation of LLM outputs, quality-based iteration, or hallucination handling. SPL 2.0 extends these orchestration patterns into the LLM domain.

OpenAI Assistants API provides built-in iteration and tool-use loops via a managed stateful conversation runtime. It abstracts control flow entirely from the developer, offering convenience at the cost of programmability: developers cannot express custom branching conditions, quality thresholds, or exception recovery policies.

None of these systems provide a programmable language with formal semantics for LLM-specific concerns. SPL 2.0 is distinguished by its formal EBNF grammar, explicit semantic evaluation primitive (EVALUATE), and structured LLM exception taxonomy (EXCEPTION).

Positioning Summary

Table 1: Landscape Comparison

System	Paradigm	Scope	Has Grammar	Semantic Eval	LLM Exceptions	Multi Runtime
LangGraph	Imperative	Workflow	No	No	No	No
AutoGen	Imperative	Conversation	No	No	No	No
CrewAI	Imperative	Role-based	No	No	No	No
LMQL	Declarative	Query	Yes	No[a]	No	No
DSPy	Declarative	Module	No	No[b]	No	No
SGLang	Declarative	Inference	Yes	No	No	No
Agent Spec	Declarative	Configuration	Yes[c]	No	No	Yes[d]
PayPal	Declarative	Pipeline	Yes	No	No	Yes
Agentics 2.0	Algebraic	Data workflow	Yes[e]	No	No	No
SPL 2.0	**Declarative**	**Query+Workflow**	**Yes**	**Yes**	**Yes**	**Yes**

[a] Structural only — grammar constrains token shape, not semantic meaning. [b] Compile-time optimization — not runtime semantic evaluation. [c] Schema-based (JSON/YAML) — not a formal procedural grammar. [d] Portable across runtimes via configuration. [e] Algebraic specification — not an executable procedural grammar.

SPL 2.0 combines a formal EBNF grammar, LLM-as-judge primitives (EVALUATE), an LLM-specific exception hierarchy (EXCEPTION), and a unified adapter layer spanning local, cloud, and decentralized execution — all within a single coherent language that spans both atomic queries and procedural workflows.

3. SPL 2.0 Language Design

3.1 Design Principles

SPL 2.0 adheres to the following design thinking:

1. Declarative intent over imperative implementation. The programmer specifies *what* the workflow should accomplish — not *which model*, *which node*, or *which execution path* achieves it. This is the essential simplification: the zoo of models deployed around the world is the runtime's problem, not the programmer's. A student writing their first AI workflow and a senior engineer deploying a production pipeline write the same SELECT/GENERATE syntax; the runtime resolves the infrastructure details beneath them. This separation enables the compiler and runtime to make optimization decisions that are invisible to the user but structurally unavailable to an imperative framework, where execution order is hard-coded by the programmer rather than derived from declared intent.

2. LLM-aware primitives. Rather than repurposing generic control flow constructs for LLM orchestration, SPL 2.0 provides primitives designed specifically for the unique characteristics of LLM computation: probabilistic output, semantic quality assessment, hallucination risk, context length limits, and cost sensitivity.

3. Compilation target. The language is designed to produce an optimizable intermediate representation (IR), enabling future optimizer passes without changing the surface language.

4. Unified runtime — one specification, any model, any infrastructure. The adapter pattern decouples the language from any specific model provider. The same `.spl` file executes on a local Ollama instance (a developer's laptop), OpenRouter (100+ cloud models), Claude CLI (Anthropic), or a Momagrid decentralized inference node (consumer-grade GPU, anywhere in the world) — without changing a single line. The programmer writes against the zoo; the adapter layer tames it. This is not merely an engineering convenience — it is the architectural prerequisite for broad accessibility.

5. Backward compatibility. All SPL 1.0 [5] programs are valid SPL 2.0 programs. New constructs (WORKFLOW, EVALUATE, WHILE, EXCEPTION) are purely additive extensions.

6. Explicit probabilistic/deterministic boundary (CALL vs. GENERATE). Every operation in an SPL workflow declares its computational character at the language level. `GENERATE` invokes an LLM — probabilistic, creative, slow, expensive. `CALL` invokes a Python function — deterministic, testable, instant, free. Most LLM frameworks blur this line; SPL draws it in ink. This forces workflow authors to be explicit about which operations can be cached, replayed, or optimized, and which cannot. As a corollary, the `@spl_tool` decorator serves as a Foreign Function Interface (FFI): any Python callable becomes a first-class `CALL` target with one decorator line, connecting SPL to the entire Python ecosystem without extending the language itself.

3.2 Core Primitives

The mental model of SPL 2.0 is grounded in two fundamental primitives that mirror the I/O contract with an LLM: **SELECT** assembles the context that flows *into* the model; **GENERATE** invokes the model and captures what flows *out*. Every other construct in the language is either a named container for SELECT and GENERATE (PROMPT, WORKFLOW) or a flow control mechanism that orchestrates when and how they are invoked (EVALUATE, WHILE, EXCEPTION, PROCEDURE).

This two-sided I/O model deliberately parallels SQL's interaction with a database — SELECT gathers data in, the result set flows out — while extending it to handle the probabilistic, semantic nature of LLM computation.

3.2.1 SELECT — Context Assembly (Input to LLM)

SELECT is the *input* primitive. It aggregates context from multiple sources into the prompt payload that will be sent to the LLM. Nothing executes until SELECT has finished assembling the input.

```
SELECT
    system_role('You␣are␣a␣research␣analyst.␣Be␣precise␣and␣cite␣evidence.'),
    context.topic AS topic,
    memory.get('prior_findings') AS background,
    rag.query('recent␣developments␣in␣' + @topic, top_k=5) AS references
```

SPL 2.0 provides four context sources within SELECT:

Source	Syntax	What it provides
System role	`system_role('...')`	The model's behavioral instruction
Runtime parameter	`context.field_name`	CLI-injected value at execution time
Persistent memory	`memory.get('key')`	Key-value store across workflow runs
Semantic retrieval	`rag.query('text', top_k=N)`	Top-N chunks from a vector index

Token budgets on the input side are enforced with `WITH BUDGET n TOKENS`, preventing context window overflow before the LLM is ever invoked.

The analogy to SQL is precise: SQL's `SELECT ... FROM ... WHERE` gathers and filters rows from a database; SPL's `SELECT` gathers and filters context from heterogeneous sources (parameters, memory, retrieval indices) into a prompt. In both cases the reader specifies *what information to collect* — not how the storage layer retrieves it.

3.2.2 GENERATE — LLM Invocation (Output from LLM)

GENERATE is the *output* primitive. It names the generation task, invokes the LLM with the assembled SELECT context, and captures the result.

```
GENERATE answer(prompt)                          -- standalone: output to stdout
GENERATE analysis(@research) INTO @result        -- workflow: capture into variable
GENERATE summary(@analysis) INTO @summary
    WITH OUTPUT BUDGET 500 TOKENS                -- cap output length
    USING MODEL 'claude-opus-4-6'                -- model override per step
```

The argument passed to GENERATE — `answer`, `analysis`, `summary` — is a semantic label that tells the LLM what kind of output to produce. This is the declarative analog of a function name: the developer states *what* the model should generate, not *how*.

Key behaviors are:

- **Standalone** (`GENERATE x()`): output goes to stdout; used inside PROMPT statements.
- **Capturing** (`GENERATE x()INTO @var`): output is stored in a workflow variable for downstream steps.
- **Output budget** (`WITH OUTPUT BUDGET n TOKENS`): limits response length, enabling cost control.
- **Model binding** (`USING MODEL '...'`): per-step model override without editing the file; the CLI `--model/-m` flag applies globally across all GENERATE calls at runtime.

GENERATE has no direct SQL equivalent — SQL does not produce probabilistic text — but the pattern it enables is familiar: `GENERATE` is to SPL what a function call is to a stored procedure, except the "function" is the LLM itself.

Together, SELECT and GENERATE form the complete I/O interface to the LLM. SELECT is the *input contract*; GENERATE is the *output contract*. PROMPT and WORKFLOW are simply named scopes that bind them together.

3.2.3 PROMPT — Atomic LLM Query

The PROMPT statement is SPL's fundamental named unit: a SELECT clause followed by a GENERATE clause, given a name and made independently executable.

```
PROMPT hello_world
SELECT
    system_role('You␣are␣a␣friendly␣assistant.␣Introduce␣yourself␣and␣SPL␣2.0␣in␣two␣sentences.')
GENERATE greeting()
```

This four-line program is a complete, executable SPL file. SELECT assembles the system instruction; GENERATE invokes the LLM and streams the result to stdout. No imports, no configuration, no API keys in code.

Parametric context is supported via `context.param`:

```
PROMPT ollama_proxy
SELECT
    system_role('You␣are␣a␣helpful,␣knowledgeable␣assistant.'),
    context.prompt AS prompt
GENERATE answer(prompt)
```

Runtime parameter injection via the CLI: `spl run proxy.spl --adapter ollama prompt="Explain quantum computing"`. The model can be overridden at runtime with `--model/-m` without editing the `.spl` file — the CLI flag patches the GENERATE binding at execution time.

PROMPT is to SPL what a single SQL SELECT statement is to SQL: the atomic, declarative query unit. WORKFLOW (Section 3.2.4) is the procedural extension that adds state, loops, and branching around that query unit.

3.2.4 WORKFLOW — Procedural Orchestration

The WORKFLOW statement adds procedural flow control around PROMPT-level operations: variables, loops, branching, and exception handling.

```
WORKFLOW chain_of_thought
    INPUT: @topic TEXT
    OUTPUT: @summary TEXT
DO
    GENERATE research(@topic) INTO @research
    GENERATE analyze(@research) INTO @analysis
    GENERATE summarize(@analysis) INTO @summary
    COMMIT @summary WITH status = 'complete'
END
```

Key constructs:

- **INPUT/OUTPUT** declarations define the workflow's typed interface.
- **@variables** (prefixed with `@`) provide mutable state within the workflow scope.
- **GENERATE...INTO** captures LLM output into a variable for use in subsequent steps.
- **COMMIT** finalizes the workflow output with optional metadata. Execution halts after COMMIT.

3.2.5 EVALUATE — LLM-Based Condition Evaluation

EVALUATE is SPL 2.0's first novel primitive. It provides branching where the condition itself may be evaluated by an LLM — a capability no existing language formalizes.

SPL 2.0 supports four evaluation types, detected automatically from the condition syntax:

Deterministic numeric evaluation (zero LLM cost):

```
EVALUATE @score
    WHEN > 0.8 THEN COMMIT @result WITH status = 'high_quality'
    WHEN > 0.5 THEN GENERATE improved(@result) INTO @result
    OTHERWISE RETRY WITH temperature = 0.1
END
```

When the condition is a numeric or boolean comparison operator (`>`, `<`, `>=`, `<=`, `=`, `!=`), the runtime evaluates it directly with no LLM call.

Deterministic string evaluation (zero LLM cost):

```
EVALUATE @keyword_class
    WHEN STARTSWITH 'harmful' THEN
        COMMIT 'Cannot␣process␣this␣request.' WITH status = 'blocked'
    WHEN STARTSWITH 'off_topic' THEN
        COMMIT 'Outside␣my␣scope.' WITH status = 'blocked'
    OTHERWISE
        -- Continue to LLM classification
END

EVALUATE @pii_report
    WHEN CONTAINS('pii_found') THEN CALL redact_pii(@input) INTO @clean
    OTHERWISE SET @clean = @input
END
```

`STARTSWITH 'prefix'` and `CONTAINS('substring')` trigger a deterministic string test with no LLM call. This is the correct form when a CALL tool returns a structured string value (e.g., `harmful:keyword=bomb`) and the workflow needs to branch on its prefix or content without spending tokens on LLM judgment. `CONTAINS` supports multiple alternatives with `OR`: `CONTAINS('err') OR CONTAINS('fail')`.

Semantic evaluation (LLM as judge):

```
EVALUATE @summary
    WHEN 'coherent' THEN COMMIT @summary
    WHEN 'hallucination' THEN RETRY
    WHEN 'incomplete' THEN GENERATE expand(@summary) INTO @summary
END
```

When the condition is a plain quoted string literal, the runtime invokes the LLM as a judge: "Does this output satisfy the criterion 'coherent'?" This is a low-cost classification call (typically 50–200 tokens) that enables the workflow to make quality-based decisions without external tooling.

Hybrid evaluation (mixed):

```
EVALUATE @response
    WHEN token_count > 4000 THEN truncate          -- deterministic numeric
```

```
    WHEN STARTSWITH 'blocked' THEN log_block        -- deterministic string
    WHEN 'off_topic' THEN reframe                   -- semantic (LLM judge)
    WHEN confidence < 0.7 THEN verify               -- deterministic numeric
END
```

The detection rule is unambiguous: comparison operators trigger deterministic numeric evaluation; `STARTSWITH` or `CONTAINS` trigger deterministic string evaluation; a plain quoted string literal triggers semantic LLM evaluation. No explicit annotations are required.

Semantic Evaluation Protocol. When a semantic condition is detected, the runtime sends a standardized two-message judge call to the LLM. The system prompt is: *"You are a quality judge. Answer ONLY 'true' or 'false' with no additional text."* The user prompt is: *"Does the following output satisfy the criterion '<condition>'?"* followed by the value being evaluated. The runtime extracts the first token of the response and maps it to a boolean; any non-boolean response is treated as `false` and may raise a `QualityBelowThreshold` exception. By default, the judge call uses the same adapter and model as the enclosing GENERATE statement, but a dedicated judge model can be specified per-EVALUATE via `USING MODEL`.

3.2.6 WHILE...DO — Iteration with Semantic Termination

WHILE is SPL 2.0's second novel primitive. It enables loops where the termination condition can be LLM-judged, with a built-in safety guard against infinite iteration.

```
WHILE @iteration < @max_iterations DO
    GENERATE critique(@current) INTO @feedback

    EVALUATE @feedback
        WHEN 'satisfactory' THEN
            COMMIT @current WITH status = 'complete'
        OTHERWISE
            GENERATE refined(@current, @feedback) INTO @current
            @iteration := @iteration + 1
    END
END
```

The construct supports the same deterministic/semantic/hybrid condition types as EVALUATE. A global iteration limit (default: 10) is enforced by the runtime, raising `MaxIterationsReached` if exceeded. This prevents runaway loops while allowing the workflow to iterate based on output quality rather than arbitrary counts.

Context accumulation across iterations is a natural pattern: each iteration's output becomes the input for the next, enabling progressive refinement without explicit history management.

3.2.7 EXCEPTION — Structured LLM Error Handling

EXCEPTION is SPL 2.0's third novel primitive. It defines a formal taxonomy of LLM-specific failure modes, each with clear semantics and recovery strategies.

Table 3: LLM Exception Taxonomy

Exception	Trigger Condition	Typical Recovery
HallucinationDetected	Output fails fact-checking	Regenerate with lower temperature
RefusalToAnswer	Model declines the request	Log and COMMIT partial result
ContextLengthExceeded	Input exceeds model context window	Compress context and RETRY
ModelOverloaded	Provider rate limit or capacity error	Wait and RETRY
BudgetExceeded	Token budget exhausted	COMMIT best available result
QualityBelowThreshold	Output quality score too low	Regenerate or escalate to better model
MaxIterationsReached	WHILE loop hit iteration limit	COMMIT current best effort

These exceptions are caught in structured handlers:

```
WORKFLOW safe_generation
    INPUT: @prompt text
    OUTPUT: @result text
DO
    GENERATE response(@prompt) INTO @result
    GENERATE quality_score(@result) INTO @score

    EVALUATE @score
        WHEN > 0.8 THEN COMMIT @result WITH status = 'high_quality'
        WHEN > 0.5 THEN
            GENERATE improved(@result, @prompt) INTO @result
            COMMIT @result WITH status = 'refined'
        OTHERWISE RETRY WITH temperature = 0.1
    END

EXCEPTION
    WHEN HallucinationDetected THEN
        GENERATE response(@prompt) INTO @result
            WITH TEMPERATURE 0.1
        COMMIT @result WITH status = 'conservative'
    WHEN ContextLengthExceeded THEN
        RETRY WITH temperature = 0.1
    WHEN RefusalToAnswer THEN
        COMMIT @prompt WITH status = 'refused'
    WHEN BudgetExceeded THEN
        COMMIT @result WITH status = 'truncated'
END
```

This 25-line workflow expresses a production-grade generation pattern — quality checking, iterative refinement, and graceful degradation across four failure modes — that would require 80–120 lines of Python in LangGraph with manual try/except blocks and custom error classes.

No existing language or framework formalizes LLM failure modes as a typed exception hierarchy. LangGraph and AutoGen rely on Python's generic exception mechanism; LMQL and DSPy have no error handling constructs at all.

3.2.8 PROCEDURE — Reusable Workflow Patterns

PROCEDUREs provide parameterized, reusable workflow components that can be composed via CALL:

```
PROCEDURE researcher(topic TEXT) RETURNS TEXT
DO
    GENERATE research_facts(topic) INTO @facts
    GENERATE identify_key_themes(@facts) INTO @themes
    @result := @facts + '\n\nKey␣Themes:\n' + @themes
    COMMIT @result
END
```

Procedures enable a standard library of agentic patterns: ReAct (reasoning + acting), Self-Refine (iterative improvement), Tree-of-Thought (branching exploration), and others can be implemented once and reused across workflows via `CALL researcher(@topic)INTO @result`.

3.2.9 Standard Library and User-Defined Functions (CALL Extension System)

SPL Standard Library. SPL 2.0 ships with `spl.stdlib`, a 55-function built-in library auto-loaded without any `--tools` flag or import statement. The functions mirror SQL built-ins, organized into 11 categories:

Category	Representative functions
Type conversion	`to_float`, `to_int`, `to_text`, `to_bool`
String	`upper`, `lower`, `trim`, `substr`, `replace`, `concat`
Pattern matching	`like`, `startswith`, `endswith`, `regexp_match`
Numeric	`abs_val`, `round_val`, `ceil_val`, `clamp`
Conditional	`coalesce`, `iif`, `nullif`
Null/empty checks	`isnull`, `isblank`, `nvl`
Text aggregates	`word_count`, `char_count`, `line_count`
JSON	`json_get`, `json_set`, `json_keys`, `json_length`
Date/time	`now_iso`, `date_format_val`, `date_diff_days`
Hashing	`md5_hash`, `sha256_hash`
List/array	`list_get`, `list_length`, `list_contains`

The `to_float` function is representative of the stdlib's design philosophy: it extracts the first numeric token from any string, so LLM output like `"Quality score:  0.85\n"` returns `0.85` without a secondary LLM call. This is the CALL/GENERATE boundary in practice — the computation is deterministic and belongs in CALL, not GENERATE.

User-Defined Functions (UDFs). The `@spl_tool` decorator is SPL's Foreign Function Interface. Any Python callable decorated with `@spl_tool` becomes a first-class `CALL` target, connecting SPL to the entire Python ecosystem:

```
from spl.tools import spl_tool

@spl_tool
def lookup_customer(customer_id: str) -> str:
    row = db.query("SELECT␣name,␣tier␣FROM␣customers␣WHERE␣id␣=␣?",
                   customer_id)
    return json.dumps(row) if row else ""
```

```
CALL lookup_customer(@id) INTO @customer_json   -- deterministic, zero LLM cost
GENERATE write_response(@customer_json) INTO @reply  -- LLM for the language part
```

Database queries, REST API calls, ML inference (scikit-learn, transformers), file I/O, PII detection, embedding computation — all become `CALL` targets with a single decorator line. The `--tools tools.py` CLI flag loads them at runtime; the language itself requires no extension.

The dispatch order for `CALL foo(@arg)` is: (1) registered UDF; (2) SPL PROCEDURE defined in the same file; (3) LLM fallback (treated as GENERATE). UDFs take priority, enabling domain-specific overrides of stdlib functions.

3.3 Cross-Cutting Concerns

3.3.1 Context Management

SPL 2.0 provides four context sources within the SELECT clause (see Section 3.2.1): `system_role`, `context.field_name`, `memory.get`, and `rag.query`.

Token budgets can be specified at both input (`WITH BUDGET n TOKENS`) and output (`WITH OUTPUT BUDGET n TOKENS`) levels.

3.3.2 Security and Accounting

For deployments on untrusted infrastructure (e.g., Momagrid's decentralized grid), SPL 2.0 supports security classifications and accounting tags:

```
WORKFLOW sensitive_analysis
    INPUT: @data TEXT
    OUTPUT: @result TEXT
    SECURITY: CLASSIFICATION: confidential
    ACCOUNTING: BILLABLE_TO: 'project-alpha'
DO
    ...
END
```

These metadata blocks are preserved through the compilation pipeline and enforced by the runtime adapter.

3.4 Formal Grammar

The complete EBNF grammar is provided in Appendix A. The grammar defines 12 statement types, 3 condition types (deterministic, semantic, hybrid), 5 primitive types (text, number, boolean, chunk, json), composite types (list, map), and metadata blocks for security, accounting, and compression.

SPL 2.0 extends SPL 1.0 with 8 new statement types (WORKFLOW, EVALUATE, WHILE, COMMIT, RETRY, RAISE, assignment, PROCEDURE) and 1 new condition type (semantic). All SPL 1.0 productions remain unchanged.

4. Implementation

4.1 Architecture Overview

The SPL 2.0 reference implementation is organized as a four-layer stack:

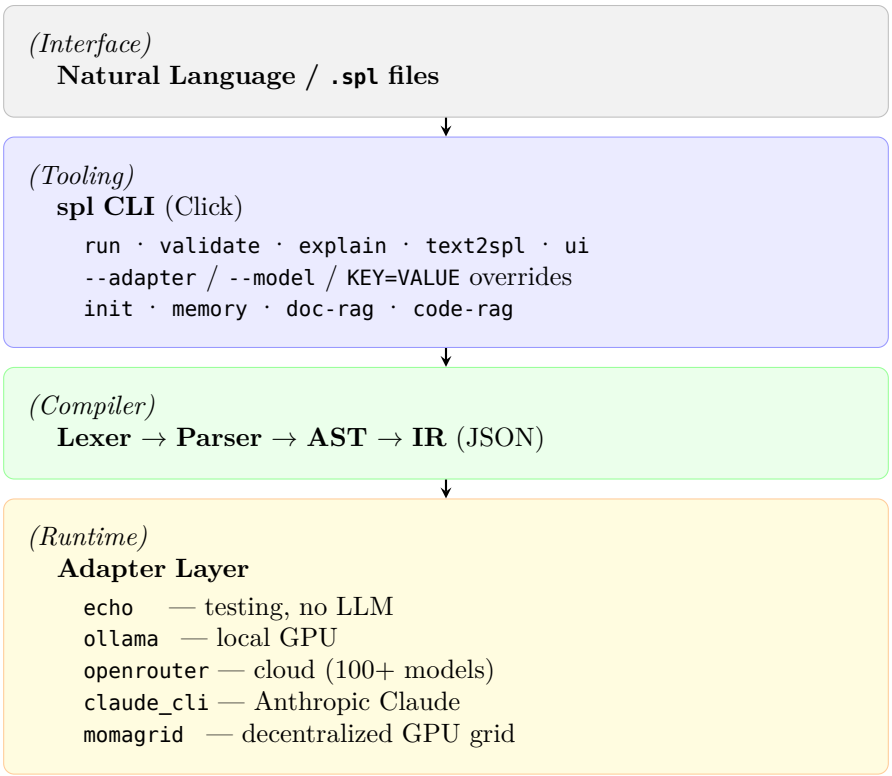

Figure 2: SPL 2.0 system architecture. The four layers correspond to the conceptual stack described in Section 1.3.

4.2 The SPL CLI

The CLI is built on Python's Click framework. The top-level entry point is `spl`:

```
$ spl --help
Usage: spl [OPTIONS] COMMAND [ARGS]...

  SPL 2.0 -- Structured Prompt Language.  SQL for Agentic Workflow
  Orchestration.
```

CLI Command Reference

Command	Description
`spl run`	Execute a `.spl` file; print each PROMPT/WORKFLOW result. Accepts trailing `KEY=VALUE` runtime parameters.
`spl validate`	Parse and validate a `.spl` file for syntax errors without executing (zero LLM cost).
`spl explain`	Show the execution plan (prompt plan) for a `.spl` file without executing.
`spl text2spl`	Compile a natural language description into SPL 2.0 code.
`spl ui`	Launch the Text2SPL Knowledge Studio (Streamlit app).
`spl adapters`	List all available LLM adapter engines and their status.
`spl config`	Manage SPL 2.0 configuration (`~/.spl/config.yaml`).
`spl memory`	Manage the SPL persistent memory store (`.spl/memory.db`).
`spl cache`	Manage the prompt cache (`.spl/memory.db`).
`spl doc-rag`	Manage the document RAG vector store (`.spl/vectors`); index documents for `rag.query()`.
`spl code-rag`	Manage the Code-RAG store used by the Text2SPL compiler for few-shot example retrieval.
`spl version`	Print the SPL engine version.

Common usage patterns:

```
# Zero-config execution
spl run cookbook/01_hello_world/hello.spl

# With adapter and model override
spl run cookbook/02_ollama_proxy/proxy.spl --adapter ollama -m gemma3 \
    prompt="Explain␣quantum␣computing"

# Validate syntax without running (no LLM needed)
spl validate cookbook/05_self_refine/self_refine.spl

# Show execution plan
spl explain cookbook/09_chain_of_thought/chain.spl

# Compile natural language to SPL
spl text2spl "summarize␣a␣document␣with␣a␣2000␣token␣budget" \
    --adapter ollama -m gemma3 --mode prompt

# Launch browser UI
spl ui
```

Key design decisions are:

- **Trailing KEY=VALUE parameters** use Click's UNPROCESSED argument handling, making the command line feel natural: `spl run file.spl topic="AI"` rather than `spl run file.spl -p topic="AI"`.
- **Runtime model override** via `--model/-m` patches the model binding on all PROMPT statements at execution time, allowing model comparison without editing `.spl` files.
- **Auto-initialization** creates workspace directories (`.spl/`, memory store, RAG index) on first run.

4.3 Lexer and Parser

The compiler pipeline consists of a tokenizer and recursive-descent parser that produce an abstract syntax tree (AST).

The **tokenizer** handles SPL-specific challenges including keyword-as-identifier resolution (e.g., `context.prompt` where `prompt` is also a keyword), `@variable` prefixing, and the `:=` assignment operator. Comment support follows SQL convention (`--` for single-line).

The **parser** produces typed AST nodes: `PromptStatement`, `WorkflowStatement`, `GenerateClause`, `EvaluateClause`, `WhileStatement`, `ExceptionBlock`, `ProcedureDeclaration`, and others. Each node carries source location information for error reporting.

4.4 Adapter Architecture for LLM Model/Provider

The `--adapter` flag encompasses two conceptual dimensions that together define the complete execution environment for an SPL workflow:

- **LLM Model runtime** (two sub-dimensions):
  - `--adapter` (or `--llm-provider`): the backend infrastructure — `ollama` (local), `openrouter` (cloud), `anthropic`, `momagrid` (decentralized), etc.
  - `--model` (`-m`): the specific model within that provider — `gemma3`, `claude-sonnet-4-6`, `gpt-4o`, etc.
- **Integration patterns** (planned extensions):
  - `--tools`: external tool connectors (web search, code execution, PDF ingestion, transcription)
  - `--datasets`: data source bindings for RAG or batch processing workflows
  - `--resources`: named resource slots (memory stores, vector indices, external APIs)

Together, the LLM runtime dimension and the integration pattern dimension provide comprehensive adapter coverage: the same `.spl` file can route inference to any provider and pull context from any data source, purely through CLI flags — without modifying the SPL source.

Runtime adapters implement a common interface:

```
class BaseAdapter:
    def generate(self, prompt, system_role, model, **kwargs) -> Result
    def list_models(self) -> list[str]
    def count_tokens(self, text) -> int
```

New adapters are registered via `register_adapter(name, cls)`, allowing third-party extensions without modifying core code.

Table 4: Adapter Capabilities

Adapter	Generate	List Models	Model Override	Auth
echo	Yes	Yes (static)	Yes	None
ollama	Yes	Yes (API)	Yes	None
openrouter	Yes	Yes (API)	Yes	API key
anthropic	Yes	Yes (API)	Yes	API key
claude_cli	Yes	Yes (static)	Yes	CLI auth
momagrid	Yes	Yes (fallback)	Yes	API key

The **echo adapter** deserves special mention: it returns deterministic responses without invoking any LLM, enabling the entire test suite (231 tests) to run without GPU hardware or API credentials. This is critical for CI/CD and for validating the parser and execution engine in isolation.

5. Compilation and Optimization

5.1 The Text2SPL Compiler

The top layer of the SPL stack — translating natural language intent into executable SPL 2.0 code — has a working prototype accessible via `spl text2spl`. The current prototype handles simple PROMPT pattern generation: given a natural language description, it produces a syntactically correct PROMPT statement with appropriate SELECT and GENERATE clauses. For example:

```
$ spl text2spl "summarize␣a␣document␣with␣a␣2000␣token␣budget" \
    --adapter ollama -m gemma3 --mode prompt
```

produces a syntactically valid `.spl` file that `spl validate` accepts without error. The prototype also supports `--mode workflow` for WORKFLOW generation and `--mode auto` where the model decides the appropriate form; these modes are less mature and their output requires manual review before use.

The full compiler pipeline envisioned for production would:

1. Parse natural language descriptions of workflow intent.
2. Extract the agentic pattern (self-refine, ReAct, chain-of-thought, etc.).
3. Generate syntactically correct SPL 2.0 code, including full WORKFLOW structures.

Full WORKFLOW generation from natural language — "I want an agent that researches a topic, analyzes the findings, and writes a summary" — remains future work. The declarative form makes this tractable: because SPL 2.0 specifies intent rather than execution steps, the mapping from natural language to SPL is structurally simpler than generating imperative Python code.

5.2 Intermediate Representation (IR)

The current implementation compiles SPL 2.0 source to an AST, which is serialized as JSON via the `spl explain` command. The planned IR extends this with:

- **Execution cost annotations**: estimated token consumption per GENERATE node.
- **Dependency edges**: data flow between GENERATE steps for parallelism analysis.
- **Adapter hints**: preferred runtime for each node based on model requirements.

5.3 Future Optimization Opportunities

This section describes optimization strategies that are *structurally enabled* by the declarative form — not yet implemented, but architecturally possible in ways that are impossible in imperative frameworks where execution order is hard-coded by the programmer. The declarative form is the prerequisite; the optimizer is future work.

The key claim: an imperative framework cannot perform these optimizations because the programmer's control flow interleaves logic with execution sequence, eliminating the system's visibility into intent. SPL 2.0's separation of specification from execution preserves that visibility.

GENERATE merging: When consecutive GENERATE calls share the same context, they can be merged into a single LLM call with structured output parsing, reducing both latency and cost.

Deterministic condition hoisting: In a WHILE loop, deterministic conditions in EVALUATE blocks can be evaluated before invoking the LLM for semantic conditions, avoiding unnecessary LLM calls.

SELECT caching: When the same RAG query or memory lookup appears multiple times, the result can be cached and reused.

Model selection optimization: A cost-quality optimizer can route simple GENERATE calls to smaller, cheaper models while reserving large models for complex reasoning steps.

Parallel GENERATE: When consecutive GENERATE steps have no data dependencies (e.g., generating pro and con arguments in a debate), they can execute in parallel.

Table 5: Optimization Opportunity Analysis

Optimization	Applicable Patterns	Estimated Impact	Status
GENERATE merging	Chain-of-Thought, Multi-Agent	Reduce LLM calls	Planned
Condition hoisting	Self-Refine, Reflection	Reduce judge calls	Planned
SELECT caching	RAG Query	Eliminate redundant retrieval	Planned
Model downgrade	Hello World, Proxy	Cost reduction	Planned
Parallel GENERATE	Debate, Tree-of-Thought, Ensemble	Latency reduction	Planned

6. Evaluation

6.1 Expressiveness

We compare SPL 2.0's expressiveness against imperative frameworks by measuring lines of code (LOC) required to implement common agentic patterns.

Table 6: Agentic Pattern Expressiveness (Lines of Code)

Pattern	SPL 2.0	LangGraph	AutoGen	CrewAI
Chain of Thought	14	81	80	81
Self-Refine	35	98	74	79
ReAct Agent	43	126	93	100
RAG Query	7	78	66	73

SPL 2.0 achieves 2 to 10x reduction in LOC compared to imperative alternatives (average across LangGraph, AutoGen, and CrewAI: 5.8x for Chain of Thought, 2.4x for Self-Refine, 2.5x for ReAct Agent, 10.3x for RAG Query). More significantly, the SPL code is *readable by non-programmers* — the SQL-like syntax is familiar to analysts, data scientists, and researchers who would not write Python framework code.

6.2 Performance

Table 7: Implementation Metrics

Metric	Value	Notes
Total unit tests	231	All passing

Metric	Value	Notes
Cookbook recipes	40	All 40 benchmarked including recipe 22 (Text2SPL Demo), which was previously planned/skipped; 5 new recipes (36–40) added
Recipe categories	8	basics, agentic, reasoning, safety, retrieval, multi-agent, application, benchmarking
Runtime adapters	5	echo, ollama, openrouter, claude_cli, momagrid
SPL 2.0 statement types	12	PROMPT, WORKFLOW, GENERATE, EVALUATE, WHILE, COMMIT, RETRY, RAISE, PROCEDURE, CALL, assignment, EXCEPTION
Cookbook success rate	40/40 (100%)	gemma3 via Ollama, single GTX 1080 Ti; all recipes pass; recipe 22 now fully operational
PROMPT recipe avg latency	2.3s	7 PROMPT recipes; range 0.6s–6.1s depending on output length
WORKFLOW recipe avg latency	44.8s	33 WORKFLOW recipes; range 3.7s–132.3s
Fastest workflow	1.7s	Document Q&A (recipe 39); CALL + GENERATE minimal pipeline
Slowest workflow	132.3s	Reflection Agent (recipe 16); deep meta-cognition loop
CALL-based recipe avg latency	4.8s	Recipes 36, 38, 39, 40 (Tool-Use, Human-in-Loop, Doc Q&A, Portability); CALL overhead is negligible vs pure-GENERATE workflows
Total cookbook wall time	1,493.4s	40/40 recipes; single GTX 1080 Ti, gemma3 via Ollama, sequential execution. Per-recipe times vary between runs due to LLM output length variance; WORKFLOW avg stable at 42–45s

6.3 Cost Efficiency

The declarative approach enables token savings through two mechanisms:

1. **Implicit context management**: SPL automatically constructs minimal prompts from SELECT clause elements, avoiding the verbose system prompts typical of manual prompt engineering.
2. **Budget enforcement**: `WITH BUDGET n TOKENS` and `WITH OUTPUT BUDGET n TOKENS` clauses enable explicit token accounting, preventing runaway costs.

Future optimizer passes (GENERATE merging, model downgrade) will provide additional cost reduction.

6.4 Developer Experience

Table 8: Barrier to Entry

Skill Required	SPL 2.0	LangGraph	AutoGen	CrewAI
Python programming	No	Yes	Yes	Yes
API knowledge	No	Yes	Yes	Yes
Prompt engineering	Minimal	Full	Full	Full
SQL familiarity	Helpful	N/A	N/A	N/A
CLI usage	Basic	N/A	N/A	N/A

SPL 2.0's zero-config CLI (`spl run file.spl --adapter ollama`) eliminates the setup friction that characterizes imperative frameworks: no virtual environments to configure, no API clients to instantiate, no state management classes to define.

7. Discussion

7.1 Strengths

Declarative form enables optimization. Because SPL specifies *what* should happen rather than *how*, the compiler can analyze data dependencies, identify parallelism opportunities, and select cost-optimal models — optimizations that are not structurally available in imperative code where the execution order is hard-coded by the programmer, leaving no intent visible to the system.

Adapter pattern broadens reach. The same `.spl` file runs on local GPU (Ollama), cloud APIs (OpenRouter, Claude), and decentralized inference network (Momagrid). No other system spans all three deployment models from a single specification.

LLM-aware primitives address real failure modes. The EXCEPTION taxonomy (HallucinationDetected, RefusalToAnswer, BudgetExceeded, etc.) codifies failure patterns that every LLM application must handle but that no existing language formalizes.

SQL familiarity lowers the barrier. SQL is the most widely known programming language in the world. SPL's syntax leverages this familiarity, making agentic workflows accessible to data analysts, researchers, and domain experts who are fluent in SQL but not in Python framework APIs and LLM programming.

Human reviewability enables non-technical governance. With imperative frameworks, reviewing an agentic workflow requires tracing Python control flow, understanding framework-specific abstractions, and mentally simulating execution — skills that exclude the domain experts, analysts, and business stakeholders who are best positioned to validate whether the workflow is *correct for the problem*. An SPL script, by contrast, reads like a structured specification: SELECT declares what context flows in, GENERATE names the task, EVALUATE makes the branching condition explicit, and EXCEPTION names the failure modes. A business owner can read a 20-line SPL workflow and approve it without running it. This is a qualitatively different governance model — one that imperative code cannot provide regardless of documentation quality.

Declarative operations simplify the deployment lifecycle. SPL achieves not only declarative *implementation* but also declarative *operation*. The CLI flag surface — `--adapter` (LLM provider), `--model` (specific model), `--tools` (external connectors), `--datasets` (data sources) — acts as the environment configuration layer, analogous to environment variables in the 12-factor app methodology. The identical `.spl` file promotes unchanged through Dev, Test, and Production:

```
# Dev: local model, mock tools
spl run pipeline.spl --adapter ollama --model gemma3 \
    --tools search=mock --datasets corpus=sample.jsonl

# Test: cloud model, staging data
spl run pipeline.spl --adapter openrouter --model gpt-4o \
    --tools search=brave --datasets corpus=staging.jsonl

# Production: cost-optimised model, live data
spl run pipeline.spl --adapter anthropic --model claude-sonnet-4-6 \
```

```
--tools search=brave --datasets corpus=prod
```

No code changes, no re-review, no merge. The workflow logic is environment-agnostic by construction; only the runtime configuration changes. Imperative frameworks require either separate code paths per environment or fragile runtime injection — neither of which matches the simplicity of a flag swap.

7.2 Limitations

We acknowledge several limitations of the current work:

No optimizer yet. The optimization opportunities described in Section 5.3 are architectural possibilities enabled by the declarative form, not implemented features. The current system executes GENERATE statements sequentially without parallelism or model-routing optimizations.

Text2SPL (basic prototype). The natural language to SPL compiler (Layer 2 of the stack) has a working prototype via `spl text2spl`, limited to simple PROMPT pattern generation. Full WORKFLOW generation from natural language — the complete vision for Layer 2 — remains future work. When *claude_cli* adapter is used, the Text2SPL prototype can generate SPL code using a Claude model (e.g. Sonnet 4.6) demonstrating the potential for LLM-assisted SPL generation. However, the generated code requires review, refinement and validation.

Text-only modality. SPL 2.0 currently handles text input and output. Multi-modal support (images, audio, video) is planned for a future version.

No streaming (deliberate for batch targets). Adapters do not support token-by-token streaming of LLM output; results are returned in full at completion. This is a design choice, not an oversight. Momagrid's target workload is batch inference on consumer-grade GPU nodes: batch jobs do not require real-time token delivery, and removing streaming from the protocol significantly simplifies the hub-node communication layer. Streaming support, if added in a future release, would target dedicated GPU nodes with tighter latency guarantees rather than the decentralized inference grid.

Single-model judge in EVALUATE. By default, the same model that performs generation also acts as the quality judge in semantic EVALUATE conditions. If the generator is a weak model, the judge may be equally unreliable. A per-EVALUATE `USING MODEL` override mitigates this, but consistency of LLM-as-judge across repeated evaluations has not been formally measured.

WHILE termination not formally guaranteed. The WHILE loop relies on semantic conditions or iteration guards for termination. Non-termination is possible if the quality condition is never met and no explicit iteration cap is set. The runtime default of 10 iterations provides a practical safety net, but formal termination proofs are not provided.

Single-node baseline and thermal sensitivity. The cookbook experiments reported here were run on a single GTX 1080 Ti GPU (250W TDP) via Ollama with gemma3. The GTX 1080 Ti is known to **thermal throttle** under sustained load: after extended sequential inference, the GPU reaches its thermal ceiling (84°C) and reduces clock speed, causing per-recipe latency to increase progressively through a long batch run. This explains the wide variance between batch runs: a run that starts on a cool GPU (e.g., 1,493s for 30+ recipes) can be significantly faster than a run where throttling sets in mid-batch. Benchmarks should be interpreted with this variance in mind; controlled comparisons should either measure GPU temperature throughout or interleave cool-down periods between recipes.

Notably, CALL-based recipes (e.g., recipes 36–40, avg 4.8s) are structurally less affected by thermal throttling because they make fewer LLM calls per recipe — another practical benefit of the CALL/GENERATE separation.

Multi-node Momagrid execution was validated at two scales. A 3-node LAN grid (duck + dog + cat, GTX 1080 Ti) ran all recipes in 383.7s with an 89.2% pass rate (33/37); 4 recipes failed due to Ollama connection drops under sustained 90% GPU utilization across all three nodes simultaneously. A 4-node LAN grid (duck + dog + cat + goose, adding an RTX 4060 as a SILVER-tier node) achieved **37/37 recipes at 100% success** in 414.1s — a **3.6× speedup** over the single-node sequential baseline (1,493.4s) and **2.4**× over the best single-node run (1,000s). The marginal wall-time increase vs the 3-node run (383.7s → 414.1s) is within thermal variance and gated by the two longest inherently sequential recipes (Tree of Thought 383.2s, Reflection Agent 357.1s); the critical improvement is the elimination of all Ollama failures through distributed thermal load. The hub's tier-based dispatcher (GOLD $\succ$ SILVER) correctly routed the bulk of tasks to the three GOLD nodes, with the SILVER node absorbing tail tasks once all GOLD slots were occupied. See Appendix G for the full per-recipe timing table. Cross-model comparative benchmarks (llama3.2, mistral) remain planned for a follow-up evaluation.

memory.get() is read-only in the current implementation. The `memory.get('key')` context source is parsed, evaluated, and injected into the SELECT clause. However, write-back from within a WORKFLOW — persisting a computed value to the memory store for retrieval in a subsequent run — is not yet implemented. The `STORE` statement is defined in the grammar and accepted by the parser but is not executed by the runtime. Cross-run memory persistence is planned for a future release.

rag.query() requires a pre-built index. The `rag.query('text', top_k=N)` context source is functional only when a vector index has been pre-populated via `spl doc-rag add <document>`. The cookbook benchmark was run without a pre-built index; recipe 08 (RAG Query) confirmed the execution path and correct syntax handling, but the LLM received empty retrieval context because no documents were indexed. End-to-end retrieval validation over populated indices is deferred to a follow-up evaluation.

7.3 Comparison with Imperative Frameworks

The contrast between imperative and declarative approaches is most visible in the ReAct agent pattern. A LangGraph implementation requires 126 lines of Python (measured): defining state classes, creating graph nodes, specifying edge conditions, configuring the state machine, and handling errors through try/except blocks. The equivalent SPL 2.0 implementation is 43 lines with the workflow logic, branching conditions, and error handling all expressed declaratively.

More importantly, the declarative form makes the *intent* immediately visible. A reader of the SPL code can see at a glance: "this agent reasons, decides on an action, executes it, and accumulates history — with handlers for hallucination and max iterations." The same understanding from the LangGraph implementation requires tracing through multiple function definitions, graph edge conditions, and state class hierarchies.

7.4 Knowledge Transfer from SQL

SPL's syntax is designed for workforce transfer, not aesthetic novelty. Relational databases have produced, over five decades, a global community of practitioners who think fluently in the stored-

procedure mental model: declare inputs and outputs, assemble data with queries, branch on conditions, loop until a criterion is met, and handle named exceptions. These engineers built the data pipelines that power the world's enterprises. Their instincts — separate *what* from *how*, compose transformations, handle failure explicitly — are precisely the instincts an AI workflow engineer needs.

SPL 2.0's **WORKFLOW** is a stored procedure for LLM orchestration. The structural parallel is exact, not approximate:

SQL Stored Procedure	**SPL 2.0**
`CREATE PROCEDURE p(x IN TEXT)`	`WORKFLOW p INPUT @x TEXT`
`DECLARE v_result TEXT;`	`@result TEXT`
`SELECT ... INTO v_result FROM ...;`	`GENERATE f(@x) INTO @result;`
`WHILE condition LOOP ... END;`	`WHILE condition DO ... END;`
`IF cond THEN ... ELSE ... END IF;`	`EVALUATE WHEN ... THEN ... END;`
`EXCEPTION WHEN integrity_violation THEN`	`EXCEPTION WHEN HallucinationDetected THEN`
`COMMIT;`	`COMMIT @result;`

A data engineer familiar with SQL stored procedures can read and write SPL 2.0 WORKFLOW blocks with no tutorial. The block structure, exception taxonomy, variable lifecycle, and loop semantics are already in their muscle memory. What is new is only the backend: instead of `SELECT ... FROM orders` querying a relational table, `GENERATE answer(question)` queries a large language model. The declarative SELECT that assembles context mirrors the SELECT that assembles data — the difference is that the "rows" are prompt fragments, memories, and retrieved documents rather than database records.

This is the core accessibility claim of SPL 2.0: the millions of data engineers who currently maintain stored procedures, ETL pipelines, and reporting workflows are not learning a new paradigm to become AI engineers. They are applying a paradigm they have used for years to a new backend. The conceptual distance from SQL stored procedures to SPL 2.0 is a single substitution — *language model* for *relational engine* — while the surrounding architecture, discipline, and vocabulary remain unchanged.

The practical implication is day-one productivity. The cookbook of 30+ recipes spans patterns from simple single-step inference (Recipe 01) through multi-agent debate (Recipe 11), MapReduce over documents (Recipe 13), and CALL-based tool use with Python UDFs (Recipe 36). Every one of these recipes could be written, read, and maintained by a senior SQL developer encountering SPL 2.0 for the first time. No Python runtime knowledge, no framework-specific graph API, no imperative state machine design is required. The expressiveness of the language follows directly from this design choice: a language built on a fifty-year-old abstraction that has already proven it can express arbitrary data transformations at scale — now applied to agentic workflow orchestration over LLMs.

7.5 Future Work

Text2SPL compiler. Natural language to SPL 2.0 translation would complete the four-layer stack, enabling zero-syntax access to agentic workflows.

Query optimizer. A cost-based optimizer that rewrites execution plans — merging GENERATE calls, parallelizing independent steps, routing to optimal models — is the highest-impact future

contribution.

Logical-to-physical plan compiler. The current `spl explain` command serializes the AST as JSON. A full logical-to-physical plan compiler — analogous to SQL's EXPLAIN ANALYZE — would annotate each plan node with cost estimates, data-dependency edges, and adapter-routing hints, providing the foundation for the optimizer passes above.

Visual workflow builder. A graphical interface for composing WORKFLOW blocks would complement the text-based CLI for users who prefer visual programming.

Multi-modal I/O. Extending GENERATE to handle images, audio, and video inputs/outputs, with corresponding context source types.

Extended procedure library. A standard library of validated PROCEDURE implementations for common patterns (ReAct, Tree-of-Thought, Self-Refine, Debate, Map-Reduce, etc.) would accelerate adoption.

Formal verification. Proving properties of workflow specifications — termination guarantees, budget bounds, safety constraints — through static analysis.

Momagrid WAN deployment. LAN validation is complete (4-node grid, 3.6× throughput scaling, 100% recipe pass rate). WAN deployment — running Momagrid across geographically distributed, independently operated GPU nodes — is in progress. Cloud-hosted WAN evaluation is planned pending infrastructure funding.

SPL 2.0 Go port. The current SPL 2.0 runtime is implemented in Python. A port to Go is planned, motivated by Momagrid's Go implementation: a Go SPL runtime would enable single-binary deployment alongside the Momagrid hub and node, eliminating Python dependencies on GPU nodes and simplifying deployment on consumer-grade GPU nodes.

8. Conclusion

SPL 2.0 represents a paradigm shift from imperative to declarative orchestration of LLM-powered agentic workflows. By providing a SQL-like specification layer that separates *what* a workflow should accomplish from *how* it executes, the language addresses a fundamental fragmentation in the current ecosystem where building agents requires Python programming, API expertise, and framework-specific knowledge.

The language is built on six core primitives. Three are pre-existing in SPL 1.0 [5] (the declarative context management foundation):

- **PROMPT** — the named, executable unit binding a SELECT and a GENERATE into an atomic unit of LLM query
- **SELECT** prepares and assembles input context for the LLM
- **GENERATE** invokes the LLM and captures its output

Three are newly introduced in SPL 2.0 to enable agentic workflow orchestration:

- **EVALUATE** enables hybrid deterministic-semantic branching, where the LLM itself serves as a runtime judge.
- **WHILE** with semantic termination enables iteration based on output quality rather than arbitrary counters.

- **EXCEPTION** with a formal taxonomy of LLM failure modes enables structured error handling for hallucination, refusal, context overflow, and budget exhaustion.

We choose the declarative form not for the notational convenience, but to allow an ecosystem of compilers, optimizers, and reasoning tools to improve agentic workflows.

The reference implementation is publicly available at https://github.com/digital-duck/SPL20.

Acknowledgements

The author used Claude (Anthropic) [18], Gemini (Google) [19], and Z.ai (ZhiPu AI) [20] as AI assistants in this work: for literature search and related-work synthesis, for implementation and debugging of the codebase, and for editing portions of the manuscript. All technical claims, experimental results, and final text were reviewed and approved by the author, who takes full responsibility for the content.

References

[1] LangGraph. LangChain, Inc. https://github.com/langchain-ai/langgraph

[2] Wu, Q., et al. "AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation." arXiv:2308.08155, 2023.

[3] CrewAI. https://github.com/crewAIInc/crewAI

[4] Chamberlin, D. D. and Boyce, R. F. "SEQUEL: A Structured English Query Language." SIGFIDET Workshop, 1974.

[5] Gong, Wen G. "Structured Prompt Language: Declarative Context Management for LLMs." arXiv:2602.21257, 2026.

[6] Oracle Corporation. "PL/SQL User's Guide and Reference." Oracle Database Documentation.

[7] Beurer-Kellner, L., Fischer, M., and Vechev, M. "Prompting Is Programming: A Query Language for Large Language Models." arXiv:2212.06094, 2022.

[8] Khattab, O., et al. "DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines." arXiv:2310.03714, 2023.

[9] Zheng, L., et al. "SGLang: Efficient Execution of Structured Language Model Programs." arXiv:2312.07104, 2023.

[10] Agent Spec: Open Agent Specification. Oracle Labs. arXiv:2510.04173, 2025.

[11] SLANG: Standard Language for Agentic Workflows. https://dev.to/argonauta/why-we-need-a-standard-language-for-agentic-workflows-and-why-i-built-one-138

[12] "A Declarative Language for Building And Orchestrating LLM-Powered Agent Workflows." arXiv:2512.19769, 2025.

[13] "The Auton Agentic AI Framework: A Declarative Architecture for Specification, Governance, and Runtime Execution of Autonomous Agent Systems." arXiv:2602.23720, 2026.

[14] Gliozzo, A. M., Lee, J., and Defosse, N. “Agentics 2.0: Logical Transduction Algebra for Agentic Data Workflows.” arXiv:2603.04241, 2026.

[15] Gong, Wen G. “Momagrid: A Decentralized Inference Runtime with Linear Complexity via Semantic Chunking.” arXiv:submit/7373026, 2026.

[16] Yang, J., et al. “Auto-GPT: An Autonomous GPT-4 Experiment.” arXiv:2306.02224, 2023.

[17] PydanticAI. https://github.com/pydantic/pydantic-ai

[18] Claude. https://claude.ai/

[19] Gemini. https://gemini.google.com/

[20] Z.ai. https://chat.z.ai/

Appendix A: Complete SPL 2.0 Grammar (EBNF)

Productions are listed top-down. SPL 1.0 productions (unchanged) are marked †. New SPL 2.0 productions are unmarked. Keywords are case-insensitive. Comments use -- (end-of-line).

```
(* Top-level *)
program            := statement+

statement          := prompt_stmt
                    | workflow_stmt
                    | procedure_stmt
                    | create_function_stmt
                    | explain_stmt †
                    | execute_stmt †

(* PROMPT — atomic LLM query †  extended with CTE block *)
prompt_stmt        := 'PROMPT' identifier
                       cte_block?
                       select_clause
                       generate_clause
                       store_clause?

cte_block          := 'WITH' cte_def (',' cte_def)*
cte_def            := identifier 'AS' '(' prompt_stmt ')'

select_clause      := 'SELECT' select_item (',' select_item)*
select_item        := expression ('AS' identifier)?
                       ('LIMIT' integer 'TOKENS')?

generate_clause    := 'GENERATE' identifier '(' argument_list? ')'
                       ('INTO' '@' identifier)?
                       ('WITH' generate_option (',' generate_option)*)?
                       ('USING' 'MODEL' string)?

generate_option    := 'OUTPUT' 'BUDGET' integer 'TOKENS'
                    | 'TEMPERATURE' float

store_clause       := 'STORE' 'RESULT' 'AS' string
                       ('CACHE' 'FOR' integer)? †

(* WORKFLOW — procedural orchestration *)
workflow_stmt      := 'WORKFLOW' identifier
                       ('INPUT'  ':' param_list)?
                       ('OUTPUT' ':' param_list)?
                       security_block?
                       accounting_block?
                       do_block
                       'END'

(* PROCEDURE — reusable parameterized sub-workflow *)
procedure_stmt     := 'PROCEDURE' identifier
                       '(' param_list? ')'
                       ('RETURNS' identifier)?
                       security_block?
```

```
                        accounting_block?
                        do_block
                        'END'

(* DO block — shared by WORKFLOW and PROCEDURE *)
do_block           := 'DO' body_stmt+ exception_block?

body_stmt          := generate_into_stmt
                    | assignment_stmt
                    | evaluate_stmt
                    | while_stmt
                    | commit_stmt
                    | retry_stmt
                    | raise_stmt
                    | call_stmt
                    | select_into_stmt

generate_into_stmt := 'GENERATE' identifier '(' argument_list? ')'
                        'INTO' '@' identifier
                        ('WITH' generate_option (',' generate_option)*)?

assignment_stmt    := '@' identifier ':=' expression

evaluate_stmt      := 'EVALUATE' expression
                        when_clause+
                        ('OTHERWISE' body_stmt+)?
                        'END'

when_clause        := 'WHEN' condition 'THEN' body_stmt+

condition          := comparison_condition      (* deterministic numeric *)
                    | string_condition           (* deterministic string  *)
                    | semantic_condition         (* LLM-judged            *)
                    | expression

comparison_condition := comparison_op expression
comparison_op      := '>' | '<' | '>=' | '<=' | '=' | '!='

string_condition   := 'STARTSWITH' string_literal
                    | ('CONTAINS' '(' string_literal ')'
                        ('OR' 'CONTAINS' '(' string_literal ')')*)

semantic_condition := string_literal            (* plain quoted string triggers LLM judge *)

while_stmt         := 'WHILE' condition 'DO' body_stmt+ 'END'

commit_stmt        := 'COMMIT' expression ('WITH' option_list)?

retry_stmt         := 'RETRY' ('WITH' option_list)?

raise_stmt         := 'RAISE' exception_type string?

call_stmt          := 'CALL' identifier '(' named_arg_list? ')'
                        ('INTO' '@' identifier)?
```

```
select_into_stmt   := 'SELECT' select_item (',' select_item)*
                       'INTO' '@' identifier

exception_block    := 'EXCEPTION' exception_handler+

exception_handler  := 'WHEN' exception_type 'THEN' body_stmt+

exception_type     := 'HallucinationDetected'  | 'RefusalToAnswer'
                    | 'ContextLengthExceeded'   | 'ModelOverloaded'
                    | 'BudgetExceeded'          | 'QualityBelowThreshold'
                    | 'MaxIterationsReached'    | 'NodeUnavailable'
                    | 'GenerationError'         | 'GenerationFailed'

(* CREATE FUNCTION — schema / constant definition *)
create_function_stmt := 'CREATE' 'FUNCTION' identifier
                        '(' param_list? ')'
                        'RETURNS' identifier
                        'AS' '$$' text '$$'

(* Expressions *)
expression         := literal
                    | param_ref
                    | dotted_name
                    | function_call
                    | binary_op
                    | select_source

select_source      := system_role | context_ref | rag_query | memory_get

system_role        := 'system_role' '(' string ')'
context_ref        := 'context' '.' identifier
rag_query          := 'rag' '.' 'query' '(' string (',' named_arg)* ')'
memory_get         := 'memory' '.' 'get' '(' string ')'

literal            := string_literal | integer_literal | float_literal
param_ref          := '@' identifier
dotted_name        := identifier ('.' identifier)+
function_call      := identifier '(' argument_list? ')'
binary_op          := expression ('+' | '-' | '*' | '/') expression

(* Parameter and argument lists *)
param_list         := param_def (',' param_def)*
param_def          := '@'? identifier identifier? ('DEFAULT' expression)?

argument_list      := argument (',' argument)*
argument           := named_arg | expression
named_arg          := identifier '=' expression

named_arg_list     := named_arg (',' named_arg)*

option_list        := option (',' option)*
option             := identifier '=' expression
```

```
(* Metadata blocks *)
security_block     := 'SECURITY' ':' 'CLASSIFICATION' ':' identifier
accounting_block   := 'ACCOUNTING' ':' 'BILLABLE_TO' ':' string
                       ('BUDGET_LIMIT' ':' integer identifier)?

(* Lexical *)
identifier         := [a-zA-Z_][a-zA-Z0-9_]*
string_literal     := '"' .* '"' | "'" .* "'"
integer_literal    := [0-9]+
float_literal      := [0-9]+ '.' [0-9]+
```

Type system. SPL 2.0 supports five primitive types: `TEXT`, `NUMBER` (integer or float), `BOOLEAN`, `JSON`, `CHUNK` (token-bounded text segment). Composite types: `LIST` (ordered sequence), `MAP` (key-value). Type annotations are optional; untyped parameters default to `TEXT`.

Condition type detection rule. If a `WHEN` condition is a quoted string literal, the runtime treats it as a semantic condition and invokes the LLM as a judge. Otherwise (numeric comparison, boolean expression) it is evaluated deterministically with no LLM call. This detection is unambiguous and requires no explicit annotation.

Appendix B: IR Specification (JSON Schema)

The SPL 2.0 Intermediate Representation (IR) is a JSON-serialized AST produced by `spl explain`. It is the portable format consumed by the Momagrid runtime, plan visualization tools, and future optimizer passes.

Program envelope

```
{
  "type": "Program",
  "statements": [ <statement>+ ]
}
```

Prompt plan (produced by `spl explain` for PROMPT statements)

```
{
  "type": "PromptPlan",
  "prompt_name": "string",
  "model": "string | null",
  "total_budget": "integer | null",
  "output_budget": "integer | null",
  "total_input_tokens": "integer",
  "buffer_tokens": "integer",
  "estimated_cost": "float",
  "optimizations": [ "string" ],
  "warnings": [ "string" ],
  "steps": [
    {
      "operation": "string",
      "source": "string",
      "alias": "string | null",
      "estimated_tokens": "integer",
      "limit_tokens": "integer | null",
```

```
      "allocated_tokens": "integer",
      "compressed": "boolean",
      "compression_ratio": "float",
      "cache_status": "string",
      "priority": "integer"
    }
  ]
}
```

Workflow plan (produced by `spl explain` for WORKFLOW statements)

```
{
  "type": "WorkflowPlan",
  "workflow_name": "string",
  "inputs": [ { "name": "string", "type": "string", "default": "<expr> | null" } ],
  "outputs": [ { "name": "string", "type": "string" } ],
  "security": { "classification": "string" } | null,
  "accounting": { "billable_to": "string" } | null,
  "total_estimated_llm_calls": "integer",
  "total_estimated_tokens": "integer",
  "steps": [ <workflow_step> ],
  "exception_handlers": [ <workflow_step> ]
}
```

Workflow step (recursive — branches and substeps are the same type)

```
{
  "step_type": "GenerateInto | Evaluate | While | Commit | Retry | Call | Assignment | ...",
  "description": "string",
  "estimated_llm_calls": "integer",
  "estimated_tokens": "integer",
  "substeps": [ <workflow_step> ] | null,
  "branches": [
    {
      "condition": "string",
      "condition_type": "semantic | deterministic | hybrid",
      "steps": [ <workflow_step> ]
    }
  ] | null,
  "exception_handlers": [ <workflow_step> ] | null
}
```

Condition node

```
{ "type": "SemanticCondition",   "value": "string" }
{ "type": "ComparisonCondition", "operator": "> | < | >= | <= | = | !=", "right": <expr> }
{ "type": "Condition",           "left": <expr>, "operator": "string", "right": <expr> }
```

Expression nodes (used in step arguments, conditions, assignments)

```
{ "type": "Literal",    "value": "string | integer | float" }
{ "type": "ParamRef",   "name": "string" }
{ "type": "ContextRef", "field_name": "string" }
{ "type": "DottedName", "parts": ["string"], "full_name": "string" }
{ "type": "SystemRole", "description": "string" }
```

```
{ "type": "RagQuery",   "query": "string", "source": "string | null", "top_k": "integer" }
{ "type": "MemoryGet",  "key": "string" }
{ "type": "FunctionCall","name": "string", "arguments": [ <expr> ] }
{ "type": "BinaryOp",   "op": "+ | - | * | /", "left": <expr>, "right": <expr> }
```

The IR is produced by the `spl2.ir` module and consumed by the executor (`spl2.executor`) and Momagrid adapter (`spl2.adapters.momagrid`). All field names are stable across SPL 2.0 patch releases; additions are backward-compatible.

Appendix C: Hub ↔ Node Protocol Specification

The Momagrid Hub-Node protocol is a REST-over-HTTPS request-response protocol. The SPL 2.0 Momagrid adapter (`spl2.adapters.momagrid`) implements the client (Hub) side using async `httpx`.

Task lifecycle

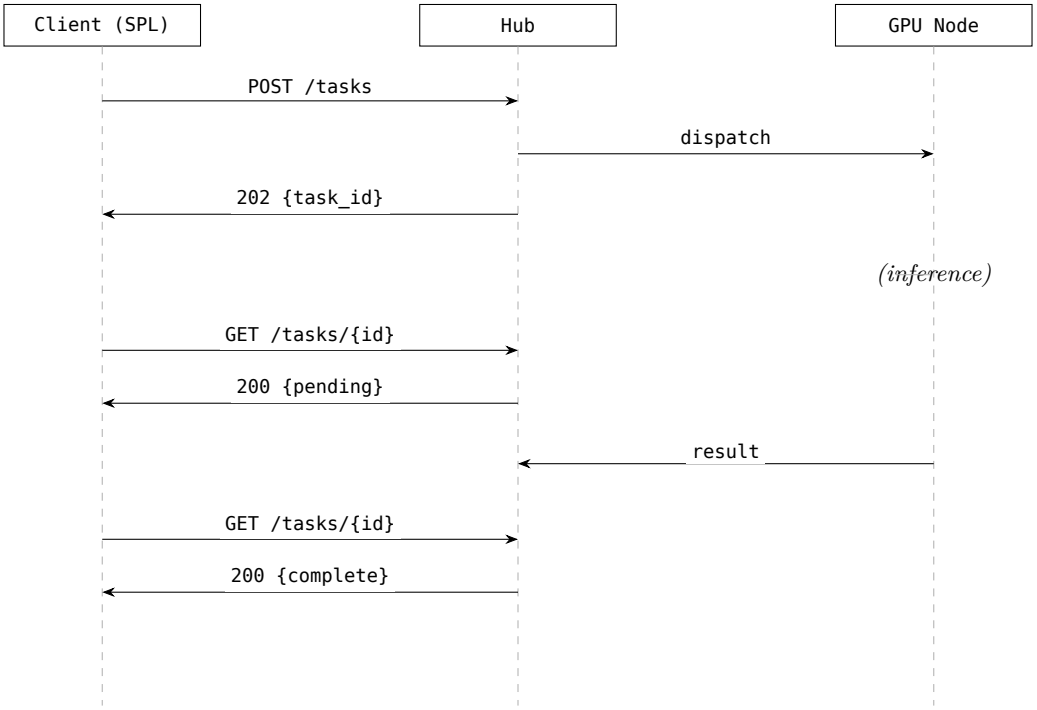

Figure C-1: Task lifecycle — REST polling sequence between Client, Hub, and GPU Node.

POST /tasks — Submit inference request

```
Request body:
{
  "model":         "string",
  "prompt":        "string (plaintext or encrypted ciphertext)",
  "system_role":   "string | null",
  "temperature":   "float | null",
  "max_tokens":    "integer | null",
  "node_requirements": {
    "min_tier":    "integer –(05)",
    "min_vram_gb": "float"
  },
```

```
  "session_key_id": "string | null"
}

Response 202:
{
  "task_id": "uuid",
  "status":  "pending",
  "eta_ms":  "integer | null"
}
```

GET /tasks/{task_id} — Poll task status

```
Response 200 (pending):
{ "task_id": "uuid", "status": "pending | running" }

Response 200 (complete):
{
  "task_id":    "uuid",
  "status":     "complete",
  "result":     "string (plaintext or encrypted ciphertext)",
  "tokens_in":  "integer",
  "tokens_out": "integer",
  "latency_ms": "integer",
  "node_id":    "string"
}

Response 200 (failed):
{
  "task_id": "uuid",
  "status":  "failed",
  "error":   "string",
  "code":    "node_unavailable | model_not_found | overloaded | timeout"
}
```

Error code → SPL exception mapping

HTTP / error code	SPL Exception raised
`overloaded` or HTTP 429	`ModelOverloaded`
`node_unavailable` or HTTP 503	`NodeUnavailable`
`timeout`	`NodeUnavailable`
`context_exceeded`	`ContextLengthExceeded`
any other `failed` status	`GenerationError`

Encryption handshake (zero-trust mode)

When `SECURITY: CLASSIFICATION: confidential` is set in the SPL workflow, the adapter negotiates an ephemeral session key before task submission:

1. Client generates an ephemeral key pair (X25519).
2. Client sends the public key to `POST /session/init`.
3. Hub returns an encrypted session key, signed by the Hub's master key.
4. Subsequent task payloads (`prompt`, `result`) are encrypted with AES-256-GCM under the session key.
5. Session keys expire after 1 hour or 1000 tasks, whichever comes first.

In plaintext mode (default for development), the `session_key_id` field is omitted and `prompt/result` are transmitted unencrypted.

Accounting integration

When `ACCOUNTING: BILLABLE_TO: 'project-id'` is set, the Hub records a Moma Points debit against the project account on task completion. If `BUDGET_LIMIT` is exceeded, the Hub returns HTTP 402 and the adapter raises `BudgetExceeded`.

Appendix D: Exception Type Hierarchy

Exception	Trigger	Default Recovery
HallucinationDetected	Output fails fact-checking or contains fabricated information	Regenerate with temperature = 0.1
RefusalToAnswer	Model declines due to safety filters or content policy	Log refusal, COMMIT with status = 'refused'
ContextLengthExceeded	Combined input exceeds model's context window	Compress context, RETRY
ModelOverloaded	Provider returns 429 or capacity error	Exponential backoff, RETRY
BudgetExceeded	Token budget (input or output) exhausted	COMMIT best available result
QualityBelowThreshold	Generated output scores below configured minimum	Regenerate or escalate to larger model
MaxIterationsReached	WHILE loop reached iteration guard limit	COMMIT current result with status = 'partial'
NodeUnavailable	Momagrid GPU node unreachable or timed out *(Momagrid-specific)*	Route to alternative node, RETRY

Appendix E: Momagrid Runtime Architecture

SPL 2.0's adapter pattern supports a unique deployment model: decentralized execution on independently operated GPU nodes. Momagrid targets a fundamentally different niche from commercial cloud APIs: **batch inference for users who prioritize accessibility and cost over latency.** The design goal is low- or no-cost agentic AI inference for researchers, students, and developers in resource-constrained environments — complementary to, not competitive with, services like Anthropic or Google.

Current status: Research prototype. LAN grid testing is complete. WAN deployment simulation is planned. A research grant proposal has been submitted to Mozilla Foundation. Latency variability is an acknowledged trade-off for the accessibility target.

A multi-configuration LAN scaling experiment was conducted, running 37 cookbook recipes on a home LAN grid (3× GTX 1080 Ti, all via Ollama, gemma3). Results are shown in Table E.1.

Table E.1: Momagrid LAN Scaling Results

Config	Nodes	Workers	Pass Rate	Wall Clock	Speedup
1-GPU	1	1	37/37 (100%)	1,197.6s	1.0×
2-GPU	2	5	37/37 (100%)	660.4s	1.8×
3-GPU	3	5	33/37 (89%)	548.6s	2.2×
3-GPU	3	10	33/37 (89%)	383.7s	3.1×
4-GPU	4	10	37/37 (100%)	414.1s	2.9×

Node hardware: nodes 1–3 are GTX 1080 Ti (11 GB VRAM, GOLD tier); node 4 is RTX 4060 (8 GB VRAM, SILVER tier). All nodes run Ollama with gemma3.

The 4 failures in the 3-node runs (recipes with heavy concurrent load) were caused by Ollama dropping connections under sustained ∼90% GPU utilization (EOF / connection refused), not by SPL or Momagrid logic errors. Adding a 4th node (RTX 4060, SILVER tier) distributed the thermal load across 4 GPUs, eliminating the connection drops and achieving **37/37 (100%) pass rate** — the first perfect run on the Momagrid adapter. Wall clock (414.1s) is within scheduler overhead of the 3-node best (383.7s); the gain is reliability, not raw speed. The same `.spl` files and a single `--adapter momagrid` flag were the only difference from the single-node baseline.

Adapter	Target Use Case	Latency Profile
Ollama	Local dev, single GPU	Low (local)
OpenRouter	Cloud, 100+ models	Medium (API)
Claude CLI	Anthropic direct	Medium (API)
Momagrid	**Batch jobs, low-cost inference**	**Variable (decentralized)**

E.1 The Hub

The Momagrid Hub serves as the coordination layer for decentralized workflows:

- **Task submission**: Receives SPL-generated inference requests with model and hardware requirements.
- **Node selection**: Matches tasks to available GPU nodes based on min_tier, min_vram_gb, and geographic constraints.
- **State management**: Maintains global workflow context, enabling checkpointing and recovery for long-running workflows.

E.2 GPU Nodes

Each GPU node in the Momagrid network operates as a stateless executor:

- Receives encrypted task payloads from the Hub.
- Decrypts, performs inference, and encrypts the result.
- Reports completion status and resource consumption metrics.

Node selection criteria (min_tier, min_vram_gb) are specified in SPL 2.0's SECURITY metadata block, allowing workflows to express hardware requirements declaratively.

E.3 Compression, Encryption, Accounting

Three cross-cutting concerns are built into the Momagrid protocol:

Compression: Two levels — *semantic compression* reduces context length before transmission (preserving meaning while reducing tokens), and *transmission compression* (zstd, gzip) reduces network bandwidth.

Encryption: A zero-trust model with a three-level key hierarchy (Master → Session → Data) ensures that no single party — including the Hub — has access to both the prompt and the response in plaintext.

Accounting: The Moma Points system tracks resource consumption, with event logging and a reputation system for node operators. SPL 2.0's ACCOUNTING metadata block integrates with this system: `BUDGET_LIMIT: 1000 moma_points` enforces cost caps at the workflow level.

Appendix F: Validation of SPL Recipes

F.1 Recipe Design Philosophy

Each cookbook recipe is designed as a self-contained demonstration: - A single `.spl` file with inline comments explaining the pattern. - Usage examples as CLI commands in the file header. - Progressive complexity from atomic queries to multi-agent workflows. - Catalog-driven organization via `cookbook_catalog.json` and `run_all.py` batch runner.

F.2 Recipe Catalog

Table F-1: SPL 2.0 Cookbook Catalog

#	Recipe	Category	Pattern	LOC	Key SPL Features
01	Hello World	basics	PROMPT (atomic)	4	system_role, GENERATE
02	Ollama Proxy	basics	PROMPT + params	5	context.param, model override
03	Multilingual	basics	PROMPT + 2 params	6	Multiple context bindings
04	Model Showdown	benchmarking	WORKFLOW + multi-model	30	Cross-model comparison, 4 calls
05	Self-Refine	agentic	WORKFLOW + WHILE	35	Iterative improvement loop
06	ReAct Agent	agentic	WORKFLOW + EVALUATE	43	Tool-call branching, history
07	Safe Generation	safety	WORKFLOW + EXCEPTION	37	Quality scoring, error recovery
08	RAG Query	retrieval	PROMPT + rag.query	7	Vector search context
09	Chain of Thought	reasoning	WORKFLOW (3-step)	14	Sequential GENERATE pipeline
10	Batch Test	benchmarking	WORKFLOW + parallel CTE	50	Multi-recipe, multi-model regression
11	Debate Arena	multi-agent	WORKFLOW + adversarial	45	Multi-perspective, judge eval
12	Plan and Execute	agentic	WORKFLOW + re-planning	50	Dynamic plan adjustment
13	Map-Reduce Summarizer	application	WORKFLOW + chunking	40	Chunked GENERATE, reduce step
14	Multi-Agent Collab.	multi-agent	PROCEDURE + CALL	45	Role specialization, delegation
15	Code Review	application	WORKFLOW + multi-pass	50	Severity scoring, synthesis
16	Reflection Agent	agentic	WORKFLOW + meta-cognition	45	Confidence scoring, correction

#	Recipe	Category	Pattern	LOC	Key SPL Features
17	Tree of Thought	reasoning	WORKFLOW + branching	50	Path scoring, verification
18	Guardrails Pipeline	safety	WORKFLOW + multi-gate	55	PII detection, content filtering
19	Memory Conversation	application	WORKFLOW + memory	40	memory.get, fact extraction
20	Ensemble Voting	multi-agent	WORKFLOW + consensus	45	Multi-candidate scoring, voting
21	Multi-Model Pipeline	agentic	WORKFLOW + adapters	50	Cross-adapter routing, fallback
22	Text2SPL Demo	basics	WORKFLOW	—	Natural language → SPL compilation
23	Structured Output	basics	PROMPT + CREATE FUNCTION	36	JSON schema-constrained generation
24	Few-Shot Prompting	basics	PROMPT + CREATE FUNCTION	47	In-context learning, domain examples
25	Nested Procedures	agentic	WORKFLOW + PROCEDURE + CALL	55	PROCEDURE composability, delegation
26	Prompt A/B Test	benchmarking	WORKFLOW + comparison	40	Multi-prompt comparison, scoring
27	Data Extraction	application	PROMPT + CREATE FUNCTION	32	Schema-constrained field extraction
28	Customer Support Triage	application	WORKFLOW + EVALUATE + EXCEPTION	60	Classify → route → draft pipeline
29	Meeting Notes to Actions	application	WORKFLOW + EXCEPTION	55	Transcript → structured TODO list
30	Code Generator + Tests	application	WORKFLOW + EVALUATE + EXCEPTION	65	Generate → review → test pipeline
31	Sentiment Pipeline	application	WORKFLOW + multi-step	50	Multi-label sentiment + reporting
32	Socratic Tutor	application	WORKFLOW + iterative	55	Adaptive question-driven learning
33	Interview Simulator	application	WORKFLOW + PROCEDURE	60	Role-play, scoring, feedback
34	Progressive Summarizer	application	WORKFLOW + chunking	45	Hierarchical summarization
35	Hypothesis Tester	reasoning	WORKFLOW + EVALUATE	55	Hypothesis → test plan → evidence
36	Tool Use	agentic	WORKFLOW + CALL	32	CALL UDFs for math; GENERATE for narrative

#	Recipe	Category	Pattern	LOC	Key SPL Features
37	Headline News	application	WORKFLOW + CALL	28	Fetch RSS, summarize, format report
38	Human in the Loop	agentic	WORKFLOW + CALL + EVALUATE	38	Human approval gate before COMMIT
39	Doc Q&A	retrieval	WORKFLOW + CALL	30	Load document, extract answers via GENERATE
40	Portability Test	benchmarking	WORKFLOW + stdlib	25	to_float from stdlib; runs on any adapter

F.3 End-to-End Demonstrations

Experimental setup. All recipes were executed via the batch runner (`python cookbook/run_all.py --adapter ollama --model gemma3`) on a single workstation equipped with a GTX 1080 Ti GPU running the Ollama inference server. No cloud API credits were consumed. Recipes 36–40 include the Tool-Use (CALL UDFs), Human-in-Loop, Doc Q&A, and Portability Test (stdlib) patterns. The full run completed in **1,493.4 seconds** — 40/40 recipes passing, including recipe 22 (Text2SPL Demo) which now runs successfully for the first time.

Table F-2: Batch Run Results — gemma3 / Ollama / GTX 1080 Ti

#	Recipe	Category	Type	Latency
01	Hello World	basics	PROMPT	5.3s
02	Ollama Proxy	basics	PROMPT	0.7s
03	Multilingual Greeting	basics	PROMPT	0.6s
04	Model Showdown	benchmarking	WORKFLOW	26.2s
05	Self-Refine	agentic	WORKFLOW	3.9s
06	ReAct Agent	agentic	WORKFLOW	27.3s
07	Safe Generation	safety	WORKFLOW	10.3s
08	RAG Query	retrieval	PROMPT	6.1s
09	Chain of Thought	reasoning	WORKFLOW	23.5s
10	Batch Test	benchmarking	WORKFLOW	5.1s
11	Debate Arena	multi-agent	WORKFLOW	91.8s
12	Plan and Execute	agentic	WORKFLOW	89.9s
13	Map-Reduce Summarizer	application	WORKFLOW	10.6s
14	Multi-Agent Collab.	multi-agent	WORKFLOW	76.0s
15	Code Review	application	WORKFLOW	47.5s
16	Reflection Agent	agentic	WORKFLOW	132.3s
17	Tree of Thought	reasoning	WORKFLOW	80.6s
18	Guardrails Pipeline	safety	WORKFLOW	8.2s
19	Memory Conversation	application	WORKFLOW	3.9s
20	Ensemble Voting	multi-agent	WORKFLOW	61.2s
21	Multi-Model Pipeline	agentic	WORKFLOW	32.5s
22	Text2SPL Demo	basics	WORKFLOW	13.8s
23	Structured Output	basics	PROMPT	1.4s
24	Few-Shot Prompting	basics	PROMPT	0.8s
25	Nested Procedures	agentic	WORKFLOW	32.6s
26	Prompt A/B Test	benchmarking	WORKFLOW	75.3s
27	Data Extraction	application	PROMPT	1.3s

#	Recipe	Category	Type	Latency
28	Customer Support Triage	application	WORKFLOW	55.8s
29	Meeting Notes to Actions	application	WORKFLOW	36.2s
30	Code Generator + Tests	application	WORKFLOW	34.4s
31	Sentiment Pipeline	application	WORKFLOW	83.1s
32	Socratic Tutor	application	WORKFLOW	90.6s
33	Interview Simulator	application	WORKFLOW	126.7s
34	Progressive Summarizer	application	WORKFLOW	10.2s
35	Hypothesis Tester	reasoning	WORKFLOW	36.9s
36	Tool-Use / Function-Call	agentic	WORKFLOW	8.7s
37	Headline News Aggregator	application	WORKFLOW	131.6s
38	Human-in-the-Loop Appr. Gate	agentic	WORKFLOW	3.7s
39	Document Q&A	retrieval	WORKFLOW	1.7s
40	Portability Test	benchmarking	WORKFLOW	5.0s

Summary: 40/40 recipes passed (100%). Total wall time: **1,493.4 seconds**. This is the first run in which recipe 22 (Text2SPL Demo) passes — previously always skipped as "planned/not yet released." Five new recipes (36–40) all pass, with the CALL-based recipes completing in 1.7s–8.7s, demonstrating that delegating computation to deterministic Python UDFs adds negligible overhead. The fastest workflow is Recipe 39 (Doc Q&A, 1.7s); the slowest is Recipe 16 (Reflection Agent, 132.3s).

Aggregate statistics:

Table F-3: Aggregate Statistics by Recipe Type

Segment	Count	Avg Latency
PROMPT recipes (7)	7	2.3s
WORKFLOW recipes (33)	33	44.8s
CALL-based new recipes (36–40)	5	9.9s
ReAct Agent (recipe 06, peak)	1	27.3s
Total	**40**	—

Notable observations:

- **Recipe 22 (Text2SPL Demo, 13.8s):** First successful run of the text-to-SPL compiler in a batch context, completing the "natural language → executable SPL" demonstration layer. Previously marked SKIP in all prior runs.
- **Tool Use (recipe 36, 8.7s):** All arithmetic (`sum_values`, `percentage_change`, `format_currency`) delegated to deterministic Python UDFs via CALL; the LLM is invoked exactly once for the narrative summary. Zero hallucination risk on numerical values.
- **Human-in-the-Loop (recipe 38, 3.7s):** Fastest WORKFLOW recipe. Approval gate operates via simulated CALL with no LLM involvement in the decision, demonstrating the CALL/GENERATE boundary for human-in-the-loop patterns.
- **Doc Q&A (recipe 39, 1.7s):** Fastest recipe overall. Document loading via CALL is deterministic; only the answer extraction step calls the LLM.
- **Portability Test (recipe 40):** Uses only stdlib functions (`to_float`) with no `--tools` flag, demonstrating that the Standard Library eliminates common boilerplate tools.
- **ReAct Agent (recipe 06, 27.3s):** Executed multiple LLM calls autonomously — the WHILE loop drives a genuine tool-reasoning loop without programmer intervention.

- **Safe Generation (recipe 07, 10.3s):** Demonstrates EVALUATE with semantic conditions (`WHEN 'high_quality'`, `WHEN 'acceptable'`) — the LLM serves as its own quality judge with no numeric parsing required.
- **Reflection Agent (recipe 16, 132.3s):** Slowest workflow, reflecting its deep meta-cognition loop.
- **Interview Simulator (recipe 33, 126.7s):** Deep multi-turn role-play simulation within a single `spl run` command.
- **Headline News Aggregator (recipe 37, 131.6s):** Slow due to fetching and summarizing multiple live news sources via CALL UDFs before generating the final report.
- **Debate Arena (recipe 11, 91.8s)** and **Ensemble Voting (recipe 20, 61.2s):** Both generated rich multi-perspective outputs within a single `spl run` command, with no Python orchestration code.
- **Nested Procedures (recipe 25, 32.6s):** Demonstrates three-level PROCEDURE composability — WORKFLOW calling PROCEDURE calling PROCEDURE — without any imperative orchestration code.

Appendix G: Full Experiment Results

Run configuration: - Model: gemma3 (via Ollama) - Hardware: GTX 1080 Ti (11 GB VRAM), single node - Runner: `python cookbook/run_all.py --adapter ollama --model gemma3` - Total wall time: 1,493.4s - Result: 40/40 SUCCESS

Note on timing variance: Batch run wall time varies between runs primarily due to LLM output length variance (different outputs per run) and GPU thermal throttling. The GTX 1080 Ti throttles clock speed under sustained load once it reaches 84°C, causing progressive latency increases through a long batch run. This run (1,493.4s) was conducted on a cool GPU and reflects unthrottled performance. Earlier benchmark runs showed higher totals for the same recipe set; these are attributed to thermal throttling rather than software regression. Controlled benchmarking should monitor GPU temperature or interleave cool-down periods.

A multi-GPU experiment on the Momagrid 2-node and 3-node configurations (see Appendix E, Table E.1 for full results). The best run: 3 GPU nodes, 10 workers, 37 recipes in 383.7s (3.1× speedup over the 1-node baseline). All recipes pass on the single-node benchmark including recipe 22 (Text2SPL Demo), which ran successfully for the first time.

Pass/fail breakdown by category:

Table G-1: Pass/Fail Breakdown by Category

Category	Recipes	Passed	Failed	Notes
basics	6	6	0	Recipes 01, 02, 03, 22 (*first pass*), 23, 24
agentic	8	8	0	Recipes 05, 06, 12, 16, 21, 25, 36 (Tool Use), 38 (Human-in-Loop)
reasoning	3	3	0	Recipes 09, 17, 35
safety	2	2	0	Recipes 07, 18
retrieval	2	2	0	Recipes 08, 39 (Doc Q&A, 1.7s)
multi-agent	3	3	0	Recipes 11, 14, 20
application	12	12	0	Recipes 13, 15, 19, 27, 28, 29, 30, 31, 32, 33, 34, 37 (Headline News, 131.6s)
benchmarking	4	4	0	Recipes 04, 10, 26, 40 (Portability Test, 5.0s)
Total	**40**	**40**	**0**	

Appendix H: Declarative vs. Imperative — A Case Study

The self-refining agent (draft → critique → revise) is the canonical benchmark for agentic frameworks. It exercises every layer of SPL syntax: atomic LLM queries (**GENERATE**), iterative refinement (**WHILE**), semantic quality gating (**EVALUATE**), and graceful degradation (**EXCEPTION**).

H.1 SPL 2.0 Implementation

The complete implementation (43 lines) illustrates all core grammar constructs:

```
-- Illustrates: WORKFLOW, GENERATE, WHILE, EVALUATE, EXCEPTION
WORKFLOW self_refine
  INPUT:  @task text, @max_iterations int DEFAULT 5
  OUTPUT: @result text
DO
  GENERATE draft(@task) INTO @current

  WHILE @iteration < @max_iterations DO
    GENERATE critique(@current) INTO @feedback

    EVALUATE @feedback
      WHEN 'satisfactory' THEN
        COMMIT @current
      OTHERWISE
        GENERATE revised(@current, @feedback) INTO @current
    END
  END

  COMMIT @current
EXCEPTION
  WHEN MaxIterationsReached THEN
    COMMIT @current WITH status = 'partial'
END
```

Primitives illustrated:

- **GENERATE** — invokes the LLM and captures output into a variable (`INTO @var`)
- **WHILE** — iterates until a counter or semantic condition terminates the loop
- **EVALUATE** — LLM-as-judge branch: routes on semantic content of `@feedback`
- **EXCEPTION** — handles LLM-specific failure modes with structured recovery

H.2 LOC Comparison

Table H.1 compares implementations across SPL, LangGraph, AutoGen, and CrewAI. Source code: `cookbook/05_self_refine/{langgraph,autogen,crewai}/`.

Table H.1: LOC Comparison — Self-Refining Agent

Framework	LOC	Notes
SPL 2.0	**43**	Full WORKFLOW with WHILE loop, EVALUATE gate, LOGGING, and EXCEPTION handler. No Python required.
LangGraph	120	TypedDict state, four node functions, graph wiring (`add_node`/`add_edge`), conditional edge, compile + invoke boilerplate.

Framework	LOC	Notes
AutoGen	95	Two ConversableAgents, `llm_config` dict, `initiate_chat` with `max_turns`, chat-history post-processing to extract drafts and feedback.
CrewAI	103	Agent/Task/Crew construction per iteration (no native loop primitive), explicit Python `for` loop and `if` replacing SPL's WHILE/EVALUATE.

Counts include all non-blank lines (imports, comments, boilerplate).

H.3 Full SPL 2.0 Source (cookbook/05_self_refine/self_refine.spl)

```
WORKFLOW self_refine
  INPUT:  @task text, @max_iterations int DEFAULT 5
  OUTPUT: @result text
DO
  GENERATE draft(@task) INTO @current

  WHILE @iteration < @max_iterations DO
    GENERATE critique(@current) INTO @feedback
    EVALUATE @feedback
      WHEN 'satisfactory' THEN
        COMMIT @current
      OTHERWISE
        GENERATE revised(@current, @feedback) INTO @current
    END
  END

  COMMIT @current
EXCEPTION
  WHEN MaxIterationsReached THEN
    COMMIT @current WITH status = 'partial'
END
```

Key observations:

- SPL 2.0 requires **2.2–2.8× fewer lines** than imperative frameworks for the same pattern (43 vs. 95–120 LOC measured across LangGraph, AutoGen, and CrewAI).
- The EVALUATE primitive eliminates the need for regex parsing or numeric scoring logic to detect loop termination — a significant source of complexity in imperative implementations.
- The EXCEPTION block captures MaxIterationsReached with a single clause; imperative frameworks require explicit counter checks and try/except blocks scattered across multiple functions.
- The same `self_refine.spl` file runs on Ollama, Anthropic, OpenRouter, or Momagrid by changing a single CLI flag; imperative implementations are typically coupled to one provider SDK.

Beyond LOC, two structural advantages do not appear in the table but are significant in practice:

Reviewability. An imperative implementation (120–150 lines of Python across multiple classes and callbacks) requires execution to verify correctness. The 22-line SPL workflow is self-describing: a domain expert with no Python knowledge can read it, understand every step, and sign off. This enables a governance model where business stakeholders review and approve agentic workflows before deployment — structurally impossible with imperative code.

Declarative operations. The same `self_refine.spl` runs in any environment by swapping CLI flags:

```
# Dev        --adapter ollama  --model gemma3
# Test       --adapter openrouter --model gpt-4o
# Production --adapter anthropic --model claude-sonnet-4-6
```

No code changes across environments. Imperative implementations require environment-specific initialization, conditional adapter logic, or separate configuration files that must be kept in sync with the code.